

Tutorials

- 1: [Hello Minikube](#)
- 2: [Learn Kubernetes Basics](#)
 - 2.1: [Create a Cluster](#)
 - 2.1.1: [Using Minikube to Create a Cluster](#)
 - 2.1.2: [Interactive Tutorial - Creating a Cluster](#)
 - 2.2: [Deploy an App](#)
 - 2.2.1: [Using kubectl to Create a Deployment](#)
 - 2.2.2: [Interactive Tutorial - Deploying an App](#)
 - 2.3: [Explore Your App](#)
 - 2.3.1: [Viewing Pods and Nodes](#)
 - 2.3.2: [Interactive Tutorial - Exploring Your App](#)
 - 2.4: [Expose Your App Publicly](#)
 - 2.4.1: [Using a Service to Expose Your App](#)
 - 2.4.2: [Interactive Tutorial - Exposing Your App](#)
 - 2.5: [Scale Your App](#)
 - 2.5.1: [Running Multiple Instances of Your App](#)
 - 2.5.2: [Interactive Tutorial - Scaling Your App](#)
 - 2.6: [Update Your App](#)
 - 2.6.1: [Performing a Rolling Update](#)
 - 2.6.2: [Interactive Tutorial - Updating Your App](#)
- 3: [Configuration](#)
 - 3.1: [Example: Configuring a Java Microservice](#)
 - 3.1.1: [Externalizing config using MicroProfile, ConfigMaps and Secrets](#)
 - 3.1.2: [Interactive Tutorial - Configuring a Java Microservice](#)
 - 3.2: [Configuring Redis using a ConfigMap](#)
- 4: [Stateless Applications](#)
 - 4.1: [Exposing an External IP Address to Access an Application in a Cluster](#)
 - 4.2: [Example: Deploying PHP Guestbook application with Redis](#)
- 5: [Stateful Applications](#)
 - 5.1: [StatefulSet Basics](#)
 - 5.2: [Example: Deploying WordPress and MySQL with Persistent Volumes](#)
 - 5.3: [Example: Deploying Cassandra with a StatefulSet](#)
 - 5.4: [Running ZooKeeper, A Distributed System Coordinator](#)
- 6: [Clusters](#)
 - 6.1: [Restrict a Container's Access to Resources with AppArmor](#)
 - 6.2: [Restrict a Container's Syscalls with seccomp](#)
- 7: [Services](#)
 - 7.1: [Using Source IP](#)

of the Kubernetes documentation contains tutorials. A tutorial shows how to accomplish a goal rather than a single [task \(/docs/tasks/\)](#). Typically a tutorial has several sections, each of which has steps. Before walking through each tutorial, you may want to bookmark the [Standardized Glossary \(ance/glossary/\)](#) page for later references.

[Kubernetes Basics \(/docs/tutorials/kubernetes-basics/\)](/docs/tutorials/kubernetes-basics/) is an in-depth interactive tutorial that helps you understand the Kubernetes system and try out some basic Kubernetes features.

Introduction to Kubernetes (edX) (<https://www.edx.org/course/introduction-kubernetes-linuxfoundationx-#>).

[minikube \(/docs/tutorials/hello-minikube/\)](https://minikube.github.io/docs/tutorials/hello-minikube/)

uration

le: [Configuring a Java Microservice \(/docs/tutorials/configuration/configure-java-microservice/\)](/docs/tutorials/configuration/configure-java-microservice/)
uring Redis Using a ConfigMap [\(/docs/tutorials/configuration/configure-redis-using-configmap/\)](/docs/tutorials/configuration/configure-redis-using-configmap/)

less Applications

ng an External IP Address to Access an Application in a Cluster [\(/docs/tutorials/stateless-tion/expose-external-ip-address/\)](/docs/tutorials/stateless-tion/expose-external-ip-address/)
le: [Deploying PHP Guestbook application with Redis \(/docs/tutorials/stateless-tion/guestbook/\)](/docs/tutorials/stateless-tion/guestbook/)

ul Applications

lSet Basics [\(/docs/tutorials/stateful-application/basic-stateful-set/\)](/docs/tutorials/stateful-application/basic-stateful-set/)
le: [WordPress and MySQL with Persistent Volumes \(/docs/tutorials/stateful-application/mysql-ress-persistent-volume/\)](/docs/tutorials/stateful-application/mysql-ress-persistent-volume/)
le: [Deploying Cassandra with Stateful Sets \(/docs/tutorials/stateful-application/cassandra/\)](/docs/tutorials/stateful-application/cassandra/)
g ZooKeeper, A CP Distributed System [\(/docs/tutorials/stateful-application/zookeeper/\)](/docs/tutorials/stateful-application/zookeeper/)

CS

nor [\(/docs/tutorials/clusters/apparmor/\)](/docs/tutorials/clusters/apparmor/)
ip [\(/docs/tutorials/clusters/seccomp/\)](/docs/tutorials/clusters/seccomp/)

ES

source IP [\(/docs/tutorials/services/source-ip/\)](/docs/tutorials/services/source-ip/)

next

like to write a tutorial, see [Content Page Types \(/docs/contribute/style/page-content-types/\)](/docs/contribute/style/page-content-types/) for about the tutorial page type.

1 - Hello Minikube

This tutorial shows you how to run a sample app on Kubernetes using minikube and Katacoda. Katacoda provides a free, in-browser Kubernetes environment.

Note: You can also follow this tutorial if you've installed minikube locally. See [minikube start \(https://minikube.sigs.k8s.io/docs/start/\)](https://minikube.sigs.k8s.io/docs/start/) for installation instructions.

Objectives

- Deploy a sample application to minikube.
- Run the app.
- View application logs.

Before you begin

This tutorial provides a container image that uses NGINX to echo back all the requests.

Create a minikube cluster

1. Click **Launch Terminal**

Launch Terminal

Note: If you installed minikube locally, run `minikube start`. Before you run `minikube dashboard`, you should open a new terminal, start `minikube dashboard` there, and then switch back to the main terminal.

2. Open the Kubernetes dashboard in a browser:

```
minikube dashboard
```

3. Katacoda environment only: At the top of the terminal pane, click the plus sign, and then click **Select port to view on Host 1**.

4. Katacoda environment only: Type `30000`, and then click **Display Port**.

Note:

The `dashboard` command enables the dashboard add-on and opens the proxy in the default web browser. You can create Kubernetes resources on the dashboard such as Deployment and Service.

If you are running in an environment as root, see [Open Dashboard with URL](#).

By default, the dashboard is only accessible from within the internal Kubernetes virtual network. The `dashboard` command creates a temporary proxy to make the dashboard accessible from outside the Kubernetes virtual network.

To stop the proxy, run `Ctrl+C` to exit the process. After the command exits, the dashboard remains running in the Kubernetes cluster. You can run the `dashboard` command again to create another proxy to access the dashboard.

Open Dashboard with URL

If you don't want to open a web browser, run the dashboard command with the `--url` flag to emit a URL:

```
minikube dashboard --url
```

Create a Deployment

A Kubernetes *Pod* (</docs/concepts/workloads/pods/>) is a group of one or more Containers, tied together for the purposes of administration and networking. The Pod in this tutorial has only one Container. A Kubernetes *Deployment* (</docs/concepts/workloads/controllers/deployment/>) checks on the health of your Pod and restarts the Pod's Container if it terminates. Deployments are the recommended way to manage the creation and scaling of Pods.

1. Use the `kubectl create` command to create a Deployment that manages a Pod. The Pod runs a Container based on the provided Docker image.

```
kubectl create deployment hello-node --image=k8s.gcr.io/echoserver:1.4
```

2. View the Deployment:

```
kubectl get deployments
```

The output is similar to:

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
hello-node	1/1	1	1	1m

3. View the Pod:

```
kubectl get pods
```

The output is similar to:

NAME	READY	STATUS	RESTARTS	AGE
hello-node-5f76cf6ccf-br9b5	1/1	Running	0	1m

4. View cluster events:

```
kubectl get events
```

5. View the `kubectl` configuration:

```
kubectl config view
```

Note: For more information about `kubectl` commands, see the [kubectl overview](/docs/reference/kubectl/overview/) (</docs/reference/kubectl/overview/>).

Create a Service

By default, the Pod is only accessible by its internal IP address within the Kubernetes cluster. To make the `hello-node` Container accessible from outside the Kubernetes virtual network, you have to expose the Pod as a Kubernetes *Service* (</docs/concepts/services-networking/service/>).

1. Expose the Pod to the public internet using the `kubectl expose` command:

```
kubectl expose deployment hello-node --type=LoadBalancer --port=8080
```

The `--type=LoadBalancer` flag indicates that you want to expose your Service outside of the cluster.

The application code inside the image `k8s.gcr.io/echoserver` only listens on TCP port 8080. If you used `kubectl expose` to expose a different port, clients could not connect to that other port.

2. View the Service you created:

```
kubectl get services
```

The output is similar to:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-node	LoadBalancer	10.108.144.78	<pending>	8080:30369/TCP	21s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	23m

On cloud providers that support load balancers, an external IP address would be provisioned to access the Service. On minikube, the `LoadBalancer` type makes the Service accessible through the `minikube service` command.

3. Run the following command:

```
minikube service hello-node
```

4. Katacoda environment only: Click the plus sign, and then click **Select port to view on Host 1**.

5. Katacoda environment only: Note the 5-digit port number displayed opposite to `8080` in services output. This port number is randomly generated and it can be different for you. Type your number in the port number text box, then click Display Port. Using the example from earlier, you would type `30369`.

This opens up a browser window that serves your app and shows the app's response.

Enable addons

The minikube tool includes a set of built-in [addons \(/docs/concepts/cluster-administration/addons/\)](/docs/concepts/cluster-administration/addons/) that can be enabled, disabled and opened in the local Kubernetes environment.

1. List the currently supported addons:

```
minikube addons list
```

The output is similar to:

```
addon-manager: enabled
dashboard: enabled
default-storageclass: enabled
efk: disabled
freshpod: disabled
gvisor: disabled
helm-tiller: disabled
ingress: disabled
ingress-dns: disabled
logviewer: disabled
metrics-server: disabled
nvidia-driver-installer: disabled
nvidia-gpu-device-plugin: disabled
registry: disabled
registry-creds: disabled
storage-provisioner: enabled
storage-provisioner-gluster: disabled
```

2. Enable an addon, for example, `metrics-server` :

```
minikube addons enable metrics-server
```

The output is similar to:

```
The 'metrics-server' addon is enabled
```

3. View the Pod and Service you created:

```
kubectl get pod,svc -n kube-system
```

The output is similar to:

NAME	READY	STATUS	RESTARTS	AGE
pod/coredns-5644d7b6d9-mh9ll	1/1	Running	0	34m
pod/coredns-5644d7b6d9-pqd2t	1/1	Running	0	34m
pod/metrics-server-67fb648c5	1/1	Running	0	26s
pod/etcd-minikube	1/1	Running	0	34m
pod/influxdb-grafana-b29w8	2/2	Running	0	26s
pod/kube-addon-manager-minikube	1/1	Running	0	34m
pod/kube-apiserver-minikube	1/1	Running	0	34m
pod/kube-controller-manager-minikube	1/1	Running	0	34m
pod/kube-proxy-rnlps	1/1	Running	0	34m
pod/kube-scheduler-minikube	1/1	Running	0	34m
pod/storage-provisioner	1/1	Running	0	34m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/metrics-server	ClusterIP	10.96.241.45	<none>	80/TCP
service/kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP,53/TCP
service/monitoring-grafana	NodePort	10.99.24.54	<none>	80:30002/TCP
service/monitoring-influxdb	ClusterIP	10.111.169.94	<none>	8083/TCP,8086/TCP

4. Disable metrics-server :

```
minikube addons disable metrics-server
```

The output is similar to:

```
metrics-server was successfully disabled
```

Clean up

Now you can clean up the resources you created in your cluster:

```
kubectl delete service hello-node
kubectl delete deployment hello-node
```

Optionally, stop the Minikube virtual machine (VM):

```
minikube stop
```

Optionally, delete the Minikube VM:

```
minikube delete
```

What's next

- Learn more about [Deployment objects \(/docs/concepts/workloads/controllers/deployment/\)](/docs/concepts/workloads/controllers/deployment/).
- Learn more about [Deploying applications \(/docs/tasks/run-application/run-stateless-application-deployment/\)](/docs/tasks/run-application/run-stateless-application-deployment/).
- Learn more about [Service objects \(/docs/concepts/services-networking/service/\)](/docs/concepts/services-networking/service/).

2 - Learn Kubernetes Basics

Basics

ides a walkthrough of the basics of the Kubernetes cluster orchestration
dule contains some background information on major Kubernetes
cepts, and includes an interactive online tutorial. These interactive
manage a simple cluster and its containerized applications for yourself.

tive tutorials, you can learn to:

ntainerized application on a cluster.

ployment.

containerized application with a new software version.


containerized application.

Katacoda to run a virtual terminal in your web browser that runs
l-scale local deployment of Kubernetes that can run anywhere. There's
l any software or configure anything; each interactive tutorial runs
ur web browser itself.

What Kubernetes do for you?


o services, users expect applications to be available 24/7, and developers
new versions of those applications several times a day. Containerization
ftware to serve these goals, enabling applications to be released and
downtime. Kubernetes helps you make sure those containerized
where and when you want, and helps them find the resources and tools
k. Kubernetes is a production-ready, open source platform designed
umulated experience in container orchestration, combined with best-of-
the community.

Basics Modules




</kubernetes-cluster/cluster-intro/>
Kubernetes cluster

</kubernetes-cluster/cluster-intro/>




[\(/docs/tutorials/kubernetes-basics/deploy-app/deploy-intro/\)](/docs/tutorials/kubernetes-basics/deploy-app/deploy-intro/)
2. Deploy an app

[\(/docs/tutorials/kubernetes-basics/deploy-app/deploy-intro/\)](/docs/tutorials/kubernetes-basics/deploy-app/deploy-intro/)




[\(/docs/tutorials/kubernetes-basics/explore/explore-intro/\)](/docs/tutorials/kubernetes-basics/explore/explore-intro/)
3. Explore your app

[\(/docs/tutorials/kubernetes-basics/explore/explore-intro/\)](/docs/tutorials/kubernetes-basics/explore/explore-intro/)




</kubernetes-expose-intro/>
Expose app publicly

</kubernetes-expose-intro/>



[\(/docs/tutorials/kubernetes-basics/scale/scale-intro/\)](/docs/tutorials/kubernetes-basics/scale/scale-intro/)
5. Scale up your app

[\(/docs/tutorials/kubernetes-basics/scale/scale-intro/\)](/docs/tutorials/kubernetes-basics/scale/scale-intro/)



[\(/docs/tutorials/kubernetes-basics/update/update-intro/\)](/docs/tutorials/kubernetes-basics/update/update-intro/)
6. Update your app

[\(/docs/tutorials/kubernetes-basics/update/update-intro/\)](/docs/tutorials/kubernetes-basics/update/update-intro/)

2.1 - Create a Cluster

2.1.1 - Using Minikube to Create a Cluster

a Kubernetes cluster is.

Minikube is.

ernetes cluster using an online terminal.

Clusters


ordinates a highly available cluster of computers that are work as a single unit. The abstractions in Kubernetes allow you to containerized applications to a cluster without tying them specifically to machines. To make use of this new model of deployment, applications need to be packaged in a way that decouples them from individual machines. Containerized applications are more portable than in past deployment models, where applications were deployed directly onto specific machines as packages deeply integrated with the operating system. **Kubernetes automates the distribution and scheduling of containers across a cluster in a more efficient way.**

Kubernetes is an open-source platform and is production-ready.

A cluster consists of two types of resources:

Control Plane coordinates the cluster
Worker Nodes are the workers that run applications

Diagram



The Control Plane is responsible for managing the cluster. The Control Plane manages all activities in your cluster, such as scheduling applications, maintaining applications' desired state, scaling applications, and rolling out updates.

A Node is a physical computer that serves as a worker machine in the cluster. Each node has a Kubelet, which is an agent for managing the node and communicating with the Kubernetes control plane. Nodes also have tools for handling container operations, such as Docker or containerd. A Kubernetes cluster that handles production traffic should have a minimum of three nodes.

When you deploy applications on Kubernetes, you tell the control plane to start containers. The control plane schedules the containers to run on the appropriate nodes. **The nodes communicate with the control plane**

Summary:

- Kubernetes cluster
- Minikube

Kubernetes is a production-grade, open-source platform that orchestrates the placement (scheduling) and execution of application containers within and across computer clusters.

Control Planes manage the cluster and the nodes that are used to host the running applications.

netes API (/docs/concepts/overview/kubernetes-api/),

I plane exposes. End users can also use the Kubernetes API
t with the cluster.

ster can be deployed on either physical or virtual machines.
th Kubernetes development, you can use Minikube. Minikube
ubernetes implementation that creates a VM on your local
loys a simple cluster containing only one node. Minikube is
x, macOS, and Windows systems. The Minikube CLI provides
ing operations for working with your cluster, including start,
delete. For this tutorial, however, you'll use a provided online
rikube pre-installed.

ow what Kubernetes is, let's go to the online tutorial and start

ive Tutorial > _(/docs/tutorials/kubernetes-basics/create-cluster/cluster-interactive/)

2.1.2 - Interactive Tutorial - Creating a Cluster

Kubernetes cluster

inutes

teractive scenario is to deploy a local development Kubernetes cluster

a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER sign will execute that command in the

ratulations on completing the
file

hear what you thought of the module

ise@gmail.com (not shared) [Switch accounts](#)



Kubernetes cluster

inutes

teractive scenario is to deploy a local development Kubernetes cluster

a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER sign will execute that command in the

ratulations on completing the
file

hear what you thought of the module

ise@gmail.com (not shared) [Switch accounts](#)



Kubernetes cluster

minutes

Interactive scenario is to deploy a local Kubernetes cluster using minikube

a pre-configured Linux environment that can be used as can type commands). Clicking on the blocks of code ⌘ sign will execute that command in the terminal.

Prerequisites on setting the module

Share what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

Kubernetes cluster

minutes

Interactive scenario is to deploy a local Kubernetes cluster using minikube

a pre-configured Linux environment that can be used as can type commands). Clicking on the blocks of code ⌘ sign will execute that command in the terminal.

Prerequisites on setting the module

Share what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

Kubernetes cluster

minutes

Interactive scenario is to deploy a local Kubernetes cluster using minikube

a pre-configured Linux environment that can be used as can type commands). Clicking on the blocks of code R sign will execute that command in the terminal.

Exercises on
Installing the module

Share what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

Kubernetes cluster

minutes

Interactive scenario is to deploy a local Kubernetes cluster using minikube

a pre-configured Linux environment that can be used as can type commands). Clicking on the blocks of code R sign will execute that command in the terminal.

Exercises on
Installing the module

Share what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

Kubernetes cluster

minutes

Interactive scenario is to deploy a local Kubernetes cluster using minikube

a pre-configured Linux environment that can be used as a regular console (you can type commands). Clicking on the blocks of code followed by the ENTER sign will execute that command in the terminal.

Congratulations on completing the module

Share what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

Kubernetes cluster

minutes

Interactive scenario is to deploy a local development Kubernetes cluster

a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER sign will execute that command in the

Congratulations on completing the module

Share what you thought of the module

oise@gmail.com (not shared) [Switch accounts](#)



[Tutorials/kubernetes-basics/](#)

Module 2 > [\(/docs/tutorials/kubernetes-basics/deploy-app/deploy-intro/\)](#)

2.2 - Deploy an App

2.2.1 - Using kubectl to Create a Deployment

t application Deployments.
r first app on Kubernetes with kubectl.

Deployments

running Kubernetes cluster, you can deploy your applications on top of it. To do so, you create a Kubernetes configuration. The Deployment instructs Kubernetes how to create instances of your application. Once you've created a Deployment, the Kubernetes control plane schedules the application instances on individual Nodes in the cluster.

When new instances are created, a Kubernetes Deployment continuously monitors those instances. If the Node hosting an instance is down or is deleted, the Deployment controller replaces the instance on another Node in the cluster. **This provides a self-healing mechanism to address machine failure or maintenance.**


In the traditional application world, installation scripts would often be used to start an application. Before Kubernetes, they did not allow recovery from machine failure. By both creating application instances and keeping them running across Nodes, Deployments provide a fundamentally different approach to application management.

Summary:

- Deployments
- Kubectl

A Deployment is responsible for creating and updating instances of your application

Deploying your first app on Kubernetes



and manage a Deployment by using the Kubernetes command **kubectl**. Kubectl uses the Kubernetes API to interact with the cluster. In this module, you'll learn the most common Kubectl commands for creating Deployments that run your applications on a Kubernetes cluster.

When you create a Deployment, you'll need to specify the container image for the application and the number of replicas that you want to run. You can learn more about scaling a Deployment later by updating your Deployment; Modules [5](#) ([Kubernetes-basics/scale/scale-intro/](#)) and [6](#) ([Kubernetes-basics/update/update-intro/](#)) of the bootcamp can scale and update your Deployments.

In this Deployment, you'll use a hello-node application packaged in a container that uses NGINX to echo back all the requests. (If you didn't have a hello-node application and deploying it using a container,

Applications need to be packaged into one of the supported container formats in order to be deployed on Kubernetes

irst by following the instructions from the [Hello Minikube](#)
[tutorials/hello-minikube/](#))).

ow what Deployments are, let's go to the online tutorial and
pp!

ive Tutorial > [\(/docs/tutorials/kubernetes-basics/deploy-app/deploy-interactive/\)](#)

2.2.2 - Interactive Tutorial - Deploying an App

execution unit of a Kubernetes application. Each Pod represents a part of a
inning on your cluster. [Learn more about Pods](#)
[orkloads/pods/](#)).

1 app

inutes

**enario is to help you deploy your first app on Kubernetes using
earn the basics about kubectl cli and how to interact with your**

a pre-configured Linux environment that can be used as a regular console (you can type
on the blocks of code followed by the ENTER sign will execute that command in the

ratulations on completing the
ile

hear what you thought of the module

oise@gmail.com (not shared) [Switch accounts](#)



1 app

inutes

**enario is to help you deploy your first app on Kubernetes using
earn the basics about kubectl cli and how to interact with your**

a pre-configured Linux environment that can be used as a regular console (you can type
on the blocks of code followed by the ENTER sign will execute that command in the

ratulations on completing the
ile

hear what you thought of the module

oise@gmail.com (not shared) [Switch accounts](#)



1 app

minutes

enario is to help you deploy your first app on kubectl. You will learn the basics about w to interact with your application.

a pre-configured Linux environment that can be used as can type commands). Clicking on the blocks of code R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

1 app

minutes

enario is to help you deploy your first app on kubectl. You will learn the basics about w to interact with your application.

a pre-configured Linux environment that can be used as can type commands). Clicking on the blocks of code R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

1 app

minutes

enario is to help you deploy your first app on kubectl. You will learn the basics about w to interact with your application.

a pre-configured Linux environment that can be used as can type commands). Clicking on the blocks of code R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

1 app

minutes

enario is to help you deploy your first app on kubectl. You will learn the basics about w to interact with your application.

a pre-configured Linux environment that can be used as can type commands). Clicking on the blocks of code R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

1 app

minutes

enario is to help you deploy your first app on
kubectl. You will learn the basics about
w to interact with your application.

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

1 app

minutes

enario is to help you deploy your first app on Kubernetes using
earn the basics about kubectl cli and how to interact with your

a pre-configured Linux environment that can be used as a regular console (you can type
on the blocks of code followed by the ENTER sign will execute that command in the

ratulations on completing the file

hear what you thought of the module

ise@gmail.com (not shared) [Switch accounts](#)



odule 1 (</docs/tutorials/kubernetes-basics/create-cluster/cluster-intro/>)

[utorials/kubernetes-basics/](/docs/tutorials/kubernetes-basics/))

odule 3 > (</docs/tutorials/kubernetes-basics/explore/explore-intro/>)

2.3 - Explore Your App

2.3.1 - Viewing Pods and Nodes

t Kubernetes Pods.
t Kubernetes Nodes.
ot deployed applications.

ies Pods

d a Deployment in Module 2 (</docs/tutorials/kubernetes-p/deploy-intro/>), Kubernetes created a **Pod** to host your rce. A Pod is a Kubernetes abstraction that represents a group pplication containers (such as Docker), and some shared ise containers. Those resources include:

age, as Volumes
, as a unique cluster IP address
i about how to run each container, such as the container on or specific ports to use

application-specific "logical host" and can contain different iners which are relatively tightly coupled. For example, a Pod th the container with your Node.js app as well as a different eds the data to be published by the Node.js webserver. The od share an IP Address and port space, are always co-located d, and run in a shared context on the same Node.


nic unit on the Kubernetes platform. When we create a (Kubernetes, that Deployment creates Pods with containers pposed to creating containers directly). Each Pod is tied to the scheduled, and remains there until termination (according to deletion. In case of a Node failure, identical Pods are ier available Nodes in the cluster.

Summary:

- Pods
- Nodes
- Kubectl main commands

A Pod is a group of one or more application containers (such as Docker) and includes shared storage (volumes), IP address and information about how to run them.

review




s on a **Node**. A Node is a worker machine in Kubernetes and irtual or a physical machine, depending on the cluster. Each d by the control plane. A Node can have multiple pods, and the ol plane automatically handles scheduling the pods across cluster. The control plane's automatic scheduling takes into able resources on each Node.

s Node runs at least:

rocess responsible for communication between the control plane and the Node; it manages the Pods and the unning on a machine.
runtime (like Docker) responsible for pulling the container a registry, unpacking the container, and running the

Containers should only be scheduled together in a single Pod if they are tightly coupled and need to share resources such as disk.

Overview



Working with kubectl

For more information on <https://kubernetes.io/docs/tutorials/kubernetes-basics/deploy-app/deploy-intro/>, you can use the kubectl command-line interface. You'll continue to use it in Module 3 to learn about deployed applications and their environments. The most common actions can be done with the following kubectl commands:

- `kubectl get` - list resources
- `kubectl describe` - show detailed information about a resource
- `kubectl logs` - print the logs from a container in a pod
- `kubectl exec` - execute a command on a container in a pod

Use these commands to see when applications were deployed, what resources they use, where they are running and what their status is.

For more information about our cluster components and the command line, see the [kubectl documentation](#).

A node is a worker machine in Kubernetes and may be a VM or physical machine, depending on the cluster. Multiple Pods can run on one Node.

[Interactive Tutorial](https://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-interactive/) > [_ \(/docs/tutorials/kubernetes-basics/explore/explore-interactive/\)](https://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-interactive/)

2.3.2 - Interactive Tutorial - Exploring Your App

our app

inutes

ou will learn how to troubleshoot Kubernetes the kubectl get, describe, logs and exec

a pre-configured Linux environment that can be used as can type commands). Clicking on the blocks of code R sign will execute that command in the terminal.

atulations on
eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

our app

inutes

ou will learn how to troubleshoot Kubernetes the kubectl get, describe, logs and exec

a pre-configured Linux environment that can be used as can type commands). Clicking on the blocks of code R sign will execute that command in the terminal.

atulations on
eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

our app

inutes

ou will learn how to troubleshoot Kubernetes applications using the
be, logs and exec commands.

a pre-configured Linux environment that can be used as a regular console (you can type
on the blocks of code followed by the ENTER sign will execute that command in the

ratulations on completing the
le

hear what you thought of the module

oise@gmail.com (not shared) [Switch accounts](#)



our app

inutes

ou will learn how to troubleshoot Kubernetes
the kubectl get, describe, logs and exec

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R sign will execute that command in the terminal.

ratulations on
eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

our app

minutes

you will learn how to troubleshoot Kubernetes
the kubectl get, describe, logs and exec

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
⌘R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

our app

minutes

you will learn how to troubleshoot Kubernetes
the kubectl get, describe, logs and exec

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
⌘R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

our app

minutes

You will learn how to troubleshoot Kubernetes applications using the kubectl get, describe, logs and exec commands.

a pre-configured Linux environment that can be used as a regular console (you can type commands on the blocks of code followed by the ENTER sign will execute that command in the terminal).

Congratulations on completing the module

Share what you thought of the module

prase@gmail.com

(not shared)

Switch accounts

our app

minutes

You will learn how to troubleshoot Kubernetes applications using the kubectl get, describe, logs and exec commands.

a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER sign will execute that command in the


Congratulations on completing the module

Share what you thought of the module

prase@gmail.com

(not shared)

Switch accounts



[Module 2](#) [../docs/tutorials/kubernetes-basics/deploy-app/deploy-intro/](#)

[tutorials/kubernetes-basics/](#)

[Module 4 >](#) [../docs/tutorials/kubernetes-basics/expose/expose-intro/](#)

2.4 - Expose Your App Publicly

2.4.1 - Using a Service to Expose Your App

to a Service in Kubernetes

How labels and LabelSelector objects relate to a Service

Application outside a Kubernetes cluster using a Service

f Kubernetes Services

(/docs/concepts/workloads/pods/) are mortal. Pods in fact docs/concepts/workloads/pods/pod-lifecycle/). When a ;, the Pods running on the Node are also lost. A ReplicaSet workloads/controllers/replicaset/) might then dynamically back to desired state via creation of new Pods to keep your ng. As another example, consider an image-processing aplicas. Those replicas are exchangeable; the front-end ot care about backend replicas or even if a Pod is lost and aid, each Pod in a Kubernetes cluster has a unique IP address, : same Node, so there needs to be a way of automatically ges among Pods so that your applications continue to

rnetes is an abstraction which defines a logical set of Pods which to access them. Services enable a loose coupling between Pods. A Service is defined using YAML (preferred), configuration/overview/#general-configuration-tips) or JSON, as objects. The set of Pods targeted by a Service is usually *LabelSelector* (see below for why you might want a Service; `selector` in the spec).

nd has a unique IP address, those IPs are not exposed outside
ut a Service. Services allow your applications to receive traffic.
xposed in different ways by specifying a `type` in the

- `ClusterIP` - Exposes the Service on an internal IP in the cluster. Makes the Service only reachable from within the cluster.
- `NodePort` - Exposes the Service on the same port of each selected Node in the cluster using NAT. Makes a Service accessible from outside the cluster using `NodePort`. Superset of `ClusterIP`.
- `LoadBalancer` - Creates an external load balancer in the current cloud (if available) and assigns a fixed, external IP to the Service. Superset of `NodePort`.

re - Maps the Service to the contents of the externalName
oo.bar.example.com), by returning a CNAME record with its
roxying of any kind is set up. This type requires v1.7 or higher
; , or CoreDNS version 0.0.8 or higher.

about the different types of Services can be found in the [/docs/tutorials/services/source-ip/](#) tutorial. Also see [cations with Services \(/docs/concepts/services-ect-applications-service\)](#).

that there are some use cases with Services that involve not
r in the spec. A Service created without `selector` will also
r responding Endpoints object. This allows users to manually
specific endpoints. Another possibility why there may be no
e strictly using `type: ExternalName`.

d Labels

traffic across a set of Pods. Services are the abstraction that and replicate in Kubernetes without impacting your

Summary

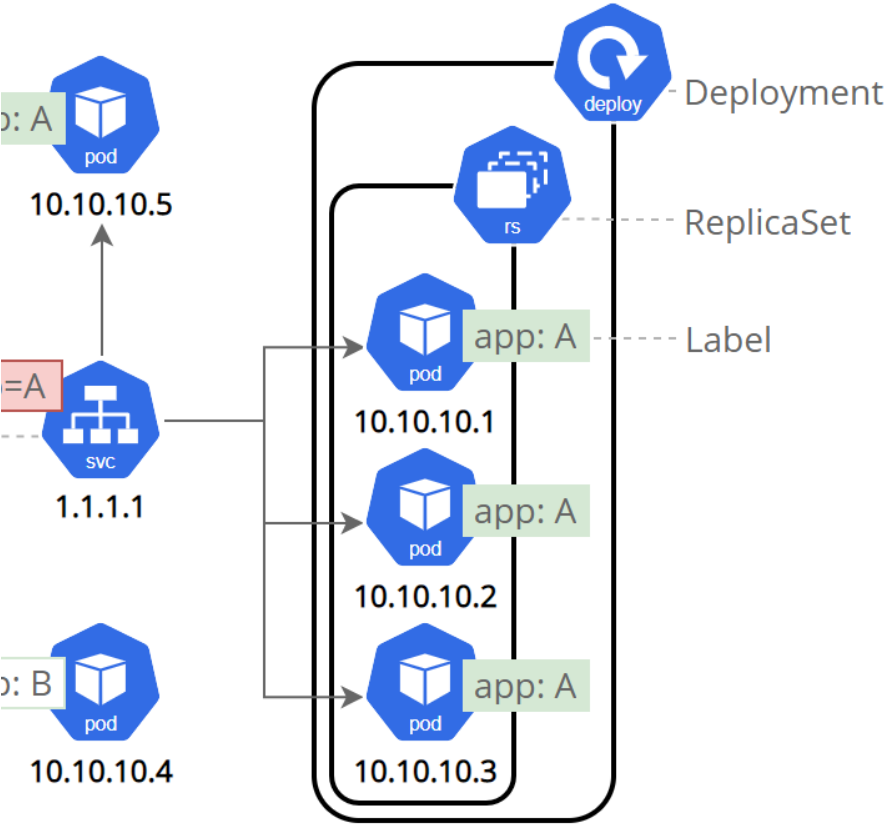
- Exposing Pods to external traffic
- Load balancing traffic across multiple Pods
- Using labels

A Kubernetes Service is an abstraction layer which defines a logical set of Pods and enables external traffic exposure, load balancing and service discovery for those Pods.

Discovery and routing among dependent Pods (such as the
sidecar components in an application) is handled by Kubernetes

group of Pods using labels and selectors
([overview/working-with-objects/labels](#)), a grouping primitive
for performing operations on objects in Kubernetes. Labels are key/value
pairs attached to objects and can be used in any number of ways:

- Grouping objects for development, test, and production
- Adding identification tags
- Selecting objects using tags



Labels are attached to objects at creation time or later on. They can be used to select objects. Let's expose our application now using a Service and a Deployment.

[Expose Tutorial](#) > [./docs/tutorials/kubernetes-basics/expose/expose-interactive/](#)

2.4.2 - Interactive Tutorial - Exposing Your App

our app publicly

inutes

ou will learn how to expose Kubernetes applications outside the cluster
expose command. You will also learn how to view and apply labels to
ubectl label command.

a pre-configured Linux environment that can be used as a regular console (you can type
on the blocks of code followed by the ENTER sign will execute that command in the

ratulations on completing the
ile

hear what you thought of the module

oise@gmail.com (not shared) [Switch accounts](#)



our app publicly

inutes

ou will learn how to expose Kubernetes applications outside the cluster
expose command. You will also learn how to view and apply labels to
ubectl label command.

a pre-configured Linux environment that can be used as a regular console (you can type
on the blocks of code followed by the ENTER sign will execute that command in the

ratulations on completing the
ile

hear what you thought of the module

oise@gmail.com (not shared) [Switch accounts](#)



our app publicly

inutes

**ou will learn how to expose Kubernetes
le the cluster using the kubectl expose
ll also learn how to view and apply labels to
ubectl label command.**

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

our app publicly

inutes

**ou will learn how to expose Kubernetes
le the cluster using the kubectl expose
ll also learn how to view and apply labels to
ubectl label command.**

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

our app publicly

minutes

**You will learn how to expose Kubernetes
in the cluster using the kubectl expose
command. You will also learn how to view and apply labels to
a pod using the kubectl label command.**

This is a pre-configured Linux environment that can be used as a playground for learning Kubernetes (you can type commands). Clicking on the blocks of code will execute that command in the terminal.

Exercises on Deploying the module

Share what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

our app publicly

minutes

**You will learn how to expose Kubernetes
in the cluster using the kubectl expose
command. You will also learn how to view and apply labels to
a pod using the kubectl label command.**

This is a pre-configured Linux environment that can be used as a playground for learning Kubernetes (you can type commands). Clicking on the blocks of code will execute that command in the terminal.

Exercises on Deploying the module

Share what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

our app publicly

minutes

**You will learn how to expose Kubernetes
in the cluster using the kubectl expose
command. You will also learn how to view and apply labels to
a pod using the kubectl label command.**

This is a pre-configured Linux environment that can be used as a sandbox for learning Kubernetes (you can type commands). Clicking on the blocks of code will execute that command in the terminal.

Tutorials on exposing the module

Share what you thought of the module

se@gmail.com

(not shared)

Switch accounts

our app publicly

minutes

**You will learn how to expose Kubernetes
in the cluster using the kubectl expose
command. You will also learn how to view and apply labels to
a pod using the kubectl label command.**

This is a pre-configured Linux environment that can be used as a sandbox for learning Kubernetes (you can type commands). Clicking on the blocks of code will execute that command in the terminal.

Tutorials on exposing the module

Share what you thought of the module

se@gmail.com

(not shared)

Switch accounts

Module 3 (/docs/tutorials/kubernetes-basics/explore/explore-intro/)

tutorials/kubernetes-basics/)

Module 5 > (/docs/tutorials/kubernetes-basics/scale/scale-intro/)

2.5 - Scale Your App

2.5.1 - Running Multiple Instances of Your App

o using kubectl.

application


modules we created a [Deployment](#) ([workloads/controllers/deployment/](#)), and then exposed it with a [Service](#) ([/docs/concepts/services-networking/service/](#)). The Deployment created only one Pod for running our application. When traffic increased, we need to scale the application to keep up with user demand. This is accomplished by changing the number of replicas in a Deployment.

Summary:

- Scaling a Deployment

You can create from the start a Deployment with multiple instances using the --replicas parameter for the kubectl create deployment command

Overview



A Deployment will ensure new Pods are created and scheduled to available resources. Scaling will increase the number of Pods to the desired number. Kubernetes also supports [autoscaling](#) ([/docs/user-guides/pod-autoscaling/](#)) of Pods, but it is outside of the scope of this tutorial. Scaling to zero is also possible, and it will terminate all Pods of the Deployment.

Running multiple instances of an application will require a way to distribute the incoming traffic. Services have an integrated load-balancer that will direct traffic to all Pods of an exposed Deployment. Services will discover the running Pods using endpoints, to ensure the traffic is distributed to all available Pods.

By running multiple instances of an Application running, you would be able

Scaling is accomplished by changing the number of replicas in a Deployment.

ates without downtime. We'll cover that in the next module.
he online terminal and scale our application.

ive Tutorial ➤ [./docs/tutorials/kubernetes-basics/scale/scale-interactive/](/docs/tutorials/kubernetes-basics/scale/scale-interactive/)

2.5.2 - Interactive Tutorial - Scaling Your App

your app

minutes

Interactive scenario is to scale a deployment and to see the load balancing in action

a pre-configured Linux environment that can be used as a regular console (you can type commands). Clicking on the blocks of code followed by the ENTER sign will execute that command in the terminal.

Congratulations on completing the module

Share what you thought of the module

user@gmail.com (not shared) [Switch accounts](#)

your app

minutes

Interactive scenario is to scale a deployment with kubectl scale and to see the load balancing in action

a pre-configured Linux environment that can be used as a regular console (you can type commands). Clicking on the blocks of code followed by the ENTER sign will execute that command in the terminal.

Congratulations on completing the module

Share what you thought of the module

user@gmail.com (not shared) [Switch accounts](#)



/our app

inutes

teractive scenario is to scale a deployment
and to see the load balancing in action

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

/our app

inutes

teractive scenario is to scale a deployment
and to see the load balancing in action

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

/our app

inutes

teractive scenario is to scale a deployment
and to see the load balancing in action

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

/our app

inutes

teractive scenario is to scale a deployment
and to see the load balancing in action

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

your app

minutes

Interactive scenario is to scale a deployment and to see the load balancing in action

a pre-configured Linux environment that can be used as a regular console (you can type commands). Clicking on the blocks of code followed by the ENTER sign will execute that command in the terminal.

Congratulations on completing the module

Share what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

your app

minutes

Interactive scenario is to scale a deployment with kubectl scale and to see the load balancing in action

a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER sign will execute that command in the terminal).

Congratulations on completing the module

Share what you thought of the module

oise@gmail.com (not shared) [Switch accounts](#)



Module 4 (</docs/tutorials/kubernetes-basics/expose/expose-interactive/>)

[tutorials/kubernetes-basics/](/docs/tutorials/kubernetes-basics/))

Module 6 > (</docs/tutorials/kubernetes-basics/update/update-intro/>)

2.6 - Update Your App

2.6.1 - Performing a Rolling Update


olling update using kubectl.

an application

lications to be available all the time and developers are
oy new versions of them several times a day. In Kubernetes
rolling updates. **Rolling updates** allow Deployments' update
n zero downtime by incrementally updating Pods instances
he new Pods will be scheduled on Nodes with available

odule we scaled our application to run multiple instances.
nent for performing updates without affecting application
fault, the maximum number of Pods that can be unavailable
e and the maximum number of new Pods that can be created,
ons can be configured to either numbers or percentages (of
etes, updates are versioned and any Deployment update can
previous (stable) version.

updates overview



tion Scaling, if a Deployment is exposed publicly, the Service
the traffic only to available Pods during the update. An
n instance that is available to the users of the application.

llow the following actions:

application from one environment to another (via container
tes)
previous versions
Integration and Continuous Delivery of applications with zero

Summary:

- Updating an app

*Rolling updates allow
Deployments' update to take
place with zero downtime by
incrementally updating Pods
instances with new ones.*

*If a Deployment is exposed
publicly, the Service will load-
balance the traffic only to
available Pods during the update.*

nteractive tutorial, we'll update our application to a new
perform a rollback.

2.6.2 - Interactive Tutorial - Updating Your App

our app

inutes

enario is to update a deployed application with kubectl set image and ie rollout undo command.

a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER sign will execute that command in the

ratulations on completing the
file

hear what you thought of the module

oise@gmail.com (not shared) [Switch accounts](#)



our app

inutes

enario is to update a deployed application with kubectl set image and ie rollout undo command.

a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER sign will execute that command in the

ratulations on completing the
file

hear what you thought of the module

oise@gmail.com (not shared) [Switch accounts](#)



our app

minutes

Scenario is to update a deployed application with kubectl set image and the rollout undo command.

a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER sign will execute that command in the

Congratulations on completing the module

Share what you thought of the module

pradise@gmail.com (not shared) [Switch accounts](#)



our app

minutes

Scenario is to update a deployed application with kubectl set image and to rollback with the rollout undo command.

a pre-configured Linux environment that can be used as a regular console (you can type commands). Clicking on the blocks of code followed by the ENTER sign will execute that command in the terminal.

Congratulations on completing the module

Share what you thought of the module

pradise@gmail.com (not shared) [Switch accounts](#)

our app

minutes

enario is to update a deployed application
nage and to rollback with the rollout undo

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

our app

minutes

enario is to update a deployed application
nage and to rollback with the rollout undo

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R sign will execute that command in the terminal.

atulations on eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

our app

minutes

Scenario is to update a deployed application image and to rollback with the rollout undo

a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER sign will execute that command in the terminal).

Congratulations on completing the module

Hear what you thought of the module

se@gmail.com

(not shared)

Switch accounts

our app

minutes

Scenario is to update a deployed application with kubectl set image and the rollout undo command.

a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER sign will execute that command in the terminal).

Congratulations on completing the file

Hear what you thought of the module

oise@gmail.com

(not shared)

Switch accounts

Module 5 (/docs/tutorials/kubernetes-basics/scale/scale-interactive/)

Kubernetes Basics (/docs/tutorials/kubernetes-basics/)

3 - Configuration

3.1 - Example: Configuring a Java Microservice

3.1.1 - Externalizing config using MicroProfile, ConfigMaps and Secrets

In this tutorial you will learn how and why to externalize your microservice's configuration. Specifically, you will learn how to use Kubernetes ConfigMaps and Secrets to set environment variables and then consume them using MicroProfile Config.

Before you begin

Creating Kubernetes ConfigMaps & Secrets

There are several ways to set environment variables for a Docker container in Kubernetes, including: Dockerfile, kubernetes.yml, Kubernetes ConfigMaps, and Kubernetes Secrets. In the tutorial, you will learn how to use the latter two for setting your environment variables whose values will be injected into your microservices. One of the benefits for using ConfigMaps and Secrets is that they can be re-used across multiple containers, including being assigned to different environment variables for the different containers.

ConfigMaps are API Objects that store non-confidential key-value pairs. In the Interactive Tutorial you will learn how to use a ConfigMap to store the application's name. For more information regarding ConfigMaps, you can find the documentation [here \(/docs/tasks/configure-pod-container/configure-pod-configmap/\)](/docs/tasks/configure-pod-container/configure-pod-configmap/).

Although Secrets are also used to store key-value pairs, they differ from ConfigMaps in that they're intended for confidential/sensitive information and are stored using Base64 encoding. This makes secrets the appropriate choice for storing such things as credentials, keys, and tokens, the former of which you'll do in the Interactive Tutorial. For more information on Secrets, you can find the documentation [here \(/docs/concepts/configuration/secret/\)](/docs/concepts/configuration/secret/).

Externalizing Config from Code

Externalized application configuration is useful because configuration usually changes depending on your environment. In order to accomplish this, we'll use Java's Contexts and Dependency Injection (CDI) and MicroProfile Config. MicroProfile Config is a feature of MicroProfile, a set of open Java technologies for developing and deploying cloud-native microservices.

CDI provides a standard dependency injection capability enabling an application to be assembled from collaborating, loosely-coupled beans. MicroProfile Config provides apps and microservices a standard way to obtain config properties from various sources, including the application, runtime, and environment. Based on the source's defined priority, the properties are automatically combined into a single set of properties that the application can access via an API. Together, CDI & MicroProfile will be used in the Interactive Tutorial to retrieve the externally provided properties from the Kubernetes ConfigMaps and Secrets and get injected into your application code.

Many open source frameworks and runtimes implement and support MicroProfile Config. Throughout the interactive tutorial, you'll be using Open Liberty, a flexible open-source Java runtime for building and running cloud-native apps and microservices. However, any MicroProfile compatible runtime could be used instead.

Objectives

- Create a Kubernetes ConfigMap and Secret
- Inject microservice configuration using MicroProfile Config

Example: Externalizing config using MicroProfile, ConfigMaps and Secrets

[Start Interactive Tutorial \(/docs/tutorials/configuration/configure-java-microservice/configure-java-microservice-interactive/\)](#)

3.1.2 - Interactive Tutorial - Configuring a Java Microservice

We're preparing the environment for you.

netes Microservice

minutes

This interactive scenario is to deploy two Java microservices to Kubernetes and change their configuration using MicroProfile Config, Kubernetes ConfigMaps and Secrets.

This is a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER key will execute that command in the terminal).

Congratulations on completing the tutorial!

We're preparing the environment for you.

netes Microservice

minutes

This interactive scenario is to deploy two Java microservices to Kubernetes and change their configuration using MicroProfile Config, Kubernetes ConfigMaps and Secrets.

This is a pre-configured Linux environment that can be used as a regular console (you can type on the blocks of code followed by the ENTER key will execute that command in the terminal).

Congratulations on

We're preparing the environment for

netes Microservice

minutes

This interactive scenario is to deploy
microservices to Kubernetes and change
configuration using MicroProfile Config,
ConfigMaps and Secrets.

This is a pre-configured Linux environment that can be used as
a terminal (you can type commands). Clicking on the blocks of code
or the Run key will execute that command in the terminal.

atulations on

We're preparing the environment for

netes Microservice

minutes

This interactive scenario is to deploy
microservices to Kubernetes and change
configuration using MicroProfile Config,
ConfigMaps and Secrets.

This is a pre-configured Linux environment that can be used as
a terminal (you can type commands). Clicking on the blocks of code
or the Run key will execute that command in the terminal.

atulations on

We're preparing the environment for

Kubernetes Microservice

minutes

This interactive scenario is to deploy
microservices to Kubernetes and change
configuration using MicroProfile Config,
ConfigMaps and Secrets.

This is a pre-configured Linux environment that can be used as
a terminal (you can type commands). Clicking on the blocks of code
or the **Run** key will execute that command in the terminal.

Prerequisites on

Kubernetes Microservice

minutes

This interactive scenario is to deploy
microservices to Kubernetes and change
configuration using MicroProfile Config,
ConfigMaps and Secrets.

This is a pre-configured Linux environment that can be used as
a terminal (you can type commands). Clicking on the blocks of code
or the **Run** key will execute that command in the terminal.

Prerequisites on
Installing the module

Learn what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

is interactive scenario is to deploy
oservices to Kubernetes and change
ation using MicroProfile Config,
onfigMaps and Secrets.

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R key will execute that command in the terminal.

atulations on
eting the module

ear what you thought of the module

se@gmail.com (not shared) [Switch accounts](#)

e’re preparing the environment for

is interactive scenario is to deploy
oservices to Kubernetes and change
ation using MicroProfile Config,
onfigMaps and Secrets.

a pre-configured Linux environment that can be used as
can type commands). Clicking on the blocks of code
R key will execute that command in the terminal.

atulations on

3.2 - Configuring Redis using a ConfigMap

This page provides a real world example of how to configure Redis using a ConfigMap and builds upon the [Configure Containers Using a ConfigMap \(/docs/tasks/configure-pod-container/configure-pod-configmap/\)](/docs/tasks/configure-pod-container/configure-pod-configmap/) task.

Objectives

- Create a ConfigMap with Redis configuration values
- Create a Redis Pod that mounts and uses the created ConfigMap
- Verify that the configuration was correctly applied.

Before you begin

You need to have a Kubernetes cluster, and the `kubectl` command-line tool must be configured to communicate with your cluster. It is recommended to run this tutorial on a cluster with at least two nodes that are not acting as control plane hosts. If you do not already have a cluster, you can create one by using [minikube \(https://minikube.sigs.k8s.io/docs/tutorials/multi_node/\)](https://minikube.sigs.k8s.io/docs/tutorials/multi_node/) or you can use one of these Kubernetes playgrounds:

- [Katacoda \(https://www.katacoda.com/courses/kubernetes/playground\)](https://www.katacoda.com/courses/kubernetes/playground)
- [Play with Kubernetes \(http://labs.play-with-k8s.com/\)](http://labs.play-with-k8s.com/)

To check the version, enter `kubectl version`.

- The example shown on this page works with `kubectl` 1.14 and above.
- Understand [Configure Containers Using a ConfigMap \(/docs/tasks/configure-pod-container/configure-pod-configmap/\)](/docs/tasks/configure-pod-container/configure-pod-configmap/).

Real World Example: Configuring Redis using a ConfigMap

Follow the steps below to configure a Redis cache using data stored in a ConfigMap.

First create a ConfigMap with an empty configuration block:

```
cat <<EOF >./example-redis-config.yaml
apiVersion: v1
kind: ConfigMap
metadata:
  name: example-redis-config
data:
  redis-config: ""
EOF
```

Apply the ConfigMap created above, along with a Redis pod manifest:

```
kubectl apply -f example-redis-config.yaml
kubectl apply -f https://raw.githubusercontent.com/kubernetes/website/main/content/en/example
```

Examine the contents of the Redis pod manifest and note the following:

- A volume named `config` is created by `spec.volumes[1]`
- The key and path under `spec.volumes[1].items[0]` exposes the `redis-config` key from the `example-redis-config` ConfigMap as a file named `redis.conf` on the `config` volume.
- The `config` volume is then mounted at `/redis-master` by `spec.containers[0].volumeMounts[1]`.

This has the net effect of exposing the data in `data.redis-config` from the `example-redis-config` ConfigMap above as `/redis-master/redis.conf` inside the Pod.

[pods/config/redis-pod.yaml](#)
<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/config/redis-pod.yaml>

```

apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
  - name: redis
    image: redis:5.0.4
    command:
    - redis-server
    - "/redis-master/redis.conf"
    env:
    - name: MASTER
      value: "true"
    ports:
    - containerPort: 6379
    resources:
      limits:
        cpu: "0.1"
    volumeMounts:
    - mountPath: /redis-master-data
      name: data
    - mountPath: /redis-master
      name: config
  volumes:
  - name: data
    emptyDir: {}
  - name: config
    configMap:
      name: example-redis-config
      items:
      - key: redis-config
        path: redis.conf

```

Examine the created objects:

```
kubectl get pod/redis configmap/example-redis-config
```

You should see the following output:

NAME	READY	STATUS	RESTARTS	AGE
pod/redis	1/1	Running	0	8s

NAME	DATA	AGE
configmap/example-redis-config	1	14s

Recall that we left `redis-config` key in the `example-redis-config` ConfigMap blank:

```
kubectl describe configmap/example-redis-config
```

You should see an empty `redis-config` key:

```

Name:         example-redis-config
Namespace:    default
Labels:       <none>
Annotations:  <none>

Data

```

```
====  
redis-config:
```

Use `kubectl exec` to enter the pod and run the `redis-cli` tool to check the current configuration:

```
kubectl exec -it redis -- redis-cli
```

Check `maxmemory` :

```
127.0.0.1:6379> CONFIG GET maxmemory
```

It should show the default value of 0:

```
1) "maxmemory"  
2) "0"
```

Similarly, check `maxmemory-policy` :

```
127.0.0.1:6379> CONFIG GET maxmemory-policy
```

Which should also yield its default value of `noeviction` :

```
1) "maxmemory-policy"  
2) "noeviction"
```

Now let's add some configuration values to the `example-redis-config` ConfigMap:

[pods/config/example-redis-config.yaml](https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/config/example-redis-config.yaml)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/config/example-redis-config.yaml>) 

```
apiVersion: v1  
kind: ConfigMap  
metadata:  
  name: example-redis-config  
data:  
  redis-config: |  
    maxmemory 2mb  
    maxmemory-policy allkeys-lru
```

Apply the updated ConfigMap:

```
kubectl apply -f example-redis-config.yaml
```

Confirm that the ConfigMap was updated:

```
kubectl describe configmap/example-redis-config
```

You should see the configuration values we just added:

```
Name:      example-redis-config
Namespace:  default
Labels:     <none>
Annotations: <none>
```

```
Data
====
redis-config:
-----
maxmemory 2mb
maxmemory-policy allkeys-lru
```

Check the Redis Pod again using `redis-cli` via `kubectl exec` to see if the configuration was applied:

```
kubectl exec -it redis -- redis-cli
```

Check `maxmemory` :

```
127.0.0.1:6379> CONFIG GET maxmemory
```

It remains at the default value of 0:

```
1) "maxmemory"
2) "0"
```

Similarly, `maxmemory-policy` remains at the `noeviction` default setting:

```
127.0.0.1:6379> CONFIG GET maxmemory-policy
```

Returns:

```
1) "maxmemory-policy"
2) "noeviction"
```

The configuration values have not changed because the Pod needs to be restarted to grab updated values from associated ConfigMaps. Let's delete and recreate the Pod:

```
kubectl delete pod redis
kubectl apply -f https://raw.githubusercontent.com/kubernetes/website/main/content/en/example
```

Now re-check the configuration values one last time:

```
kubectl exec -it redis -- redis-cli
```

Check `maxmemory` :

```
127.0.0.1:6379> CONFIG GET maxmemory
```

It should now return the updated value of 2097152:

```
1) "maxmemory"  
2) "2097152"
```

Similarly, `maxmemory-policy` has also been updated:

```
127.0.0.1:6379> CONFIG GET maxmemory-policy
```

It now reflects the desired value of `allkeys-lru` :

```
1) "maxmemory-policy"  
2) "allkeys-lru"
```

Clean up your work by deleting the created resources:

```
kubectl delete pod/redis configmap/example-redis-config
```

What's next

- Learn more about [ConfigMaps \(/docs/tasks/configure-pod-container/configure-pod-configmap/\)](/docs/tasks/configure-pod-container/configure-pod-configmap/).

4 - Stateless Applications

4.1 - Exposing an External IP Address to Access an Application in a Cluster

This page shows how to create a Kubernetes Service object that exposes an external IP address.

Before you begin

- Install [kubectL \(/docs/tasks/tools/\)](/docs/tasks/tools/).
- Use a cloud provider like Google Kubernetes Engine or Amazon Web Services to create a Kubernetes cluster. This tutorial creates an [external load balancer \(/docs/tasks/access-application-cluster/create-external-load-balancer/\)](/docs/tasks/access-application-cluster/create-external-load-balancer/), which requires a cloud provider.
- Configure `kubectl` to communicate with your Kubernetes API server. For instructions, see the documentation for your cloud provider.

Objectives

- Run five instances of a Hello World application.
- Create a Service object that exposes an external IP address.
- Use the Service object to access the running application.

Creating a service for an application running in five pods

1. Run a Hello World application in your cluster:

service/load-balancer-example.yaml

(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/service/load-balancer-example.yaml>)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app.kubernetes.io/name: load-balancer-example
  name: hello-world
spec:
  replicas: 5
  selector:
    matchLabels:
      app.kubernetes.io/name: load-balancer-example
  template:
    metadata:
      labels:
        app.kubernetes.io/name: load-balancer-example
    spec:
      containers:
        - image: gcr.io/google-samples/node-hello:1.0
          name: hello-world
          ports:
            - containerPort: 8080
```

```
kubectl apply -f https://k8s.io/examples/service/load-balancer-example.yaml
```

The preceding command creates a [Deployment \(/docs/concepts/workloads/controllers/deployment/\)](/docs/concepts/workloads/controllers/deployment/) and an associated [ReplicaSet \(/docs/concepts/workloads/controllers/replicaset/\)](/docs/concepts/workloads/controllers/replicaset/). The ReplicaSet has five

Pods (/docs/concepts/workloads/pods/) each of which runs the Hello World application.

2. Display information about the Deployment:

```
kubectl get deployments hello-world
kubectl describe deployments hello-world
```

3. Display information about your ReplicaSet objects:

```
kubectl get replicaset
kubectl describe replicaset
```

4. Create a Service object that exposes the deployment:

```
kubectl expose deployment hello-world --type=LoadBalancer --name=my-service
```

5. Display information about the Service:

```
kubectl get services my-service
```

The output is similar to:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
my-service	LoadBalancer	10.3.245.137	104.198.205.71	8080/TCP	54s

Note: The `type=LoadBalancer` service is backed by external cloud providers, which is not covered in this example, please refer to [this page \(/docs/concepts/services-networking/service/#loadbalancer\)](#) for the details.

Note: If the external IP address is shown as `<pending>`, wait for a minute and enter the same command again.

6. Display detailed information about the Service:

```
kubectl describe services my-service
```

The output is similar to:

```
Name: my-service
Namespace: default
Labels: app.kubernetes.io/name=load-balancer-example
Annotations: <none>
Selector: app.kubernetes.io/name=load-balancer-example
Type: LoadBalancer
IP: 10.3.245.137
LoadBalancer Ingress: 104.198.205.71
Port: <unset> 8080/TCP
NodePort: <unset> 32377/TCP
Endpoints: 10.0.0.6:8080,10.0.1.6:8080,10.0.1.7:8080 + 2 more...
Session Affinity: None
Events: <none>
```

Make a note of the external IP address (`LoadBalancer Ingress`) exposed by your service. In this example, the external IP address is 104.198.205.71. Also note the value of `Port` and `NodePort` . In this example, the `Port` is 8080 and the `NodePort` is 32377.

7. In the preceding output, you can see that the service has several endpoints: 10.0.0.6:8080,10.0.1.6:8080,10.0.1.7:8080 + 2 more. These are internal addresses of the pods that are running the Hello World application. To verify these are pod addresses, enter this command:

```
kubectl get pods --output=wide
```

The output is similar to:

NAME	...	IP	NODE
hello-world-2895499144-1jaz9	...	10.0.1.6	gke-cluster-1-default-pool-e0b8d269-1afc
hello-world-2895499144-2e5uh	...	10.0.1.8	gke-cluster-1-default-pool-e0b8d269-1afc
hello-world-2895499144-9m4h1	...	10.0.0.6	gke-cluster-1-default-pool-e0b8d269-5v7a
hello-world-2895499144-o4z13	...	10.0.1.7	gke-cluster-1-default-pool-e0b8d269-1afc
hello-world-2895499144-segjf	...	10.0.2.5	gke-cluster-1-default-pool-e0b8d269-cpuc

8. Use the external IP address (`LoadBalancer Ingress`) to access the Hello World application:

```
curl http://<external-ip>:<port>
```

where `<external-ip>` is the external IP address (`LoadBalancer Ingress`) of your Service, and `<port>` is the value of `Port` in your Service description. If you are using minikube, typing `minikube service my-service` will automatically open the Hello World application in a browser.

The response to a successful request is a hello message:

```
Hello Kubernetes!
```

Cleaning up

To delete the Service, enter this command:

```
kubectl delete services my-service
```

To delete the Deployment, the ReplicaSet, and the Pods that are running the Hello World application, enter this command:

```
kubectl delete deployment hello-world
```

What's next

Learn more about [connecting applications with services \(/docs/concepts/services-networking/connect-applications-service/\)](/docs/concepts/services-networking/connect-applications-service/).

4.2 - Example: Deploying PHP Guestbook application with Redis

This tutorial shows you how to build and deploy a simple (*not production ready*), multi-tier web application using Kubernetes and [Docker](https://www.docker.com/) (<https://www.docker.com/>). This example consists of the following components:

- A single-instance [Redis](https://www.redis.io/) (<https://www.redis.io/>) to store guestbook entries
- Multiple web frontend instances

Objectives

- Start up a Redis leader.
- Start up two Redis followers.
- Start up the guestbook frontend.
- Expose and view the Frontend Service.
- Clean up.

Before you begin

You need to have a Kubernetes cluster, and the kubectl command-line tool must be configured to communicate with your cluster. It is recommended to run this tutorial on a cluster with at least two nodes that are not acting as control plane hosts. If you do not already have a cluster, you can create one by using [minikube](https://minikube.sigs.k8s.io/docs/tutorials/multi_node/) (https://minikube.sigs.k8s.io/docs/tutorials/multi_node/) or you can use one of these Kubernetes playgrounds:

- [Katacoda](https://www.katacoda.com/courses/kubernetes/playground) (<https://www.katacoda.com/courses/kubernetes/playground>)
- [Play with Kubernetes](http://labs.play-with-k8s.com/) (<http://labs.play-with-k8s.com/>)


Your Kubernetes server must be at or later than version v1.14. To check the version, enter `kubectl version`.

Start up the Redis Database

The guestbook application uses Redis to store its data.

Creating the Redis Deployment

The manifest file, included below, specifies a Deployment controller that runs a single replica Redis Pod.

[application/guestbook/redis-leader-deployment.yaml](https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/guestbook/redis-leader-deployment.yaml)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/guestbook/redis-leader-deployment.yaml>) 

```
# SOURCE: https://cloud.google.com/kubernetes-engine/docs/tutorials/guestbook
apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis-leader
  labels:
    app: redis
    role: leader
    tier: backend
spec:
  replicas: 1
  selector:
    matchLabels:
      app: redis
  template:
    metadata:
      labels:
        app: redis
```

```
    role: leader
    tier: backend
spec:
  containers:
  - name: leader
    image: "docker.io/redis:6.0.5"
    resources:
      requests:
        cpu: 100m
        memory: 100Mi
    ports:
    - containerPort: 6379
```

1. Launch a terminal window in the directory you downloaded the manifest files.
2. Apply the Redis Deployment from the `redis-leader-deployment.yaml` file:

```
kubectl apply -f https://k8s.io/examples/application/guestbook/redis-leader-deployment.y
```

3. Query the list of Pods to verify that the Redis Pod is running:

```
kubectl get pods
```

The response should be similar to this:

NAME	READY	STATUS	RESTARTS	AGE
redis-leader-fb76b4755-xjr2n	1/1	Running	0	13s

4. Run the following command to view the logs from the Redis leader Pod:

```
kubectl logs -f deployment/redis-leader
```

Creating the Redis leader Service

The guestbook application needs to communicate to the Redis to write its data. You need to apply a [Service \(/docs/concepts/services-networking/service/\)](/docs/concepts/services-networking/service/) to proxy the traffic to the Redis Pod. A Service defines a policy to access the Pods.

[application/guestbook/redis-leader-service.yaml](https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/guestbook/redis-leader-service.yaml)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/guestbook/redis-leader-service.yaml>) 

```
# SOURCE: https://cloud.google.com/kubernetes-engine/docs/tutorials/guestbook
apiVersion: v1
kind: Service
metadata:
  name: redis-leader
  labels:
    app: redis
    role: leader
    tier: backend
spec:
  ports:
  - port: 6379
    targetPort: 6379
  selector:
    app: redis
    role: leader
    tier: backend
```

1. Apply the Redis Service from the following `redis-leader-service.yaml` file:

```
kubectl apply -f https://k8s.io/examples/application/guestbook/redis-leader-service.yaml
```

2. Query the list of Services to verify that the Redis Service is running:

```
kubectl get service
```

The response should be similar to this:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.0.0.1	<none>	443/TCP	1m
redis-leader	ClusterIP	10.103.78.24	<none>	6379/TCP	16s

Note: This manifest file creates a Service named `redis-leader` with a set of labels that match the labels previously defined, so the Service routes network traffic to the Redis Pod.

Set up Redis followers

Although the Redis leader is a single Pod, you can make it highly available and meet traffic demands by adding a few Redis followers, or replicas.

[application/guestbook/redis-follower-deployment.yaml](#)

(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/guestbook/redis-follower-deployment.yaml>)

```
# SOURCE: https://cloud.google.com/kubernetes-engine/docs/tutorials/guestbook
apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis-follower
  labels:
    app: redis
    role: follower
    tier: backend
spec:
  replicas: 2
  selector:
    matchLabels:
      app: redis
  template:
    metadata:
      labels:
        app: redis
        role: follower
        tier: backend
    spec:
      containers:
        - name: follower
          image: gcr.io/google_samples/gb-redis-follower:v2
          resources:
            requests:
              cpu: 100m
              memory: 100Mi
          ports:
            - containerPort: 6379
```

1. Apply the Redis Deployment from the following `redis-follower-deployment.yaml` file:

```
kubectl apply -f https://k8s.io/examples/application/guestbook/redis-follower-deployment
```

2. Verify that the two Redis follower replicas are running by querying the list of Pods:

```
kubectl get pods
```

The response should be similar to this:

NAME	READY	STATUS	RESTARTS	AGE
redis-follower-dddfbdcc9-82sfr	1/1	Running	0	37s
redis-follower-dddfbdcc9-qrt5k	1/1	Running	0	38s
redis-leader-fb76b4755-xjr2n	1/1	Running	0	11m

Creating the Redis follower service

The guestbook application needs to communicate with the Redis followers to read data. To make the Redis followers discoverable, you must set up another [Service \(/docs/concepts/services-networking/service/\)](#).

[application/guestbook/redis-follower-service.yaml](#)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/guestbook/redis-follower-service.yaml>)

```
# SOURCE: https://cloud.google.com/kubernetes-engine/docs/tutorials/guestbook
apiVersion: v1
kind: Service
metadata:
  name: redis-follower
  labels:
    app: redis
    role: follower
    tier: backend
spec:
  ports:
    # the port that this service should serve on
    - port: 6379
  selector:
    app: redis
    role: follower
    tier: backend
```

1. Apply the Redis Service from the following `redis-follower-service.yaml` file:

```
kubectl apply -f https://k8s.io/examples/application/guestbook/redis-follower-service.ya
```

2. Query the list of Services to verify that the Redis Service is running:

```
kubectl get service
```

The response should be similar to this:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3d19h
redis-follower	ClusterIP	10.110.162.42	<none>	6379/TCP	9s
redis-leader	ClusterIP	10.103.78.24	<none>	6379/TCP	6m10s

Note: This manifest file creates a Service named `redis-follower` with a set of labels that match the labels previously defined, so the Service routes network traffic to the Redis Pod.

Set up and Expose the Guestbook Frontend

Now that you have the Redis storage of your guestbook up and running, start the guestbook web servers. Like the Redis followers, the frontend is deployed using a Kubernetes Deployment.

The guestbook app uses a PHP frontend. It is configured to communicate with either the Redis follower or leader Services, depending on whether the request is a read or a write. The frontend exposes a JSON interface, and serves a jQuery-Ajax-based UX.

Creating the Guestbook Frontend Deployment

application/guestbook/frontend-deployment.yaml

(https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/guestbook/frontend-deployment.yaml)

```
# SOURCE: https://cloud.google.com/kubernetes-engine/docs/tutorials/guestbook
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
spec:
  replicas: 3
  selector:
    matchLabels:
      app: guestbook
      tier: frontend
  template:
    metadata:
      labels:
        app: guestbook
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v5
          env:
            - name: GET_HOSTS_FROM
              value: "dns"
          resources:
            requests:
              cpu: 100m
              memory: 100Mi
          ports:
            - containerPort: 80
```

1. Apply the frontend Deployment from the `frontend-deployment.yaml` file:

```
kubectl apply -f https://k8s.io/examples/application/guestbook/frontend-deployment.yaml
```

2. Query the list of Pods to verify that the three frontend replicas are running:

```
kubectl get pods -l app=guestbook -l tier=frontend
```

The response should be similar to this:

NAME	READY	STATUS	RESTARTS	AGE
frontend-85595f5bf9-5tqhb	1/1	Running	0	47s
frontend-85595f5bf9-qbzwm	1/1	Running	0	47s
frontend-85595f5bf9-zchwc	1/1	Running	0	47s

Creating the Frontend Service

The Redis Services you applied is only accessible within the Kubernetes cluster because the default type for a Service is [ClusterIP \(/docs/concepts/services-networking/service/#publishing-services-service-types\)](#). ClusterIP provides a single IP address for the set of Pods the Service is pointing to. This IP address is accessible only within the cluster.

If you want guests to be able to access your guestbook, you must configure the frontend Service to be externally visible, so a client can request the Service from outside the Kubernetes cluster. However a Kubernetes user you can use `kubectl port-forward` to access the service even though it uses a `ClusterIP`.

Note: Some cloud providers, like Google Compute Engine or Google Kubernetes Engine, support external load balancers. If your cloud provider supports load balancers and you want to use it, uncomment `type: LoadBalancer`.

`application/guestbook/frontend-service.yaml`
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/guestbook/frontend-service.yaml>)

```
# SOURCE: https://cloud.google.com/kubernetes-engine/docs/tutorials/guestbook
apiVersion: v1
kind: Service
metadata:
  name: frontend
  labels:
    app: guestbook
    tier: frontend
spec:
  # if your cluster supports it, uncomment the following to automatically create
  # an external load-balanced IP for the frontend service.
  # type: LoadBalancer
  #type: LoadBalancer
  ports:
    # the port that this service should serve on
    - port: 80
  selector:
    app: guestbook
    tier: frontend
```

1. Apply the frontend Service from the `frontend-service.yaml` file:

```
kubectl apply -f https://k8s.io/examples/application/guestbook/frontend-service.yaml
```

2. Query the list of Services to verify that the frontend Service is running:

```
kubectl get services
```

The response should be similar to this:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
frontend	ClusterIP	10.97.28.230	<none>	80/TCP	19s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3d19h
redis-follower	ClusterIP	10.110.162.42	<none>	6379/TCP	5m48s
redis-leader	ClusterIP	10.103.78.24	<none>	6379/TCP	11m

Viewing the Frontend Service via `kubectl port-forward`

1. Run the following command to forward port `8080` on your local machine to port `80` on the service.

```
kubectl port-forward svc/frontend 8080:80
```

The response should be similar to this:


```
Forwarding from 127.0.0.1:8080 -> 80
Forwarding from [::1]:8080 -> 80
```

2. load the page <http://localhost:8080> (<http://localhost:8080>) in your browser to view your guestbook.

Viewing the Frontend Service via LoadBalancer

If you deployed the `frontend-service.yaml` manifest with type: `LoadBalancer` you need to find the IP address to view your Guestbook.

1. Run the following command to get the IP address for the frontend Service.

```
kubectl get service frontend
```

The response should be similar to this:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
frontend	LoadBalancer	10.51.242.136	109.197.92.229	80:32372/TCP	1m

2. Copy the external IP address, and load the page in your browser to view your guestbook.

Note: Try adding some guestbook entries by typing in a message, and clicking Submit. The message you typed appears in the frontend. This message indicates that data is successfully added to Redis through the Services you created earlier.

Scale the Web Frontend

You can scale up or down as needed because your servers are defined as a Service that uses a Deployment controller.

1. Run the following command to scale up the number of frontend Pods:

```
kubectl scale deployment frontend --replicas=5
```

2. Query the list of Pods to verify the number of frontend Pods running:

```
kubectl get pods
```

The response should look similar to this:

NAME	READY	STATUS	RESTARTS	AGE
frontend-85595f5bf9-5df5m	1/1	Running	0	83s
frontend-85595f5bf9-7zmg5	1/1	Running	0	83s
frontend-85595f5bf9-cpskg	1/1	Running	0	15m
frontend-85595f5bf9-l2l54	1/1	Running	0	14m
frontend-85595f5bf9-l9c8z	1/1	Running	0	14m
redis-follower-dddfbdcc9-82sfr	1/1	Running	0	97m
redis-follower-dddfbdcc9-qrt5k	1/1	Running	0	97m
redis-leader-fb76b4755-xjr2n	1/1	Running	0	108m

3. Run the following command to scale down the number of frontend Pods:

```
kubectl scale deployment frontend --replicas=2
```

4. Query the list of Pods to verify the number of frontend Pods running:

```
kubectl get pods
```

The response should look similar to this:

NAME	READY	STATUS	RESTARTS	AGE
frontend-85595f5bf9-cpskg	1/1	Running	0	16m
frontend-85595f5bf9-l9c8z	1/1	Running	0	15m
redis-follower-dddfbdcc9-82sfr	1/1	Running	0	98m
redis-follower-dddfbdcc9-qrt5k	1/1	Running	0	98m
redis-leader-fb76b4755-xjr2n	1/1	Running	0	109m

Cleaning up

Deleting the Deployments and Services also deletes any running Pods. Use labels to delete multiple resources with one command.

1. Run the following commands to delete all Pods, Deployments, and Services.

```
kubectl delete deployment -l app=redis
kubectl delete service -l app=redis
kubectl delete deployment frontend
kubectl delete service frontend
```

The response should look similar to this:

```
deployment.apps "redis-follower" deleted
deployment.apps "redis-leader" deleted
deployment.apps "frontend" deleted
service "frontend" deleted
```

2. Query the list of Pods to verify that no Pods are running:

```
kubectl get pods
```

The response should look similar to this:

```
No resources found in default namespace.
```

What's next

- Complete the [Kubernetes Basics \(/docs/tutorials/kubernetes-basics/\)](/docs/tutorials/kubernetes-basics/) Interactive Tutorials
- Use Kubernetes to create a blog using [Persistent Volumes for MySQL and Wordpress \(/docs/tutorials/stateful-application/mysql-wordpress-persistent-volume/#visit-your-new-wordpress-blog\)](/docs/tutorials/stateful-application/mysql-wordpress-persistent-volume/#visit-your-new-wordpress-blog).
- Read more about [connecting applications \(/docs/concepts/services-networking/connect-applications-service/\)](/docs/concepts/services-networking/connect-applications-service/)
- Read more about [Managing Resources \(/docs/concepts/cluster-administration/manage-deployment/#using-labels-effectively\)](/docs/concepts/cluster-administration/manage-deployment/#using-labels-effectively)

5 - Stateful Applications

5.1 - StatefulSet Basics

This tutorial provides an introduction to managing applications with StatefulSets (</docs/concepts/workloads/controllers/statefulset/>). It demonstrates how to create, delete, scale, and update the Pods of StatefulSets.

Before you begin

Before you begin this tutorial, you should familiarize yourself with the following Kubernetes concepts:

- [Pods](/docs/concepts/workloads/pods/) (</docs/concepts/workloads/pods/>).
- [Cluster DNS](/docs/concepts/services-networking/dns-pod-service/) (</docs/concepts/services-networking/dns-pod-service/>).
- [Headless Services](/docs/concepts/services-networking/service/#headless-services) (</docs/concepts/services-networking/service/#headless-services>).
- [PersistentVolumes](/docs/concepts/storage/persistent-volumes/) (</docs/concepts/storage/persistent-volumes/>).
- [PersistentVolume Provisioning](https://github.com/kubernetes/examples/tree/master/staging/persistent-volume-provisioning/) (<https://github.com/kubernetes/examples/tree/master/staging/persistent-volume-provisioning/>).
- [StatefulSets](/docs/concepts/workloads/controllers/statefulset/) (</docs/concepts/workloads/controllers/statefulset/>).
- The `kubectl` (</docs/reference/kubectl/kubectl/>) command line tool

Note: This tutorial assumes that your cluster is configured to dynamically provision PersistentVolumes. If your cluster is not configured to do so, you will have to manually provision two 1 GiB volumes prior to starting this tutorial.

Objectives

StatefulSets are intended to be used with stateful applications and distributed systems. However, the administration of stateful applications and distributed systems on Kubernetes is a broad, complex topic. In order to demonstrate the basic features of a StatefulSet, and not to conflate the former topic with the latter, you will deploy a simple web application using a StatefulSet.

After this tutorial, you will be familiar with the following.

- How to create a StatefulSet
- How a StatefulSet manages its Pods
- How to delete a StatefulSet
- How to scale a StatefulSet
- How to update a StatefulSet's Pods

Creating a StatefulSet

Begin by creating a StatefulSet using the example below. It is similar to the example presented in the [StatefulSets](/docs/concepts/workloads/controllers/statefulset/) (</docs/concepts/workloads/controllers/statefulset/>) concept. It creates a [headless Service](/docs/concepts/services-networking/service/#headless-services) (</docs/concepts/services-networking/service/#headless-services>), `nginx`, to publish the IP addresses of Pods in the StatefulSet, `web`.

[application/web/web.yaml](https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/web/web.yaml)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/web/web.yaml>)

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
  labels:
    app: nginx
spec:
```

```

ports:
- port: 80
  name: web
clusterIP: None
selector:
  app: nginx
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  serviceName: "nginx"
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: k8s.gcr.io/nginx-slim:0.8
        ports:
        - containerPort: 80
          name: web
        volumeMounts:
        - name: www
          mountPath: /usr/share/nginx/html
  volumeClaimTemplates:
  - metadata:
      name: www
    spec:
      accessModes: [ "ReadWriteOnce" ]
      resources:
        requests:
          storage: 1Gi

```

Download the example above, and save it to a file named `web.yaml`

You will need to use two terminal windows. In the first terminal, use `kubectl get` (</docs/reference/generated/kubectl/kubectl-commands/#get>) to watch the creation of the StatefulSet's Pods.

```
kubectl get pods -w -l app=nginx
```

In the second terminal, use `kubectl apply` (</docs/reference/generated/kubectl/kubectl-commands/#apply>) to create the headless Service and StatefulSet defined in `web.yaml`.

```
kubectl apply -f web.yaml
```

```
service/nginx created
statefulset.apps/web created
```

The command above creates two Pods, each running an [NGINX](https://www.nginx.com) (<https://www.nginx.com>) webserver. Get the `nginx` Service...

```
kubectl get service nginx
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
nginx	ClusterIP	None	<none>	80/TCP	12s

...then get the `web` StatefulSet, to verify that both were created successfully:

```
kubectl get statefulset web
```

NAME	DESIRED	CURRENT	AGE
web	2	1	20s

Ordered Pod Creation

For a StatefulSet with n replicas, when Pods are being deployed, they are created sequentially, ordered from $\{0..n-1\}$. Examine the output of the `kubectl get` command in the first terminal. Eventually, the output will look like the example below.

```
kubectl get pods -w -l app=nginx
```

NAME	READY	STATUS	RESTARTS	AGE
web-0	0/1	Pending	0	0s
web-0	0/1	Pending	0	0s
web-0	0/1	ContainerCreating	0	0s
web-0	1/1	Running	0	19s
web-1	0/1	Pending	0	0s
web-1	0/1	Pending	0	0s
web-1	0/1	ContainerCreating	0	0s
web-1	1/1	Running	0	18s

Notice that the `web-1` Pod is not launched until the `web-0` Pod is *Running* (see [Pod Phase \(/docs/concepts/workloads/pods/pod-lifecycle/#pod-phase\)](/docs/concepts/workloads/pods/pod-lifecycle/#pod-phase)) and *Ready* (see `type` in [Pod Conditions \(/docs/concepts/workloads/pods/pod-lifecycle/#pod-conditions\)](/docs/concepts/workloads/pods/pod-lifecycle/#pod-conditions)).

Pods in a StatefulSet

Pods in a StatefulSet have a unique ordinal index and a stable network identity.

Examining the Pod's Ordinal Index

Get the StatefulSet's Pods:

```
kubectl get pods -l app=nginx
```

NAME	READY	STATUS	RESTARTS	AGE
web-0	1/1	Running	0	1m
web-1	1/1	Running	0	1m

As mentioned in the [StatefulSets \(/docs/concepts/workloads/controllers/statefulset/\)](/docs/concepts/workloads/controllers/statefulset/) concept, the Pods in a StatefulSet have a sticky, unique identity. This identity is based on a unique ordinal index that is assigned to each Pod by the StatefulSet controller (</docs/concepts/architecture/controller/>). The Pods' names take the form `<statefulset name>-<ordinal index>`. Since the `web` StatefulSet has two replicas, it creates two Pods, `web-0` and `web-1`.

Using Stable Network Identities

Each Pod has a stable hostname based on its ordinal index. Use `kubectl exec` (</docs/reference/generated/kubectl/kubectl-commands/#exec>) to execute the `hostname` command in each Pod:

```
for i in 0 1; do kubectl exec "web-$i" -- sh -c 'hostname'; done
```

```
web-0
web-1
```

Use `kubectl run` [\(/docs/reference/generated/kubectl/kubectl-commands/#run\)](#) to execute a container that provides the `nslookup` command from the `dnsutils` package. Using `nslookup` on the Pods' hostnames, you can examine their in-cluster DNS addresses:

```
kubectl run -i --tty --image busybox:1.28 dns-test --restart=Never --rm
```

which starts a new shell. In that new shell, run:

```
# Run this in the dns-test container shell
nslookup web-0.nginx
```

The output is similar to:

```
Server:      10.0.0.10
Address 1: 10.0.0.10 kube-dns.kube-system.svc.cluster.local

Name:      web-0.nginx
Address 1: 10.244.1.6

nslookup web-1.nginx
Server:      10.0.0.10
Address 1: 10.0.0.10 kube-dns.kube-system.svc.cluster.local

Name:      web-1.nginx
Address 1: 10.244.2.6
```

(and now exit the container shell: `exit`)

The CNAME of the headless service points to SRV records (one for each Pod that is Running and Ready). The SRV records point to A record entries that contain the Pods' IP addresses.

In one terminal, watch the StatefulSet's Pods:

```
kubectl get pod -w -l app=nginx
```

In a second terminal, use `kubectl delete` [\(/docs/reference/generated/kubectl/kubectl-commands/#delete\)](#) to delete all the Pods in the StatefulSet:

```
kubectl delete pod -l app=nginx
```

```
pod "web-0" deleted
pod "web-1" deleted
```

Wait for the StatefulSet to restart them, and for both Pods to transition to Running and Ready:

```
kubectl get pod -w -l app=nginx
```

NAME	READY	STATUS	RESTARTS	AGE
web-0	0/1	ContainerCreating	0	0s
NAME	READY	STATUS	RESTARTS	AGE
web-0	1/1	Running	0	2s
web-1	0/1	Pending	0	0s
web-1	0/1	Pending	0	0s
web-1	0/1	ContainerCreating	0	0s
web-1	1/1	Running	0	34s

Use `kubectl exec` and `kubectl run` to view the Pods' hostnames and in-cluster DNS entries. First, view the Pods' hostnames:

```
for i in 0 1; do kubectl exec web-$i -- sh -c 'hostname'; done
```

```
web-0
web-1
```

then, run:

```
kubectl run -i --tty --image busybox:1.28 dns-test --restart=Never --rm /bin/sh
```

which starts a new shell.
In that new shell, run:

```
# Run this in the dns-test container shell
nslookup web-0.nginx
```

The output is similar to:

```
Server:      10.0.0.10
Address 1: 10.0.0.10 kube-dns.kube-system.svc.cluster.local

Name:        web-0.nginx
Address 1: 10.244.1.7

nslookup web-1.nginx
Server:      10.0.0.10
Address 1: 10.0.0.10 kube-dns.kube-system.svc.cluster.local

Name:        web-1.nginx
Address 1: 10.244.2.8
```

(and now exit the container shell: `exit`)

The Pods' ordinals, hostnames, SRV records, and A record names have not changed, but the IP addresses associated with the Pods may have changed. In the cluster used for this tutorial, they have. This is why it is important not to configure other applications to connect to Pods in a `StatefulSet` by IP address.

If you need to find and connect to the active members of a `StatefulSet`, you should query the `CNAME` of the headless Service (`nginx.default.svc.cluster.local`). The SRV records associated with the `CNAME` will contain only the Pods in the `StatefulSet` that are `Running` and `Ready`.

If your application already implements connection logic that tests for liveness and readiness, you can use the SRV records of the Pods (`web-0.nginx.default.svc.cluster.local` , `web-1.nginx.default.svc.cluster.local`), as they are stable, and your application will be able to discover the Pods' addresses when they transition to `Running` and `Ready`.

Writing to Stable Storage

Get the `PersistentVolumeClaims` for `web-0` and `web-1` :

```
kubectl get pvc -l app=nginx
```

The output is similar to:

NAME	STATUS	VOLUME	CAPACITY	ACCESSMODES	AGE
www-web-0	Bound	pvc-15c268c7-b507-11e6-932f-42010a800002	1Gi	RW0	48s
www-web-1	Bound	pvc-15c79307-b507-11e6-932f-42010a800002	1Gi	RW0	48s

The `StatefulSet` controller created two `PersistentVolumeClaims` (</docs/concepts/storage/persistent-volumes/>) that are bound to two `PersistentVolumes` (</docs/concepts/storage/persistent-volumes/>).

As the cluster used in this tutorial is configured to dynamically provision PersistentVolumes, the PersistentVolumes were created and bound automatically.

The NGINX webserver, by default, serves an index file from `/usr/share/nginx/html/index.html` . The `volumeMounts` field in the StatefulSet's `spec` ensures that the `/usr/share/nginx/html` directory is backed by a PersistentVolume.

Write the Pods' hostnames to their `index.html` files and verify that the NGINX webserver serves the hostnames:

```
for i in 0 1; do kubectl exec "web-$i" -- sh -c 'echo "${hostname}" > /usr/share/nginx/html/i
for i in 0 1; do kubectl exec -i -t "web-$i" -- curl http://localhost/; done
```

```
web-0
web-1
```

Note:

If you instead see **403 Forbidden** responses for the above curl command, you will need to fix the permissions of the directory mounted by the `volumeMounts` (due to a [bug when using hostPath volumes](https://github.com/kubernetes/kubernetes/issues/2630) (<https://github.com/kubernetes/kubernetes/issues/2630>)), by running:

```
for i in 0 1; do kubectl exec web-$i -- chmod 755 /usr/share/nginx/html; done
```

before retrying the `curl` command above.

In one terminal, watch the StatefulSet's Pods:

```
kubectl get pod -w -l app=nginx
```

In a second terminal, delete all of the StatefulSet's Pods:

```
kubectl delete pod -l app=nginx
```

```
pod "web-0" deleted
pod "web-1" deleted
```

Examine the output of the `kubectl get` command in the first terminal, and wait for all of the Pods to transition to Running and Ready.

```
kubectl get pod -w -l app=nginx
```

NAME	READY	STATUS	RESTARTS	AGE
web-0	0/1	ContainerCreating	0	0s
NAME	READY	STATUS	RESTARTS	AGE
web-0	1/1	Running	0	2s
web-1	0/1	Pending	0	0s
web-1	0/1	Pending	0	0s
web-1	0/1	ContainerCreating	0	0s
web-1	1/1	Running	0	34s

Verify the web servers continue to serve their hostnames:

```
for i in 0 1; do kubectl exec -i -t "web-$i" -- curl http://localhost/; done
```

```
web-0
web-1
```


Even though `web-0` and `web-1` were rescheduled, they continue to serve their hostnames because the PersistentVolumes associated with their PersistentVolumeClaims are remounted to their `volumeMounts` . No matter what node `web-0` and `web-1` are scheduled on, their PersistentVolumes will be mounted to the appropriate mount points.

Scaling a StatefulSet

Scaling a StatefulSet refers to increasing or decreasing the number of replicas. This is accomplished by updating the `replicas` field. You can use either `kubectl scale` (</docs/reference/generated/kubectl/kubectl-commands/#scale>) or `kubectl patch` (</docs/reference/generated/kubectl/kubectl-commands/#patch>) to scale a StatefulSet.

Scaling Up

In one terminal window, watch the Pods in the StatefulSet:

```
kubectl get pods -w -l app=nginx
```

In another terminal window, use `kubectl scale` to scale the number of replicas to 5:

```
kubectl scale sts web --replicas=5
```

```
statefulset.apps/web scaled
```

Examine the output of the `kubectl get` command in the first terminal, and wait for the three additional Pods to transition to Running and Ready.

```
kubectl get pods -w -l app=nginx
```

NAME	READY	STATUS	RESTARTS	AGE	
web-0	1/1	Running	0	2h	
web-1	1/1	Running	0	2h	
NAME	READY	STATUS	RESTARTS	AGE	
web-2	0/1	Pending	0	0s	
web-2	0/1	Pending	0	0s	
web-2	0/1	ContainerCreating	0	0s	0s
web-2	1/1	Running	0	19s	
web-3	0/1	Pending	0	0s	
web-3	0/1	Pending	0	0s	
web-3	0/1	ContainerCreating	0	0s	0s
web-3	1/1	Running	0	18s	
web-4	0/1	Pending	0	0s	
web-4	0/1	Pending	0	0s	
web-4	0/1	ContainerCreating	0	0s	0s
web-4	1/1	Running	0	19s	

The StatefulSet controller scaled the number of replicas. As with [StatefulSet creation](#), the StatefulSet controller created each Pod sequentially with respect to its ordinal index, and it waited for each Pod's predecessor to be Running and Ready before launching the subsequent Pod.

Scaling Down

In one terminal, watch the StatefulSet's Pods:

```
kubectl get pods -w -l app=nginx
```

In another terminal, use `kubectl patch` to scale the StatefulSet back down to three replicas:

```
kubectl patch sts web -p '{"spec":{"replicas":3}}'
```

```
statefulset.apps/web patched
```

Wait for `web-4` and `web-3` to transition to `Terminating`.

```
kubectl get pods -w -l app=nginx
```

NAME	READY	STATUS	RESTARTS	AGE
web-0	1/1	Running	0	3h
web-1	1/1	Running	0	3h
web-2	1/1	Running	0	55s
web-3	1/1	Running	0	36s
web-4	0/1	ContainerCreating	0	18s
NAME	READY	STATUS	RESTARTS	AGE
web-4	1/1	Running	0	19s
web-4	1/1	Terminating	0	24s
web-4	1/1	Terminating	0	24s
web-3	1/1	Terminating	0	42s
web-3	1/1	Terminating	0	42s

Ordered Pod Termination

The controller deleted one Pod at a time, in reverse order with respect to its ordinal index, and it waited for each to be completely shutdown before deleting the next.

Get the StatefulSet's PersistentVolumeClaims:

```
kubectl get pvc -l app=nginx
```

NAME	STATUS	VOLUME	CAPACITY	ACCESSMODES	AGE
www-web-0	Bound	pvc-15c268c7-b507-11e6-932f-42010a800002	1Gi	RW0	13h
www-web-1	Bound	pvc-15c79307-b507-11e6-932f-42010a800002	1Gi	RW0	13h
www-web-2	Bound	pvc-e1125b27-b508-11e6-932f-42010a800002	1Gi	RW0	13h
www-web-3	Bound	pvc-e1176df6-b508-11e6-932f-42010a800002	1Gi	RW0	13h
www-web-4	Bound	pvc-e11bb5f8-b508-11e6-932f-42010a800002	1Gi	RW0	13h

There are still five PersistentVolumeClaims and five PersistentVolumes. When exploring a Pod's [stable storage](#), we saw that the PersistentVolumes mounted to the Pods of a StatefulSet are not deleted when the StatefulSet's Pods are deleted. This is still true when Pod deletion is caused by scaling the StatefulSet down.

Updating StatefulSets

In Kubernetes 1.7 and later, the StatefulSet controller supports automated updates. The strategy used is determined by the `spec.updateStrategy` field of the StatefulSet API Object. This feature can be used to upgrade the container images, resource requests and/or limits, labels, and annotations of the Pods in a StatefulSet. There are two valid update strategies, `RollingUpdate` and `OnDelete`.

`RollingUpdate` update strategy is the default for StatefulSets.

Rolling Update

The `RollingUpdate` update strategy will update all Pods in a StatefulSet, in reverse ordinal order, while respecting the StatefulSet guarantees.

Patch the `web` StatefulSet to apply the `RollingUpdate` update strategy:

```
kubectl patch statefulset web -p '{"spec":{"updateStrategy":{"type":"RollingUpdate"}}}'
```

```
statefulset.apps/web patched
```

In one terminal window, patch the `web` StatefulSet to change the container image again:

```
kubectl patch statefulset web --type='json' -p='[{"op": "replace", "path": "/spec/template/sp
```

```
statefulset.apps/web patched
```

In another terminal, watch the Pods in the StatefulSet:

```
kubectl get pod -l app=nginx -w
```

The output is similar to:

NAME	READY	STATUS	RESTARTS	AGE	
web-0	1/1	Running	0	7m	
web-1	1/1	Running	0	7m	
web-2	1/1	Running	0	8m	
web-2	1/1	Terminating	0	8m	
web-2	1/1	Terminating	0	8m	
web-2	0/1	Terminating	0	8m	
web-2	0/1	Terminating	0	8m	
web-2	0/1	Terminating	0	8m	
web-2	0/1	Terminating	0	8m	
web-2	0/1	Pending	0	0s	
web-2	0/1	Pending	0	0s	
web-2	0/1	ContainerCreating	0		0s
web-2	1/1	Running	0	19s	
web-1	1/1	Terminating	0	8m	
web-1	0/1	Terminating	0	8m	
web-1	0/1	Terminating	0	8m	
web-1	0/1	Terminating	0	8m	
web-1	0/1	Pending	0	0s	
web-1	0/1	Pending	0	0s	
web-1	0/1	ContainerCreating	0		0s
web-1	1/1	Running	0	6s	
web-0	1/1	Terminating	0	7m	
web-0	1/1	Terminating	0	7m	
web-0	0/1	Terminating	0	7m	
web-0	0/1	Terminating	0	7m	
web-0	0/1	Terminating	0	7m	
web-0	0/1	Terminating	0	7m	
web-0	0/1	Pending	0	0s	
web-0	0/1	Pending	0	0s	
web-0	0/1	ContainerCreating	0		0s
web-0	1/1	Running	0	10s	

The Pods in the StatefulSet are updated in reverse ordinal order. The StatefulSet controller terminates each Pod, and waits for it to transition to Running and Ready prior to updating the next Pod. Note that, even though the StatefulSet controller will not proceed to update the next Pod until its ordinal successor is Running and Ready, it will restore any Pod that fails during the update to its current version.

Pods that have already received the update will be restored to the updated version, and Pods that have not yet received the update will be restored to the previous version. In this way, the controller attempts to continue to keep the application healthy and the update consistent in the presence of intermittent failures.

Get the Pods to view their container images:

```
for p in 0 1 2; do kubectl get pod "web-$p" --template '{{range $i, $c := .spec.containers}}{{
```

```
k8s.gcr.io/nginx-slim:0.8
k8s.gcr.io/nginx-slim:0.8
k8s.gcr.io/nginx-slim:0.8
```

All the Pods in the StatefulSet are now running the previous container image.

Note: You can also use `kubectl rollout status sts/<name>` to view the status of a rolling update to a StatefulSet

Staging an Update

You can stage an update to a StatefulSet by using the `partition` parameter of the `RollingUpdate` update strategy. A staged update will keep all of the Pods in the StatefulSet at the current version while allowing mutations to the StatefulSet's `.spec.template`.

Patch the `web` StatefulSet to add a partition to the `updateStrategy` field:

```
kubectl patch statefulset web -p '{"spec":{"updateStrategy":{"type":"RollingUpdate","rollingU
```

```
statefulset.apps/web patched
```

Patch the StatefulSet again to change the container's image:

```
kubectl patch statefulset web --type='json' -p='[{"op": "replace", "path": "/spec/template/sp
```

```
statefulset.apps/web patched
```

Delete a Pod in the StatefulSet:

```
kubectl delete pod web-2
```

```
pod "web-2" deleted
```

Wait for the Pod to be Running and Ready.

```
kubectl get pod -l app=nginx -w
```

NAME	READY	STATUS	RESTARTS	AGE
web-0	1/1	Running	0	4m
web-1	1/1	Running	0	4m
web-2	0/1	ContainerCreating	0	11s
web-2	1/1	Running	0	18s

Get the Pod's container image:

```
kubectl get pod web-2 --template '{{range $i, $c := .spec.containers}}{{ $c.image }}{{end}}'
```

```
k8s.gcr.io/nginx-slim:0.8
```

Notice that, even though the update strategy is `RollingUpdate` the StatefulSet restored the Pod with its original container. This is because the ordinal of the Pod is less than the `partition` specified by the `updateStrategy`.

Rolling Out a Canary

You can roll out a canary to test a modification by decrementing the `partition` you specified [above](#).

Patch the StatefulSet to decrement the partition:

```
kubectl patch statefulset web -p '{"spec":{"updateStrategy":{"type":"RollingUpdate"},"rollingU
```

```
statefulset.apps/web patched
```

Wait for `web-2` to be Running and Ready.

```
kubectl get pod -l app=nginx -w
```

NAME	READY	STATUS	RESTARTS	AGE
web-0	1/1	Running	0	4m
web-1	1/1	Running	0	4m
web-2	0/1	ContainerCreating	0	11s
web-2	1/1	Running 0	18s	

Get the Pod's container:

```
kubectl get pod web-2 --template '{{range $i, $c := .spec.containers}}{{$.image}}{{end}}'
```

```
k8s.gcr.io/nginx-slim:0.7
```

When you changed the `partition`, the StatefulSet controller automatically updated the `web-2` Pod because the Pod's ordinal was greater than or equal to the `partition`.

Delete the `web-1` Pod:

```
kubectl delete pod web-1
```

```
pod "web-1" deleted
```

Wait for the `web-1` Pod to be Running and Ready.

```
kubectl get pod -l app=nginx -w
```

The output is similar to:

NAME	READY	STATUS	RESTARTS	AGE
web-0	1/1	Running	0	6m
web-1	0/1	Terminating	0	6m
web-2	1/1	Running	0	2m
web-1	0/1	Terminating	0	6m
web-1	0/1	Terminating	0	6m
web-1	0/1	Terminating	0	6m
web-1	0/1	Pending 0	0s	
web-1	0/1	Pending 0	0s	
web-1	0/1	ContainerCreating	0	0s
web-1	1/1	Running 0	18s	

Get the `web-1` Pod's container image:

```
kubectl get pod web-1 --template '{{range $i, $c := .spec.containers}}{{$.image}}{{end}}'
```

```
k8s.gcr.io/nginx-slim:0.8
```

web-1 was restored to its original configuration because the Pod's ordinal was less than the partition. When a partition is specified, all Pods with an ordinal that is greater than or equal to the partition will be updated when the StatefulSet's `.spec.template` is updated. If a Pod that has an ordinal less than the partition is deleted or otherwise terminated, it will be restored to its original configuration.

Phased Roll Outs

You can perform a phased roll out (e.g. a linear, geometric, or exponential roll out) using a partitioned rolling update in a similar manner to how you rolled out a [canary](#). To perform a phased roll out, set the `partition` to the ordinal at which you want the controller to pause the update.

The partition is currently set to `2`. Set the partition to `0`:

```
kubectl patch statefulset web -p '{"spec":{"updateStrategy":{"type":"RollingUpdate","rollingU
```

```
statefulset.apps/web patched
```

Wait for all of the Pods in the StatefulSet to become Running and Ready.

```
kubectl get pod -l app=nginx -w
```

The output is similar to:

NAME	READY	STATUS	RESTARTS	AGE
web-0	1/1	Running	0	3m
web-1	0/1	ContainerCreating	0	11s
web-2	1/1	Running	0	2m
web-1	1/1	Running 0	18s	
web-0	1/1	Terminating 0	3m	
web-0	1/1	Terminating 0	3m	
web-0	0/1	Terminating 0	3m	
web-0	0/1	Terminating 0	3m	
web-0	0/1	Terminating 0	3m	
web-0	0/1	Terminating 0	3m	
web-0	0/1	Pending 0	0s	
web-0	0/1	Pending 0	0s	
web-0	0/1	ContainerCreating 0	0s	
web-0	1/1	Running 0	3s	

Get the container image details for the Pods in the StatefulSet:

```
for p in 0 1 2; do kubectl get pod "web-$p" --template '{{range $i, $c := .spec.containers}}{{
```

```
k8s.gcr.io/nginx-slim:0.7
k8s.gcr.io/nginx-slim:0.7
k8s.gcr.io/nginx-slim:0.7
```

By moving the `partition` to `0`, you allowed the StatefulSet to continue the update process.

On Delete

The `OnDelete` update strategy implements the legacy (1.6 and prior) behavior. When you select this update strategy, the StatefulSet controller will not automatically update Pods when a modification is made to the StatefulSet's `.spec.template` field. This strategy can be selected by setting the `.spec.template.updateStrategy.type` to `OnDelete`.

Deleting StatefulSets

StatefulSet supports both Non-Cascading and Cascading deletion. In a Non-Cascading Delete, the StatefulSet's Pods are not deleted when the StatefulSet is deleted. In a Cascading Delete, both the StatefulSet and its Pods are deleted.

Non-Cascading Delete

In one terminal window, watch the Pods in the StatefulSet.

```
kubectl get pods -w -l app=nginx
```

Use `kubectl delete` (</docs/reference/generated/kubectl/kubectl-commands/#delete>) to delete the StatefulSet. Make sure to supply the `--cascade=orphan` parameter to the command. This parameter tells Kubernetes to only delete the StatefulSet, and to not delete any of its Pods.

```
kubectl delete statefulset web --cascade=orphan
```

```
statefulset.apps "web" deleted
```

Get the Pods, to examine their status:

```
kubectl get pods -l app=nginx
```

NAME	READY	STATUS	RESTARTS	AGE
web-0	1/1	Running	0	6m
web-1	1/1	Running	0	7m
web-2	1/1	Running	0	5m

Even though `web` has been deleted, all of the Pods are still Running and Ready. Delete `web-0` :

```
kubectl delete pod web-0
```

```
pod "web-0" deleted
```

Get the StatefulSet's Pods:

```
kubectl get pods -l app=nginx
```

NAME	READY	STATUS	RESTARTS	AGE
web-1	1/1	Running	0	10m
web-2	1/1	Running	0	7m

As the `web` StatefulSet has been deleted, `web-0` has not been relaunched.

In one terminal, watch the StatefulSet's Pods.

```
kubectl get pods -w -l app=nginx
```

In a second terminal, recreate the StatefulSet. Note that, unless you deleted the `nginx` Service (which you should not have), you will see an error indicating that the Service already exists.

```
kubectl apply -f web.yaml
```

```
statefulset.apps/web created
service/nginx unchanged
```

Ignore the error. It only indicates that an attempt was made to create the *nginx* headless Service even though that Service already exists.

Examine the output of the `kubectl get` command running in the first terminal.

```
kubectl get pods -w -l app=nginx
```

NAME	READY	STATUS	RESTARTS	AGE
web-1	1/1	Running	0	16m
web-2	1/1	Running	0	2m
NAME	READY	STATUS	RESTARTS	AGE
web-0	0/1	Pending	0	0s
web-0	0/1	Pending	0	0s
web-0	0/1	ContainerCreating	0	0s
web-0	1/1	Running	0	18s
web-2	1/1	Terminating	0	3m
web-2	0/1	Terminating	0	3m
web-2	0/1	Terminating	0	3m
web-2	0/1	Terminating	0	3m

When the `web` StatefulSet was recreated, it first relaunched `web-0`. Since `web-1` was already Running and Ready, when `web-0` transitioned to Running and Ready, it adopted this Pod. Since you recreated the StatefulSet with `replicas` equal to 2, once `web-0` had been recreated, and once `web-1` had been determined to already be Running and Ready, `web-2` was terminated.

Let's take another look at the contents of the `index.html` file served by the Pods' web servers:

```
for i in 0 1; do kubectl exec -i -t "web-$i" -- curl http://localhost/; done
```

```
web-0
web-1
```

Even though you deleted both the StatefulSet and the `web-0` Pod, it still serves the hostname originally entered into its `index.html` file. This is because the StatefulSet never deletes the PersistentVolumes associated with a Pod. When you recreated the StatefulSet and it relaunched `web-0`, its original PersistentVolume was remounted.

Cascading Delete

In one terminal window, watch the Pods in the StatefulSet.

```
kubectl get pods -w -l app=nginx
```

In another terminal, delete the StatefulSet again. This time, omit the `--cascade=orphan` parameter.

```
kubectl delete statefulset web
```

```
statefulset.apps "web" deleted
```

Examine the output of the `kubectl get` command running in the first terminal, and wait for all of the Pods to transition to Terminating.

```
kubectl get pods -w -l app=nginx
```


NAME	READY	STATUS	RESTARTS	AGE
web-0	1/1	Running	0	11m
web-1	1/1	Running	0	27m

NAME	READY	STATUS	RESTARTS	AGE
web-0	1/1	Terminating	0	12m
web-1	1/1	Terminating	0	29m
web-0	0/1	Terminating	0	12m
web-0	0/1	Terminating	0	12m
web-0	0/1	Terminating	0	12m
web-1	0/1	Terminating	0	29m
web-1	0/1	Terminating	0	29m
web-1	0/1	Terminating	0	29m

As you saw in the [Scaling Down](#) section, the Pods are terminated one at a time, with respect to the reverse order of their ordinal indices. Before terminating a Pod, the StatefulSet controller waits for the Pod's successor to be completely terminated.

Note: Although a cascading delete removes a StatefulSet together with its Pods, the cascade does not delete the headless Service associated with the StatefulSet. You must delete the `nginx` Service manually.

```
kubectl delete service nginx
```

```
service "nginx" deleted
```

Recreate the StatefulSet and headless Service one more time:

```
kubectl apply -f web.yaml
```

```
service/nginx created
statefulset.apps/web created
```

When all of the StatefulSet's Pods transition to Running and Ready, retrieve the contents of their `index.html` files:

```
for i in 0 1; do kubectl exec -i -t "web-$i" -- curl http://localhost/; done
```

```
web-0
web-1
```

Even though you completely deleted the StatefulSet, and all of its Pods, the Pods are recreated with their PersistentVolumes mounted, and `web-0` and `web-1` continue to serve their hostnames.

Finally, delete the `nginx` Service...

```
kubectl delete service nginx
```

```
service "nginx" deleted
```

...and the `web` StatefulSet:

```
kubectl delete statefulset web
```

```
statefulset "web" deleted
```

Pod Management Policy

For some distributed systems, the StatefulSet ordering guarantees are unnecessary and/or undesirable. These systems require only uniqueness and identity. To address this, in Kubernetes 1.7, we introduced `.spec.podManagementPolicy` to the StatefulSet API Object.

OrderedReady Pod Management

OrderedReady pod management is the default for StatefulSets. It tells the StatefulSet controller to respect the ordering guarantees demonstrated above.

Parallel Pod Management

Parallel pod management tells the StatefulSet controller to launch or terminate all Pods in parallel, and not to wait for Pods to become Running and Ready or completely terminated prior to launching or terminating another Pod. This option only affects the behavior for scaling operations. Updates are not affected.

[application/web/web-parallel.yaml](https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/web/web-parallel.yaml)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/web/web-parallel.yaml>). 

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  ports:
    - port: 80
      name: web
  clusterIP: None
  selector:
    app: nginx
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  serviceName: "nginx"
  podManagementPolicy: "Parallel"
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: k8s.gcr.io/nginx-slim:0.8
          ports:
            - containerPort: 80
              name: web
          volumeMounts:
            - name: www
              mountPath: /usr/share/nginx/html
  volumeClaimTemplates:
    - metadata:
        name: www
      spec:
        accessModes: [ "ReadWriteOnce" ]
        resources:
          requests:
            storage: 1Gi
```

Download the example above, and save it to a file named `web-parallel.yaml`

This manifest is identical to the one you downloaded above except that the `.spec.podManagementPolicy` of the `web StatefulSet` is set to `Parallel`.

In one terminal, watch the Pods in the StatefulSet.

```
kubectl get pod -l app=nginx -w
```

In another terminal, create the StatefulSet and Service in the manifest:

```
kubectl apply -f web-parallel.yaml
```

```
service/nginx created
statefulset.apps/web created
```

Examine the output of the `kubectl get` command that you executed in the first terminal.

```
kubectl get pod -l app=nginx -w
```

NAME	READY	STATUS	RESTARTS	AGE	
web-0	0/1	Pending	0	0s	
web-0	0/1	Pending	0	0s	
web-1	0/1	Pending	0	0s	
web-1	0/1	Pending	0	0s	
web-0	0/1	ContainerCreating	0	0s	
web-1	0/1	ContainerCreating	0	0s	
web-0	1/1	Running	0	10s	
web-1	1/1	Running	0	10s	

The StatefulSet controller launched both `web-0` and `web-1` at the same time.

Keep the second terminal open, and, in another terminal window scale the StatefulSet:

```
kubectl scale statefulset/web --replicas=4
```

```
statefulset.apps/web scaled
```

Examine the output of the terminal where the `kubectl get` command is running.

web-3	0/1	Pending	0	0s	
web-3	0/1	Pending	0	0s	
web-3	0/1	Pending	0	7s	
web-3	0/1	ContainerCreating	0	7s	
web-2	1/1	Running	0	10s	
web-3	1/1	Running	0	26s	

The StatefulSet launched two new Pods, and it did not wait for the first to become Running and Ready prior to launching the second.

Cleaning up

You should have two terminals open, ready for you to run `kubectl` commands as part of cleanup.

```
kubectl delete sts web
# sts is an abbreviation for statefulset
```

You can watch `kubectl get` to see those Pods being deleted.

```
kubectl get pod -l app=nginx -w
```

web-3	1/1	Terminating	0	9m
web-2	1/1	Terminating	0	9m
web-3	1/1	Terminating	0	9m
web-2	1/1	Terminating	0	9m
web-1	1/1	Terminating	0	44m
web-0	1/1	Terminating	0	44m
web-0	0/1	Terminating	0	44m
web-3	0/1	Terminating	0	9m
web-2	0/1	Terminating	0	9m
web-1	0/1	Terminating	0	44m
web-0	0/1	Terminating	0	44m
web-2	0/1	Terminating	0	9m
web-2	0/1	Terminating	0	9m
web-2	0/1	Terminating	0	9m
web-1	0/1	Terminating	0	44m
web-1	0/1	Terminating	0	44m
web-1	0/1	Terminating	0	44m
web-0	0/1	Terminating	0	44m
web-0	0/1	Terminating	0	44m
web-0	0/1	Terminating	0	44m
web-3	0/1	Terminating	0	9m
web-3	0/1	Terminating	0	9m
web-3	0/1	Terminating	0	9m

During deletion, a StatefulSet removes all Pods concurrently; it does not wait for a Pod's ordinal successor to terminate prior to deleting that Pod.

Close the terminal where the `kubectl get` command is running and delete the `nginx` Service:

```
kubectl delete svc nginx
```

Note:

You also need to delete the persistent storage media for the PersistentVolumes used in this tutorial.

Follow the necessary steps, based on your environment, storage configuration, and provisioning method, to ensure that all storage is reclaimed.

5.2 - Example: Deploying WordPress and MySQL with Persistent Volumes

This tutorial shows you how to deploy a WordPress site and a MySQL database using Minikube. Both applications use PersistentVolumes and PersistentVolumeClaims to store data.

A [PersistentVolume](/docs/concepts/storage/persistent-volumes/) (PV) is a piece of storage in the cluster that has been manually provisioned by an administrator, or dynamically provisioned by Kubernetes using a [StorageClass](/docs/concepts/storage/storage-classes/). A [PersistentVolumeClaim](/docs/concepts/storage/persistent-volumes/#persistentvolumeclaims) (PVC) is a request for storage by a user that can be fulfilled by a PV. PersistentVolumes and PersistentVolumeClaims are independent from Pod lifecycles and preserve data through restarting, rescheduling, and even deleting Pods.

Warning: This deployment is not suitable for production use cases, as it uses single instance WordPress and MySQL Pods. Consider using [WordPress Helm Chart](https://github.com/kubernetes/charts/tree/master/stable/wordpress) (<https://github.com/kubernetes/charts/tree/master/stable/wordpress>) to deploy WordPress in production.

Note: The files provided in this tutorial are using GA Deployment APIs and are specific to Kubernetes version 1.9 and later. If you wish to use this tutorial with an earlier version of Kubernetes, please update the API version appropriately, or reference earlier versions of this tutorial.

Objectives

- Create PersistentVolumeClaims and PersistentVolumes
- Create a `kustomization.yaml` with
 - a Secret generator
 - MySQL resource configs
 - WordPress resource configs
- Apply the kustomization directory by `kubectl apply -k ./`
- Clean up

Before you begin

You need to have a Kubernetes cluster, and the `kubectl` command-line tool must be configured to communicate with your cluster. It is recommended to run this tutorial on a cluster with at least two nodes that are not acting as control plane hosts. If you do not already have a cluster, you can create one by using [minikube](https://minikube.sigs.k8s.io/docs/tutorials/multi_node/) or you can use one of these Kubernetes playgrounds:

- [Katacoda](https://www.katacoda.com/courses/kubernetes/playground)
- [Play with Kubernetes](http://labs.play-with-k8s.com/)

To check the version, enter `kubectl version`. The example shown on this page works with `kubectl 1.14` and above.

Download the following configuration files:

1. [mysql-deployment.yaml](/examples/application/wordpress/mysql-deployment.yaml)
2. [wordpress-deployment.yaml](/examples/application/wordpress/wordpress-deployment.yaml)

Create PersistentVolumeClaims and PersistentVolumes

MySQL and Wordpress each require a PersistentVolume to store data. Their PersistentVolumeClaims will be created at the deployment step.

Many cluster environments have a default StorageClass installed. When a StorageClass is not specified in the PersistentVolumeClaim, the cluster's default StorageClass is used instead.

When a PersistentVolumeClaim is created, a PersistentVolume is dynamically provisioned based on the StorageClass configuration.

Warning: In local clusters, the default StorageClass uses the `hostPath` provisioner. `hostPath` volumes are only suitable for development and testing. With `hostPath` volumes, your data lives in `/tmp` on the node the Pod is scheduled onto and does not move between nodes. If a Pod dies and gets scheduled to another node in the cluster, or the node is rebooted, the data is lost.

Note: If you are bringing up a cluster that needs to use the `hostPath` provisioner, the `--enable-hostpath-provisioner` flag must be set in the `controller-manager` component.

Note: If you have a Kubernetes cluster running on Google Kubernetes Engine, please follow [this guide \(https://cloud.google.com/kubernetes-engine/docs/tutorials/persistent-disk\)](https://cloud.google.com/kubernetes-engine/docs/tutorials/persistent-disk).

Create a kustomization.yaml

Add a Secret generator

A [Secret \(/docs/concepts/configuration/secret/\)](/docs/concepts/configuration/secret/) is an object that stores a piece of sensitive data like a password or key. Since 1.14, `kubectl` supports the management of Kubernetes objects using a kustomization file. You can create a Secret by generators in `kustomization.yaml`.

Add a Secret generator in `kustomization.yaml` from the following command. You will need to replace `YOUR_PASSWORD` with the password you want to use.

```
cat <<EOF >./kustomization.yaml
secretGenerator:
- name: mysql-pass
  literals:
  - password=YOUR_PASSWORD
EOF
```

Add resource configs for MySQL and WordPress

The following manifest describes a single-instance MySQL Deployment. The MySQL container mounts the PersistentVolume at `/var/lib/mysql`. The `MYSQL_ROOT_PASSWORD` environment variable sets the database password from the Secret.

application/wordpress/mysql-deployment.yaml
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/wordpress/mysql-deployment.yaml>)

```
apiVersion: v1
kind: Service
metadata:
  name: wordpress-mysql
  labels:
    app: wordpress
spec:
  ports:
    - port: 3306
  selector:
    app: wordpress
    tier: mysql
  clusterIP: None
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mysql-pv-claim
  labels:
    app: wordpress
```

```
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 20Gi
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wordpress-mysql
  labels:
    app: wordpress
spec:
  selector:
    matchLabels:
      app: wordpress
      tier: mysql
  strategy:
    type: Recreate
  template:
    metadata:
      labels:
        app: wordpress
        tier: mysql
    spec:
      containers:
        - image: mysql:5.6
          name: mysql
          env:
            - name: MYSQL_ROOT_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: mysql-pass
                  key: password
          ports:
            - containerPort: 3306
              name: mysql
          volumeMounts:
            - name: mysql-persistent-storage
              mountPath: /var/lib/mysql
      volumes:
        - name: mysql-persistent-storage
          persistentVolumeClaim:
            claimName: mysql-pv-claim
```

The following manifest describes a single-instance WordPress Deployment. The WordPress container mounts the PersistentVolume at `/var/www/html` for website data files. The `WORDPRESS_DB_HOST` environment variable sets the name of the MySQL Service defined above, and WordPress will access the database by Service. The `WORDPRESS_DB_PASSWORD` environment variable sets the database password from the Secret `kustomize` generated.

[application/wordpress/wordpress-deployment.yaml](#)

(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/wordpress/wordpress-deployment.yaml>) 

```
apiVersion: v1
kind: Service
metadata:
  name: wordpress
  labels:
    app: wordpress
spec:
  ports:
    - port: 80
  selector:
    app: wordpress
    tier: frontend
  type: LoadBalancer
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: wp-pv-claim
  labels:
    app: wordpress
spec:
```

```

accessModes:
  - ReadWriteOnce
resources:
  requests:
    storage: 20Gi
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wordpress
  labels:
    app: wordpress
spec:
  selector:
    matchLabels:
      app: wordpress
      tier: frontend
  strategy:
    type: Recreate
  template:
    metadata:
      labels:
        app: wordpress
        tier: frontend
    spec:
      containers:
        - image: wordpress:4.8-apache
          name: wordpress
          env:
            - name: WORDPRESS_DB_HOST
              value: wordpress-mysql
            - name: WORDPRESS_DB_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: mysql-pass
                  key: password
          ports:
            - containerPort: 80
              name: wordpress
          volumeMounts:
            - name: wordpress-persistent-storage
              mountPath: /var/www/html
      volumes:
        - name: wordpress-persistent-storage
          persistentVolumeClaim:
            claimName: wp-pv-claim

```

1. Download the MySQL deployment configuration file.

```
curl -LO https://k8s.io/examples/application/wordpress/mysql-deployment.yaml
```

2. Download the WordPress configuration file.

```
curl -LO https://k8s.io/examples/application/wordpress/wordpress-deployment.yaml
```

3. Add them to `kustomization.yaml` file.

```

cat <<EOF >>./kustomization.yaml
resources:
  - mysql-deployment.yaml
  - wordpress-deployment.yaml
EOF

```

Apply and Verify

The `kustomization.yaml` contains all the resources for deploying a WordPress site and a MySQL database. You can apply the directory by


```
kubectl apply -k ./
```

Now you can verify that all objects exist.

1. Verify that the Secret exists by running the following command:

```
kubectl get secrets
```

The response should be like this:

NAME	TYPE	DATA	AGE
mysql-pass-c57bb4t7mf	Opaque	1	9s

2. Verify that a PersistentVolume got dynamically provisioned.

```
kubectl get pvc
```

Note: It can take up to a few minutes for the PVs to be provisioned and bound.

The response should be like this:

NAME	STATUS	VOLUME	CAPACITY	ACCESS
mysql-pv-claim	Bound	pvc-8cbd7b2e-4044-11e9-b2bb-42010a800002	20Gi	RWO
wp-pv-claim	Bound	pvc-8cd0df54-4044-11e9-b2bb-42010a800002	20Gi	RWO

3. Verify that the Pod is running by running the following command:

```
kubectl get pods
```

Note: It can take up to a few minutes for the Pod's Status to be **RUNNING**.

The response should be like this:

NAME	READY	STATUS	RESTARTS	AGE
wordpress-mysql-1894417608-x5dzt	1/1	Running	0	40s

4. Verify that the Service is running by running the following command:

```
kubectl get services wordpress
```

The response should be like this:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
wordpress	LoadBalancer	10.0.0.89	<pending>	80:32406/TCP	4m

Note: Minikube can only expose Services through **NodePort**. The **EXTERNAL-IP** is always pending.

5. Run the following command to get the IP Address for the WordPress Service:

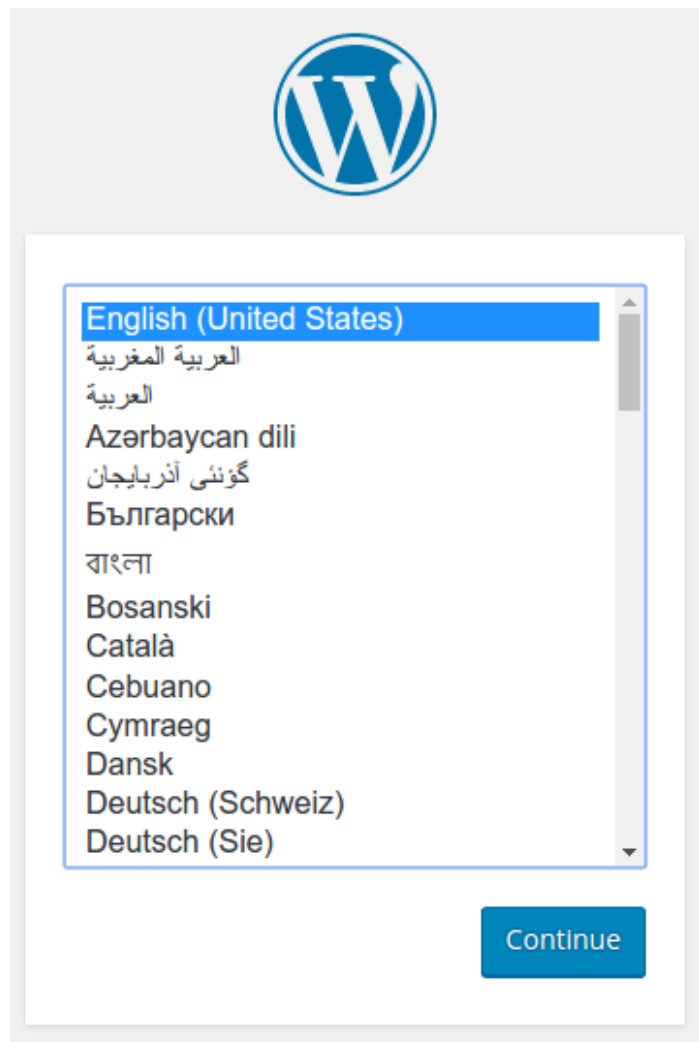
```
minikube service wordpress --url
```

The response should be like this:

```
http://1.2.3.4:32406
```

6. Copy the IP address, and load the page in your browser to view your site.

You should see the WordPress set up page similar to the following screenshot.



Warning: Do not leave your WordPress installation on this page. If another user finds it, they can set up a website on your instance and use it to serve malicious content.

Either install WordPress by creating a username and password or delete your instance.

Cleaning up

1. Run the following command to delete your Secret, Deployments, Services and PersistentVolumeClaims:

```
kubectl delete -k ./
```

What's next

- Learn more about [Introspection and Debugging](/docs/tasks/debug-application-cluster/debug-application-introspection/) (/docs/tasks/debug-application-cluster/debug-application-introspection/)
- Learn more about [Jobs](/docs/concepts/workloads/controllers/job/) (/docs/concepts/workloads/controllers/job/)
- Learn more about [Port Forwarding](/docs/tasks/access-application-cluster/port-forward-access-application-cluster/) (/docs/tasks/access-application-cluster/port-forward-access-application-cluster/)
- Learn how to [Get a Shell to a Container](/docs/tasks/debug-application-cluster/get-shell-running-container/) (/docs/tasks/debug-application-cluster/get-shell-running-container/)

5.3 - Example: Deploying Cassandra with a StatefulSet

This tutorial shows you how to run [Apache Cassandra](https://cassandra.apache.org/) on Kubernetes. Cassandra, a database, needs persistent storage to provide data durability (application *state*). In this example, a custom Cassandra seed provider lets the database discover new Cassandra instances as they join the Cassandra cluster.

StatefulSets make it easier to deploy stateful applications into your Kubernetes cluster. For more information on the features used in this tutorial, see [StatefulSet \(/docs/concepts/workloads/controllers/statefulset/\)](/docs/concepts/workloads/controllers/statefulset/).

Note:

Cassandra and Kubernetes both use the term *node* to mean a member of a cluster. In this tutorial, the Pods that belong to the StatefulSet are Cassandra nodes and are members of the Cassandra cluster (called a *ring*). When those Pods run in your Kubernetes cluster, the Kubernetes control plane schedules those Pods onto Kubernetes Nodes (</docs/concepts/architecture/nodes/>).

When a Cassandra node starts, it uses a *seed list* to bootstrap discovery of other nodes in the ring. This tutorial deploys a custom Cassandra seed provider that lets the database discover new Cassandra Pods as they appear inside your Kubernetes cluster.

Objectives

- Create and validate a Cassandra headless Service (</docs/concepts/services-networking/service/>).
- Use a StatefulSet (</docs/concepts/workloads/controllers/statefulset/>) to create a Cassandra ring.
- Validate the StatefulSet.
- Modify the StatefulSet.
- Delete the StatefulSet and its Pods (</docs/concepts/workloads/pods/>).

Before you begin

You need to have a Kubernetes cluster, and the `kubectl` command-line tool must be configured to communicate with your cluster. It is recommended to run this tutorial on a cluster with at least two nodes that are not acting as control plane hosts. If you do not already have a cluster, you can create one by using [minikube](https://minikube.sigs.k8s.io/docs/tutorials/multi_node/) or you can use one of these Kubernetes playgrounds:

- [Katacoda](https://www.katacoda.com/courses/kubernetes/playground)
- [Play with Kubernetes](http://labs.play-with-k8s.com/)

To complete this tutorial, you should already have a basic familiarity with [Pods](/docs/concepts/workloads/pods/), [Services](/docs/concepts/services-networking/service/), and [StatefulSets](/docs/concepts/workloads/controllers/statefulset/).

Additional Minikube setup instructions

Caution:

[Minikube](https://minikube.sigs.k8s.io/docs/) defaults to 2048MB of memory and 2 CPU. Running Minikube with the default resource configuration results in insufficient resource errors during this tutorial. To avoid these errors, start Minikube with the following settings:

```
minikube start --memory 5120 --cpus=4
```

Creating a headless Service for Cassandra

In Kubernetes, a [Service \(/docs/concepts/services-networking/service/\)](/docs/concepts/services-networking/service/) describes a set of [Pods \(/docs/concepts/workloads/pods/\)](/docs/concepts/workloads/pods/) that perform the same task.

The following Service is used for DNS lookups between Cassandra Pods and clients within your cluster:

[application/cassandra/cassandra-service.yaml](https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/cassandra/cassandra-service.yaml)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/cassandra/cassandra-service.yaml>)

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: cassandra
    name: cassandra
spec:
  clusterIP: None
  ports:
    - port: 9042
  selector:
    app: cassandra
```

Create a Service to track all Cassandra StatefulSet members from the `cassandra-service.yaml` file:

```
kubectl apply -f https://k8s.io/examples/application/cassandra/cassandra-service.yaml
```

Validating (optional)

Get the Cassandra Service.

```
kubectl get svc cassandra
```

The response is

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
cassandra	ClusterIP	None	<none>	9042/TCP	45s

If you don't see a Service named `cassandra`, that means creation failed. Read [Debug Services \(/docs/tasks/debug-application-cluster/debug-service/\)](/docs/tasks/debug-application-cluster/debug-service/) for help troubleshooting common issues.

Using a StatefulSet to create a Cassandra ring

The StatefulSet manifest, included below, creates a Cassandra ring that consists of three Pods.

Note: This example uses the default provisioner for Minikube. Please update the following StatefulSet for the cloud you are working with.

[application/cassandra/cassandra-statefulset.yaml](https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/cassandra/cassandra-statefulset.yaml)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/cassandra/cassandra-statefulset.yaml>)

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
```

```
name: cassandra
labels:
  app: cassandra
spec:
  serviceName: cassandra
  replicas: 3
  selector:
    matchLabels:
      app: cassandra
  template:
    metadata:
      labels:
        app: cassandra
    spec:
      terminationGracePeriodSeconds: 1800
      containers:
        - name: cassandra
          image: gcr.io/google-samples/cassandra:v13
          imagePullPolicy: Always
          ports:
            - containerPort: 7000
              name: intra-node
            - containerPort: 7001
              name: tls-intra-node
            - containerPort: 7199
              name: jmx
            - containerPort: 9042
              name: cql
          resources:
            limits:
              cpu: "500m"
              memory: 1Gi
            requests:
              cpu: "500m"
              memory: 1Gi
          securityContext:
            capabilities:
              add:
                - IPC_LOCK
          lifecycle:
            preStop:
              exec:
                command:
                  - /bin/sh
                  - -c
                  - nodetool drain
          env:
            - name: MAX_HEAP_SIZE
              value: 512M
            - name: HEAP_NEWSIZE
              value: 100M
            - name: CASSANDRA_SEEDS
              value: "cassandra-0.cassandra.default.svc.cluster.local"
            - name: CASSANDRA_CLUSTER_NAME
              value: "K8Demo"
            - name: CASSANDRA_DC
              value: "DC1-K8Demo"
            - name: CASSANDRA_RACK
              value: "Rack1-K8Demo"
            - name: POD_IP
              valueFrom:
                fieldRef:
                  fieldPath: status.podIP
          readinessProbe:
            exec:
              command:
                - /bin/bash
                - -c
                - /ready-probe.sh
            initialDelaySeconds: 15
            timeoutSeconds: 5
            # These volume mounts are persistent. They are like inline claims,
            # but not exactly because the names need to match exactly one of
            # the stateful pod volumes.
          volumeMounts:
            - name: cassandra-data
              mountPath: /cassandra_data
            # These are converted to volume claims by the controller
            # and mounted at the paths mentioned above.
            # do not use these in production until ssd GCEPersistentDisk or other ssd pd
          volumeClaimTemplates:
            - metadata:
                name: cassandra-data
              spec:
                accessModes: [ "ReadWriteOnce" ]
                storageClassName: fast
```

```
resources:
  requests:
    storage: 1Gi
---
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: fast
provisioner: k8s.io/minikube-hostpath
parameters:
  type: pd-ssd
```

Create the Cassandra StatefulSet from the `cassandra-statefulset.yaml` file:

```
# Use this if you are able to apply cassandra-statefulset.yaml unmodified
kubectl apply -f https://k8s.io/examples/application/cassandra/cassandra-statefulset.yaml
```

If you need to modify `cassandra-statefulset.yaml` to suit your cluster, download <https://k8s.io/examples/application/cassandra/cassandra-statefulset.yaml> (<https://k8s.io/examples/application/cassandra/cassandra-statefulset.yaml>) and then apply that manifest, from the folder you saved the modified version into:

```
# Use this if you needed to modify cassandra-statefulset.yaml locally
kubectl apply -f cassandra-statefulset.yaml
```

Validating the Cassandra StatefulSet

1. Get the Cassandra StatefulSet:

```
kubectl get statefulset cassandra
```

The response should be similar to:

NAME	DESIRED	CURRENT	AGE
cassandra	3	0	13s

The `StatefulSet` resource deploys Pods sequentially.

2. Get the Pods to see the ordered creation status:

```
kubectl get pods -l="app=cassandra"
```

The response should be similar to:

NAME	READY	STATUS	RESTARTS	AGE
cassandra-0	1/1	Running	0	1m
cassandra-1	0/1	ContainerCreating	0	8s

It can take several minutes for all three Pods to deploy. Once they are deployed, the same command returns output similar to:

NAME	READY	STATUS	RESTARTS	AGE
cassandra-0	1/1	Running	0	10m
cassandra-1	1/1	Running	0	9m
cassandra-2	1/1	Running	0	8m

3. Run the Cassandra `nodetool` (<https://cwiki.apache.org/confluence/display/CASSANDRA2/NodeTool>) inside the first Pod, to display the status of the ring.

```
kubectl exec -it cassandra-0 -- nodetool status
```

The response should look something like:

```
Datcenter: DC1-K8Demo
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
-- Address      Load          Tokens         Owns (effective)  Host ID
UN  172.17.0.5    83.57 KiB     32             74.0%             e2dd09e6-d9d3-477e-96c5-45094c
UN  172.17.0.4    101.04 KiB    32             58.8%             f89d6835-3a42-4419-92b3-0e62c
UN  172.17.0.6    84.74 KiB     32             67.1%             a6a1e8c2-3dc5-4417-b1a0-26507a
```

Modifying the Cassandra StatefulSet

Use `kubectl edit` to modify the size of a Cassandra StatefulSet.

1. Run the following command:

```
kubectl edit statefulset cassandra
```

This command opens an editor in your terminal. The line you need to change is the `replicas` field. The following sample is an excerpt of the StatefulSet file:

```
# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will
# reopened with the relevant failures.
#
apiVersion: apps/v1
kind: StatefulSet
metadata:
  creationTimestamp: 2016-08-13T18:40:58Z
  generation: 1
  labels:
    app: cassandra
    name: cassandra
    namespace: default
    resourceVersion: "323"
    uid: 7a219483-6185-11e6-a910-42010a8a0fc0
spec:
  replicas: 3
```

2. Change the number of replicas to 4, and then save the manifest.

The StatefulSet now scales to run with 4 Pods.

3. Get the Cassandra StatefulSet to verify your change:

```
kubectl get statefulset cassandra
```

The response should be similar to:

```
NAME          DESIRED  CURRENT  AGE
cassandra     4        4        36m
```

Cleaning up

Deleting or scaling a StatefulSet down does not delete the volumes associated with the StatefulSet. This setting is for your safety because your data is more valuable than automatically purging all related StatefulSet resources.

Warning: Depending on the storage class and reclaim policy, deleting the *PersistentVolumeClaims* may cause the associated volumes to also be deleted. Never assume you'll be able to access data if its volume claims are deleted.

1. Run the following commands (chained together into a single command) to delete everything in the Cassandra StatefulSet:

```
grace=$(kubectl get pod cassandra-0 -o=jsonpath='{.spec.terminationGracePeriodSeconds}')
&& kubectl delete statefulset -l app=cassandra \
&& echo "Sleeping ${grace} seconds" 1>&2 \
&& sleep $grace \
&& kubectl delete persistentvolumeclaim -l app=cassandra
```

2. Run the following command to delete the Service you set up for Cassandra:

```
kubectl delete service -l app=cassandra
```

Cassandra container environment variables

The Pods in this tutorial use the `gcr.io/google-samples/cassandra:v13` (<https://github.com/kubernetes/examples/blob/master/cassandra/image/Dockerfile>) image from Google's [container registry](https://cloud.google.com/container-registry/docs/) (<https://cloud.google.com/container-registry/docs/>). The Docker image above is based on `debian-base` (<https://github.com/kubernetes/release/tree/master/images/build/debian-base>) and includes OpenJDK 8.

This image includes a standard Cassandra installation from the Apache Debian repo. By using environment variables you can change values that are inserted into `cassandra.yaml`.

Environment variable	Default value
CASSANDRA_CLUSTER_NAME	'Test Cluster'
CASSANDRA_NUM_TOKENS	32
CASSANDRA_RPC_ADDRESS	0.0.0.0

What's next

- Learn how to [Scale a StatefulSet](/docs/tasks/run-application/scale-stateful-set/) (</docs/tasks/run-application/scale-stateful-set/>).
- Learn more about the *KubernetesSeedProvider* (<https://github.com/kubernetes/examples/blob/master/cassandra/java/src/main/java/io/k8s/cassandra/KubernetesSeedProvider.java>)
- See more custom [Seed Provider Configurations](#) (<https://git.k8s.io/examples/cassandra/java/README.md>)

5.4 - Running ZooKeeper, A Distributed System Coordinator

This tutorial demonstrates running [Apache Zookeeper](https://zookeeper.apache.org) (<https://zookeeper.apache.org>) on Kubernetes using [StatefulSets](/docs/concepts/workloads/controllers/statefulset/) (</docs/concepts/workloads/controllers/statefulset/>), [PodDisruptionBudgets](/docs/concepts/workloads/pods/disruptions/#pod-disruption-budget) (</docs/concepts/workloads/pods/disruptions/#pod-disruption-budget>), and [PodAntiAffinity](/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity) (</docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity>).

Before you begin

Before starting this tutorial, you should be familiar with the following Kubernetes concepts:

- [Pods](/docs/concepts/workloads/pods/) (</docs/concepts/workloads/pods/>)
- [Cluster DNS](/docs/concepts/services-networking/dns-pod-service/) (</docs/concepts/services-networking/dns-pod-service/>)
- [Headless Services](/docs/concepts/services-networking/service/#headless-services) (</docs/concepts/services-networking/service/#headless-services>)
- [PersistentVolumes](/docs/concepts/storage/volumes/) (</docs/concepts/storage/volumes/>)
- [PersistentVolume Provisioning](https://github.com/kubernetes/examples/tree/master/staging/persistent-volume-provisioning/) (<https://github.com/kubernetes/examples/tree/master/staging/persistent-volume-provisioning/>)
- [StatefulSets](/docs/concepts/workloads/controllers/statefulset/) (</docs/concepts/workloads/controllers/statefulset/>)
- [PodDisruptionBudgets](/docs/concepts/workloads/pods/disruptions/#pod-disruption-budget) (</docs/concepts/workloads/pods/disruptions/#pod-disruption-budget>)
- [PodAntiAffinity](/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity) (</docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity>)
- [kubectl CLI](/docs/reference/kubectl/kubectl/) (</docs/reference/kubectl/kubectl/>)

You must have a cluster with at least four nodes, and each node requires at least 2 CPUs and 4 GiB of memory. In this tutorial you will cordon and drain the cluster's nodes. **This means that the cluster will terminate and evict all Pods on its nodes, and the nodes will temporarily become unschedulable.** You should use a dedicated cluster for this tutorial, or you should ensure that the disruption you cause will not interfere with other tenants.

This tutorial assumes that you have configured your cluster to dynamically provision PersistentVolumes. If your cluster is not configured to do so, you will have to manually provision three 20 GiB volumes before starting this tutorial.

Objectives

After this tutorial, you will know the following.

- How to deploy a ZooKeeper ensemble using StatefulSet.
- How to consistently configure the ensemble.
- How to spread the deployment of ZooKeeper servers in the ensemble.
- How to use PodDisruptionBudgets to ensure service availability during planned maintenance.

ZooKeeper

[Apache ZooKeeper](https://zookeeper.apache.org/doc/current/) (<https://zookeeper.apache.org/doc/current/>) is a distributed, open-source coordination service for distributed applications. ZooKeeper allows you to read, write, and observe updates to data. Data are organized in a file system like hierarchy and replicated to all ZooKeeper servers in the ensemble (a set of ZooKeeper servers). All operations on data are atomic and sequentially consistent. ZooKeeper ensures this by using the [Zab](https://pdfs.semanticscholar.org/b02c/6b00bd5dbdbd951fddb00b906c82fa80f0b3.pdf) (<https://pdfs.semanticscholar.org/b02c/6b00bd5dbdbd951fddb00b906c82fa80f0b3.pdf>) consensus protocol to replicate a state machine across all servers in the ensemble.

The ensemble uses the Zab protocol to elect a leader, and the ensemble cannot write data until that election is complete. Once complete, the ensemble uses Zab to ensure that it replicates all writes to a quorum before it acknowledges and makes them visible to clients. Without respect to weighted quorums, a quorum is a majority component of the ensemble containing the current leader. For instance, if the ensemble has three servers, a component that contains the leader and one other server constitutes a quorum. If the ensemble can not achieve a quorum, the ensemble cannot write data.

ZooKeeper servers keep their entire state machine in memory, and write every mutation to a durable WAL (Write Ahead Log) on storage media. When a server crashes, it can recover its previous state by replaying the WAL. To prevent the WAL from growing without bound, ZooKeeper servers will periodically snapshot them in memory state to storage media. These snapshots can be loaded directly into memory, and all WAL entries that preceded the snapshot may be discarded.

Creating a ZooKeeper ensemble

The manifest below contains a [Headless Service \(/docs/concepts/services-networking/service/#headless-services\)](/docs/concepts/services-networking/service/#headless-services), a [Service \(/docs/concepts/services-networking/service/\)](/docs/concepts/services-networking/service/), a [PodDisruptionBudget \(/docs/concepts/workloads/pods/disruptions/#pod-disruption-budgets\)](/docs/concepts/workloads/pods/disruptions/#pod-disruption-budgets), and a [StatefulSet \(/docs/concepts/workloads/controllers/statefulset/\)](/docs/concepts/workloads/controllers/statefulset/).

[application/zookeeper/zookeeper.yaml](#)

(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/application/zookeeper/zookeeper.yaml>)



```
apiVersion: v1
kind: Service
metadata:
  name: zk-hs
  labels:
    app: zk
spec:
  ports:
    - port: 2888
      name: server
    - port: 3888
      name: leader-election
  clusterIP: None
  selector:
    app: zk
---
apiVersion: v1
kind: Service
metadata:
  name: zk-cs
  labels:
    app: zk
spec:
  ports:
    - port: 2181
      name: client
  selector:
    app: zk
---
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: zk-pdb
spec:
  selector:
    matchLabels:
      app: zk
  maxUnavailable: 1
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: zk
spec:
  selector:
    matchLabels:
      app: zk
  serviceName: zk-hs
  replicas: 3
  updateStrategy:
    type: RollingUpdate
  podManagementPolicy: OrderedReady
  template:
    metadata:
      labels:
        app: zk
    spec:
      affinity:
        podAntiAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
```

```

    - labelSelector:
      matchExpressions:
        - key: "app"
          operator: In
          values:
            - zk
      topologyKey: "kubernetes.io/hostname"
  containers:
  - name: kubernetes-zookeeper
    imagePullPolicy: Always
    image: "k8s.gcr.io/kubernetes-zookeeper:1.0-3.4.10"
    resources:
      requests:
        memory: "1Gi"
        cpu: "0.5"
    ports:
      - containerPort: 2181
        name: client
      - containerPort: 2888
        name: server
      - containerPort: 3888
        name: leader-election
    command:
      - sh
      - -c
      - "start-zookeeper \
        --servers=3 \
        --data_dir=/var/lib/zookeeper/data \
        --data_log_dir=/var/lib/zookeeper/data/log \
        --conf_dir=/opt/zookeeper/conf \
        --client_port=2181 \
        --election_port=3888 \
        --server_port=2888 \
        --tick_time=2000 \
        --init_limit=10 \
        --sync_limit=5 \
        --heap=512M \
        --max_client_cnxns=60 \
        --snap_retain_count=3 \
        --purge_interval=12 \
        --max_session_timeout=40000 \
        --min_session_timeout=4000 \
        --log_level=INFO"
    readinessProbe:
      exec:
        command:
          - sh
          - -c
          - "zookeeper-ready 2181"
        initialDelaySeconds: 10
        timeoutSeconds: 5
    livenessProbe:
      exec:
        command:
          - sh
          - -c
          - "zookeeper-ready 2181"
        initialDelaySeconds: 10
        timeoutSeconds: 5
    volumeMounts:
      - name: datadir
        mountPath: /var/lib/zookeeper
  securityContext:
    runAsUser: 1000
    fsGroup: 1000
  volumeClaimTemplates:
  - metadata:
      name: datadir
    spec:
      accessModes: [ "ReadWriteOnce" ]
      resources:
        requests:
          storage: 10Gi

```

Open a terminal, and use the `kubectl apply` [./docs/reference/generated/kubect/kubectl-commands/#apply](#) command to create the manifest.

```
kubectl apply -f https://k8s.io/examples/application/zookeeper/zookeeper.yaml
```

This creates the `zk-hs` Headless Service, the `zk-cs` Service, the `zk-pdb` PodDisruptionBudget, and the `zk` StatefulSet.

```
service/zk-hs created
service/zk-cs created
poddisruptionbudget.policy/zk-pdb created
statefulset.apps/zk created
```

Use `kubectl get` ([./docs/reference/generated/kubectl/kubectl-commands/#get](#)) to watch the StatefulSet controller create the StatefulSet's Pods.

```
kubectl get pods -w -l app=zk
```

Once the `zk-2` Pod is Running and Ready, use `CTRL-C` to terminate `kubectl`.

NAME	READY	STATUS	RESTARTS	AGE	
zk-0	0/1	Pending	0	0s	
zk-0	0/1	Pending	0	0s	
zk-0	0/1	ContainerCreating	0	0s	
zk-0	0/1	Running	0	19s	
zk-0	1/1	Running	0	40s	
zk-1	0/1	Pending	0	0s	
zk-1	0/1	Pending	0	0s	
zk-1	0/1	ContainerCreating	0	0s	
zk-1	0/1	Running	0	18s	
zk-1	1/1	Running	0	40s	
zk-2	0/1	Pending	0	0s	
zk-2	0/1	Pending	0	0s	
zk-2	0/1	ContainerCreating	0	0s	
zk-2	0/1	Running	0	19s	
zk-2	1/1	Running	0	40s	

The StatefulSet controller creates three Pods, and each Pod has a container with a [ZooKeeper](#) (<https://www-us.apache.org/dist/zookeeper/stable/>) server.

Facilitating leader election

Because there is no terminating algorithm for electing a leader in an anonymous network, Zab requires explicit membership configuration to perform leader election. Each server in the ensemble needs to have a unique identifier, all servers need to know the global set of identifiers, and each identifier needs to be associated with a network address.

Use `kubectl exec` ([./docs/reference/generated/kubectl/kubectl-commands/#exec](#)) to get the hostnames of the Pods in the `zk` StatefulSet.

```
for i in 0 1 2; do kubectl exec zk-$i -- hostname; done
```

The StatefulSet controller provides each Pod with a unique hostname based on its ordinal index. The hostnames take the form of `<statefulset name>-<ordinal index>`. Because the `replicas` field of the `zk` StatefulSet is set to `3`, the Set's controller creates three Pods with their hostnames set to `zk-0`, `zk-1`, and `zk-2`.

```
zk-0
zk-1
zk-2
```

The servers in a ZooKeeper ensemble use natural numbers as unique identifiers, and store each server's identifier in a file called `myid` in the server's data directory.

To examine the contents of the `myid` file for each server use the following command.

```
for i in 0 1 2; do echo "myid zk-$i"; kubectl exec zk-$i -- cat /var/lib/zookeeper/data/myid; done
```

Because the identifiers are natural numbers and the ordinal indices are non-negative integers, you can generate an identifier by adding 1 to the ordinal.

```
myid zk-0
1
myid zk-1
2
myid zk-2
3
```

To get the Fully Qualified Domain Name (FQDN) of each Pod in the `zk` StatefulSet use the following command.

```
for i in 0 1 2; do kubectl exec zk-$i -- hostname -f; done
```

The `zk-hs` Service creates a domain for all of the Pods, `zk-hs.default.svc.cluster.local`.

```
zk-0.zk-hs.default.svc.cluster.local
zk-1.zk-hs.default.svc.cluster.local
zk-2.zk-hs.default.svc.cluster.local
```

The A records in [Kubernetes DNS \(/docs/concepts/services-networking/dns-pod-service/\)](/docs/concepts/services-networking/dns-pod-service/) resolve the FQDNs to the Pods' IP addresses. If Kubernetes reschedules the Pods, it will update the A records with the Pods' new IP addresses, but the A records names will not change.

ZooKeeper stores its application configuration in a file named `zoo.cfg`. Use `kubectl exec` to view the contents of the `zoo.cfg` file in the `zk-0` Pod.

```
kubectl exec zk-0 -- cat /opt/zookeeper/conf/zoo.cfg
```

In the `server.1`, `server.2`, and `server.3` properties at the bottom of the file, the `1`, `2`, and `3` correspond to the identifiers in the ZooKeeper servers' `myid` files. They are set to the FQDNs for the Pods in the `zk` StatefulSet.

```
clientPort=2181
dataDir=/var/lib/zookeeper/data
dataLogDir=/var/lib/zookeeper/log
tickTime=2000
initLimit=10
syncLimit=2000
maxClientCnxns=60
minSessionTimeout= 4000
maxSessionTimeout= 40000
autopurge.snapRetainCount=3
autopurge.purgeInterval=0
server.1=zk-0.zk-hs.default.svc.cluster.local:2888:3888
server.2=zk-1.zk-hs.default.svc.cluster.local:2888:3888
server.3=zk-2.zk-hs.default.svc.cluster.local:2888:3888
```

Achieving consensus

Consensus protocols require that the identifiers of each participant be unique. No two participants in the Zab protocol should claim the same unique identifier. This is necessary to allow the processes in the system to agree on which processes have committed which data. If two Pods are launched with the same ordinal, two ZooKeeper servers would both identify themselves as the same server.

```
kubectl get pods -w -l app=zk
```

NAME	READY	STATUS	RESTARTS	AGE	
zk-0	0/1	Pending	0	0s	
zk-0	0/1	Pending	0	0s	
zk-0	0/1	ContainerCreating	0	0s	
zk-0	0/1	Running	0	19s	
zk-0	1/1	Running	0	40s	
zk-1	0/1	Pending	0	0s	
zk-1	0/1	Pending	0	0s	
zk-1	0/1	ContainerCreating	0	0s	
zk-1	0/1	Running	0	18s	
zk-1	1/1	Running	0	40s	
zk-2	0/1	Pending	0	0s	
zk-2	0/1	Pending	0	0s	
zk-2	0/1	ContainerCreating	0	0s	
zk-2	0/1	Running	0	19s	
zk-2	1/1	Running	0	40s	

The A records for each Pod are entered when the Pod becomes Ready. Therefore, the FQDNs of the ZooKeeper servers will resolve to a single endpoint, and that endpoint will be the unique ZooKeeper server claiming the identity configured in its `myid` file.

```
zk-0.zk-hs.default.svc.cluster.local
zk-1.zk-hs.default.svc.cluster.local
zk-2.zk-hs.default.svc.cluster.local
```

This ensures that the `servers` properties in the ZooKeepers' `zoo.cfg` files represents a correctly configured ensemble.

```
server.1=zk-0.zk-hs.default.svc.cluster.local:2888:3888
server.2=zk-1.zk-hs.default.svc.cluster.local:2888:3888
server.3=zk-2.zk-hs.default.svc.cluster.local:2888:3888
```

When the servers use the Zab protocol to attempt to commit a value, they will either achieve consensus and commit the value (if leader election has succeeded and at least two of the Pods are Running and Ready), or they will fail to do so (if either of the conditions are not met). No state will arise where one server acknowledges a write on behalf of another.

Sanity testing the ensemble

The most basic sanity test is to write data to one ZooKeeper server and to read the data from another.

The command below executes the `zkCli.sh` script to write `world` to the path `/hello` on the `zk-0` Pod in the ensemble.

```
kubectl exec zk-0 zkCli.sh create /hello world
```

```
WATCHER::

WatchedEvent state:SyncConnected type:None path:null
Created /hello
```

To get the data from the `zk-1` Pod use the following command.

```
kubectl exec zk-1 zkCli.sh get /hello
```

The data that you created on `zk-0` is available on all the servers in the ensemble.

```
WATCHER::
```

```
WatchedEvent state:SyncConnected type:None path:null
world
cZxid = 0x100000002
ctime = Thu Dec 08 15:13:30 UTC 2016
mZxid = 0x100000002
mtime = Thu Dec 08 15:13:30 UTC 2016
pZxid = 0x100000002
cversion = 0
dataVersion = 0
aclVersion = 0
ephemeralOwner = 0x0
dataLength = 5
numChildren = 0
```

Providing durable storage

As mentioned in the [ZooKeeper Basics](#) section, ZooKeeper commits all entries to a durable WAL, and periodically writes snapshots in memory state, to storage media. Using WALs to provide durability is a common technique for applications that use consensus protocols to achieve a replicated state machine.

Use the `kubectl delete` ([/docs/reference/generated/kubectl/kubectl-commands/#delete](#)) command to delete the `zk` StatefulSet.

```
kubectl delete statefulset zk
```

```
statefulset.apps "zk" deleted
```

Watch the termination of the Pods in the StatefulSet.

```
kubectl get pods -w -l app=zk
```

When `zk-0` is fully terminated, use `CTRL-C` to terminate `kubectl`.

zk-2	1/1	Terminating	0	9m
zk-0	1/1	Terminating	0	11m
zk-1	1/1	Terminating	0	10m
zk-2	0/1	Terminating	0	9m
zk-2	0/1	Terminating	0	9m
zk-2	0/1	Terminating	0	9m
zk-1	0/1	Terminating	0	10m
zk-1	0/1	Terminating	0	10m
zk-1	0/1	Terminating	0	10m
zk-0	0/1	Terminating	0	11m
zk-0	0/1	Terminating	0	11m
zk-0	0/1	Terminating	0	11m

Reapply the manifest in `zookeeper.yaml`.

```
kubectl apply -f https://k8s.io/examples/application/zookeeper/zookeeper.yaml
```

This creates the `zk` StatefulSet object, but the other API objects in the manifest are not modified because they already exist.

Watch the StatefulSet controller recreate the StatefulSet's Pods.

```
kubectl get pods -w -l app=zk
```

Once the `zk-2` Pod is Running and Ready, use `CTRL-C` to terminate `kubectl`.

NAME	READY	STATUS	RESTARTS	AGE	
zk-0	0/1	Pending	0	0s	
zk-0	0/1	Pending	0	0s	
zk-0	0/1	ContainerCreating	0	0s	
zk-0	0/1	Running	0	19s	
zk-0	1/1	Running	0	40s	
zk-1	0/1	Pending	0	0s	
zk-1	0/1	Pending	0	0s	
zk-1	0/1	ContainerCreating	0	0s	
zk-1	0/1	Running	0	18s	
zk-1	1/1	Running	0	40s	
zk-2	0/1	Pending	0	0s	
zk-2	0/1	Pending	0	0s	
zk-2	0/1	ContainerCreating	0	0s	
zk-2	0/1	Running	0	19s	
zk-2	1/1	Running	0	40s	

Use the command below to get the value you entered during the sanity test, from the zk-2 Pod.

```
kubectl exec zk-2 zkCli.sh get /hello
```

Even though you terminated and recreated all of the Pods in the zk StatefulSet, the ensemble still serves the original value.

```
WATCHER::

WatchedEvent state:SyncConnected type:None path:null
world
cZxid = 0x100000002
ctime = Thu Dec 08 15:13:30 UTC 2016
mZxid = 0x100000002
mtime = Thu Dec 08 15:13:30 UTC 2016
pZxid = 0x100000002
cversion = 0
dataVersion = 0
aclVersion = 0
ephemeralOwner = 0x0
dataLength = 5
numChildren = 0
```

The volumeClaimTemplates field of the zk StatefulSet's spec specifies a PersistentVolume provisioned for each Pod.

```
volumeClaimTemplates:
- metadata:
  name: datadir
  annotations:
    volume.alpha.kubernetes.io/storage-class: anything
  spec:
    accessModes: [ "ReadWriteOnce" ]
    resources:
      requests:
        storage: 20Gi
```

The StatefulSet controller generates a PersistentVolumeClaim for each Pod in the StatefulSet . Use the following command to get the StatefulSet 's PersistentVolumeClaims .

```
kubectl get pvc -l app=zk
```

When the StatefulSet recreated its Pods, it remounts the Pods' PersistentVolumes.

NAME	STATUS	VOLUME	CAPACITY	ACCESSMODES	
datadir-zk-0	Bound	pvc-bed742cd-bcb1-11e6-994f-42010a800002	20Gi	RWO	/
datadir-zk-1	Bound	pvc-bedd27d2-bcb1-11e6-994f-42010a800002	20Gi	RWO	/
datadir-zk-2	Bound	pvc-bee0817e-bcb1-11e6-994f-42010a800002	20Gi	RWO	/

The volumeMounts section of the StatefulSet 's container template mounts the PersistentVolumes in the ZooKeeper servers' data directories.


```
volumeMounts:
- name: datadir
  mountPath: /var/lib/zookeeper
```

When a Pod in the `zk StatefulSet` is (re)scheduled, it will always have the same `PersistentVolume` mounted to the ZooKeeper server's data directory. Even when the Pods are rescheduled, all the writes made to the ZooKeeper servers' WALs, and all their snapshots, remain durable.

Ensuring consistent configuration

As noted in the [Facilitating Leader Election](#) and [Achieving Consensus](#) sections, the servers in a ZooKeeper ensemble require consistent configuration to elect a leader and form a quorum. They also require consistent configuration of the Zab protocol in order for the protocol to work correctly over a network. In our example we achieve consistent configuration by embedding the configuration directly into the manifest.

Get the `zk StatefulSet`.

```
kubectl get sts zk -o yaml
```

```
...
command:
- sh
- -c
- "start-zookeeper \
  --servers=3 \
  --data_dir=/var/lib/zookeeper/data \
  --data_log_dir=/var/lib/zookeeper/data/log \
  --conf_dir=/opt/zookeeper/conf \
  --client_port=2181 \
  --election_port=3888 \
  --server_port=2888 \
  --tick_time=2000 \
  --init_limit=10 \
  --sync_limit=5 \
  --heap=512M \
  --max_client_cnxns=60 \
  --snap_retain_count=3 \
  --purge_interval=12 \
  --max_session_timeout=40000 \
  --min_session_timeout=4000 \
  --log_level=INFO"
...
```

The command used to start the ZooKeeper servers passed the configuration as command line parameter. You can also use environment variables to pass configuration to the ensemble.

Configuring logging

One of the files generated by the `zkGenConfig.sh` script controls ZooKeeper's logging. ZooKeeper uses [Log4j](https://logging.apache.org/log4j/2.x/), and, by default, it uses a time and size based rolling file appender for its logging configuration.

Use the command below to get the logging configuration from one of Pods in the `zk StatefulSet`.

```
kubectl exec zk-0 cat /usr/etc/zookeeper/log4j.properties
```

The logging configuration below will cause the ZooKeeper process to write all of its logs to the standard output file stream.

```
zookeeper.root.logger=CONSOLE
zookeeper.console.threshold=INFO
log4j.rootLogger=${zookeeper.root.logger}
log4j.appender.CONSOLE=org.apache.log4j.ConsoleAppender
log4j.appender.CONSOLE.Threshold=${zookeeper.console.threshold}
log4j.appender.CONSOLE.layout=org.apache.log4j.PatternLayout
log4j.appender.CONSOLE.layout.ConversionPattern=%d{ISO8601} [myid:%X{myid}] - %-5p [%t:%C{1}]@%s
```

This is the simplest possible way to safely log inside the container. Because the applications write logs to standard out, Kubernetes will handle log rotation for you. Kubernetes also implements a sane retention policy that ensures application logs written to standard out and standard error do not exhaust local storage media.

Use `kubectl logs` (</docs/reference/generated/kubectl/kubectl-commands/#logs>) to retrieve the last 20 log lines from one of the Pods.

```
kubectl logs zk-0 --tail 20
```

You can view application logs written to standard out or standard error using `kubectl logs` and from the Kubernetes Dashboard.

```
2016-12-06 19:34:16,236 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:16,237 [myid:1] - INFO [Thread-1136:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:26,155 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:26,155 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:26,156 [myid:1] - INFO [Thread-1137:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:26,222 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:26,222 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:26,226 [myid:1] - INFO [Thread-1138:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:36,151 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:36,152 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:36,152 [myid:1] - INFO [Thread-1139:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:36,230 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:36,231 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:36,231 [myid:1] - INFO [Thread-1140:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:46,149 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:46,149 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:46,149 [myid:1] - INFO [Thread-1141:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:46,230 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:46,230 [myid:1] - INFO [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
2016-12-06 19:34:46,230 [myid:1] - INFO [Thread-1142:NIOServerCxn@1008] - Closed socket connection from client /10.0.0.1:49152
```

Kubernetes integrates with many logging solutions. You can choose a logging solution that best fits your cluster and applications. For cluster-level logging and aggregation, consider deploying a [sidecar container](/docs/concepts/cluster-administration/logging#sidecar-container) (</docs/concepts/cluster-administration/logging#sidecar-container-with-logging-agent>) to rotate and ship your logs.

Configuring a non-privileged user

The best practices to allow an application to run as a privileged user inside of a container are a matter of debate. If your organization requires that applications run as a non-privileged user you can use a [SecurityContext](/docs/tasks/configure-pod-container/security-context/) (</docs/tasks/configure-pod-container/security-context/>) to control the user that the entry point runs as.

The `zk` `StatefulSet`'s `Pod` template contains a `SecurityContext`.

```
securityContext:
  runAsUser: 1000
  fsGroup: 1000
```

In the Pods' containers, UID 1000 corresponds to the `zookeeper` user and GID 1000 corresponds to the `zookeeper` group.

Get the ZooKeeper process information from the `zk-0` `Pod`.

```
kubectl exec zk-0 -- ps -elf
```

As the `runAsUser` field of the `securityContext` object is set to 1000, instead of running as root, the ZooKeeper process runs as the zookeeper user.

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	STIME	TTY	TIME	CMD
4	S	zookeeper+	1	0	0	80	0	-	1127	-	20:46	?	00:00:00	sh -c zkGenConfig.s
0	S	zookeeper+	27	1	0	80	0	-	1155556	-	20:46	?	00:00:19	/usr/lib/jvm/java-8

By default, when the Pod's PersistentVolumes is mounted to the ZooKeeper server's data directory, it is only accessible by the root user. This configuration prevents the ZooKeeper process from writing to its WAL and storing its snapshots.

Use the command below to get the file permissions of the ZooKeeper data directory on the `zk-0` Pod.

```
kubectl exec -ti zk-0 -- ls -ld /var/lib/zookeeper/data
```

Because the `fsGroup` field of the `securityContext` object is set to 1000, the ownership of the Pods' PersistentVolumes is set to the zookeeper group, and the ZooKeeper process is able to read and write its data.

```
drwxr-sr-x 3 zookeeper zookeeper 4096 Dec  5 20:45 /var/lib/zookeeper/data
```

Managing the ZooKeeper process

The [ZooKeeper documentation](https://zookeeper.apache.org/doc/current/zookeeperAdmin.html#sc_supervision) (https://zookeeper.apache.org/doc/current/zookeeperAdmin.html#sc_supervision) mentions that "You will want to have a supervisory process that manages each of your ZooKeeper server processes (JVM)." Utilizing a watchdog (supervisory process) to restart failed processes in a distributed system is a common pattern. When deploying an application in Kubernetes, rather than using an external utility as a supervisory process, you should use Kubernetes as the watchdog for your application.

Updating the ensemble

The `zk` `StatefulSet` is configured to use the `RollingUpdate` update strategy. You can use `kubectl patch` to update the number of `cpus` allocated to the servers.

```
kubectl patch sts zk --type='json' -p='[{"op": "replace", "path": "/spec/template/spec/containers/0/resources/requests/cpu", "value": "100m"}']'
```

```
statefulset.apps/zk patched
```

Use `kubectl rollout status` to watch the status of the update.

```
kubectl rollout status sts/zk
```

```
waiting for statefulset rolling update to complete 0 pods at revision zk-5db4499664...
Waiting for 1 pods to be ready...
Waiting for 1 pods to be ready...
waiting for statefulset rolling update to complete 1 pods at revision zk-5db4499664...
Waiting for 1 pods to be ready...
Waiting for 1 pods to be ready...
waiting for statefulset rolling update to complete 2 pods at revision zk-5db4499664...
Waiting for 1 pods to be ready...
Waiting for 1 pods to be ready...
statefulset rolling update complete 3 pods at revision zk-5db4499664...
```

This terminates the Pods, one at a time, in reverse ordinal order, and recreates them with the new configuration. This ensures that quorum is maintained during a rolling update.

Use the `kubectl rollout history` command to view a history or previous configurations.

```
kubectl rollout history sts/zk
```

```
statefulsets "zk"
REVISION
1
2
```

Use the `kubectl rollout undo` command to roll back the modification.

```
kubectl rollout undo sts/zk
```

```
statefulset.apps/zk rolled back
```

Handling process failure

[Restart Policies \(/docs/concepts/workloads/pods/pod-lifecycle/#restart-policy\)](#) control how Kubernetes handles process failures for the entry point of the container in a Pod. For Pods in a `StatefulSet`, the only appropriate `RestartPolicy` is `Always`, and this is the default value. For stateful applications you should **never** override the default policy.

Use the following command to examine the process tree for the ZooKeeper server running in the `zk-0` Pod.

```
kubectl exec zk-0 -- ps -ef
```

The command used as the container's entry point has PID 1, and the ZooKeeper process, a child of the entry point, has PID 27.

UID	PID	PPID	C	STIME	TTY	TIME	CMD
zookeeper+	1	0	0	15:03	?	00:00:00	sh -c zkGenConfig.sh && zkServer.sh start-foreground
zookeeper+	27	1	0	15:03	?	00:00:03	/usr/lib/jvm/java-8-openjdk-amd64/bin/java -Dzoo

In another terminal watch the Pods in the `zk` `StatefulSet` with the following command.

```
kubectl get pod -w -l app=zk
```

In another terminal, terminate the ZooKeeper process in Pod `zk-0` with the following command.

```
kubectl exec zk-0 -- pkill java
```

The termination of the ZooKeeper process caused its parent process to terminate. Because the `RestartPolicy` of the container is `Always`, it restarted the parent process.

NAME	READY	STATUS	RESTARTS	AGE
zk-0	1/1	Running	0	21m
zk-1	1/1	Running	0	20m
zk-2	1/1	Running	0	19m
NAME	READY	STATUS	RESTARTS	AGE
zk-0	0/1	Error	0	29m
zk-0	0/1	Running	1	29m
zk-0	1/1	Running	1	29m

If your application uses a script (such as `zkServer.sh`) to launch the process that implements the application's business logic, the script must terminate with the child process. This ensures that Kubernetes will restart the application's container when the process implementing the application's business logic fails.

Testing for liveness

Configuring your application to restart failed processes is not enough to keep a distributed system healthy. There are scenarios where a system's processes can be both alive and unresponsive, or otherwise unhealthy. You should use liveness probes to notify Kubernetes that your application's processes are unhealthy and it should restart them.

The Pod template for the zk StatefulSet specifies a liveness probe.

```
livenessProbe:
  exec:
    command:
      - sh
      - -c
      - "zookeeper-ready 2181"
    initialDelaySeconds: 15
    timeoutSeconds: 5
```

The probe calls a bash script that uses the ZooKeeper ruok four letter word to test the server's health.

```
OK=$(echo ruok | nc 127.0.0.1 $1)
if [ "$OK" == "imok" ]; then
  exit 0
else
  exit 1
fi
```

In one terminal window, use the following command to watch the Pods in the zk StatefulSet.

```
kubectl get pod -w -l app=zk
```

In another window, using the following command to delete the zookeeper-ready script from the file system of Pod zk-0 .

```
kubectl exec zk-0 -- rm /usr/bin/zookeeper-ready
```

When the liveness probe for the ZooKeeper process fails, Kubernetes will automatically restart the process for you, ensuring that unhealthy processes in the ensemble are restarted.

```
kubectl get pod -w -l app=zk
```

NAME	READY	STATUS	RESTARTS	AGE
zk-0	1/1	Running	0	1h
zk-1	1/1	Running	0	1h
zk-2	1/1	Running	0	1h
NAME	READY	STATUS	RESTARTS	AGE
zk-0	0/1	Running	0	1h
zk-0	0/1	Running	1	1h
zk-0	1/1	Running	1	1h

Testing for readiness

Readiness is not the same as liveness. If a process is alive, it is scheduled and healthy. If a process is ready, it is able to process input. Liveness is a necessary, but not sufficient, condition for readiness. There are cases, particularly during initialization and termination, when a process can be alive but not ready.

If you specify a readiness probe, Kubernetes will ensure that your application's processes will not receive network traffic until their readiness checks pass.

For a ZooKeeper server, liveness implies readiness. Therefore, the readiness probe from the `zookeeper.yaml` manifest is identical to the liveness probe.

```
readinessProbe:
  exec:
    command:
      - sh
      - -c
      - "zookeeper-ready 2181"
  initialDelaySeconds: 15
  timeoutSeconds: 5
```

Even though the liveness and readiness probes are identical, it is important to specify both. This ensures that only healthy servers in the ZooKeeper ensemble receive network traffic.

Tolerating Node failure

ZooKeeper needs a quorum of servers to successfully commit mutations to data. For a three server ensemble, two servers must be healthy for writes to succeed. In quorum based systems, members are deployed across failure domains to ensure availability. To avoid an outage, due to the loss of an individual machine, best practices preclude co-locating multiple instances of the application on the same machine.

By default, Kubernetes may co-locate Pods in a `StatefulSet` on the same node. For the three server ensemble you created, if two servers are on the same node, and that node fails, the clients of your ZooKeeper service will experience an outage until at least one of the Pods can be rescheduled.

You should always provision additional capacity to allow the processes of critical systems to be rescheduled in the event of node failures. If you do so, then the outage will only last until the Kubernetes scheduler reschedules one of the ZooKeeper servers. However, if you want your service to tolerate node failures with no downtime, you should set `podAntiAffinity`.

Use the command below to get the nodes for Pods in the `zk` `StatefulSet`.

```
for i in 0 1 2; do kubectl get pod zk-$i --template {{.spec.nodeName}}; echo ""; done
```

All of the Pods in the `zk` `StatefulSet` are deployed on different nodes.

```
kubernetes-node-cxpk
kubernetes-node-a5aq
kubernetes-node-2g2d
```

This is because the Pods in the `zk` `StatefulSet` have a `PodAntiAffinity` specified.

```
affinity:
  podAntiAffinity:
    requiredDuringSchedulingIgnoredDuringExecution:
      - labelSelector:
          matchExpressions:
            - key: "app"
              operator: In
              values:
                - zk
        topologyKey: "kubernetes.io/hostname"
```

The `requiredDuringSchedulingIgnoredDuringExecution` field tells the Kubernetes Scheduler that it should never co-locate two Pods which have `app` label as `zk` in the domain defined by the `topologyKey`. The `topologyKey` `kubernetes.io/hostname` indicates that the domain is an individual node. Using different rules, labels, and selectors, you can extend this technique to spread your ensemble across physical, network, and power failure domains.

Surviving maintenance

In this section you will cordon and drain nodes. If you are using this tutorial on a shared cluster, be sure that this will not adversely affect other tenants.

The previous section showed you how to spread your Pods across nodes to survive unplanned node failures, but you also need to plan for temporary node failures that occur due to planned maintenance.

Use this command to get the nodes in your cluster.

```
kubectl get nodes
```

Use `kubectl cordon` [\(/docs/reference/generated/kubectl/kubectl-commands/#cordon\)](/docs/reference/generated/kubectl/kubectl-commands/#cordon) to cordon all but four of the nodes in your cluster.

```
kubectl cordon <node-name>
```

Use this command to get the `zk-pdb` `PodDisruptionBudget` .

```
kubectl get pdb zk-pdb
```

The `max-unavailable` field indicates to Kubernetes that at most one Pod from `zk` `StatefulSet` can be unavailable at any time.

NAME	MIN-AVAILABLE	MAX-UNAVAILABLE	ALLOWED-DISRUPTIONS	AGE
zk-pdb	N/A	1	1	

In one terminal, use this command to watch the Pods in the `zk` `StatefulSet` .

```
kubectl get pods -w -l app=zk
```

In another terminal, use this command to get the nodes that the Pods are currently scheduled on.

```
for i in 0 1 2; do kubectl get pod zk-$i --template {{.spec.nodeName}}; echo ""; done
```

```
kubernetes-node-pb41
kubernetes-node-ixsl
kubernetes-node-i4c4
```

Use `kubectl drain` [\(/docs/reference/generated/kubectl/kubectl-commands/#drain\)](/docs/reference/generated/kubectl/kubectl-commands/#drain) to cordon and drain the node on which the `zk-0` Pod is scheduled.

```
kubectl drain $(kubectl get pod zk-0 --template {{.spec.nodeName}}) --ignore-daemonsets --for
```

```
node "kubernetes-node-pb41" cordoned

WARNING: Deleting pods not managed by ReplicationController, ReplicaSet, Job, or DaemonSet: f
pod "zk-0" deleted
node "kubernetes-node-pb41" drained
```

As there are four nodes in your cluster, `kubectl drain` , succeeds and the `zk-0` is rescheduled to another node.

NAME	READY	STATUS	RESTARTS	AGE
zk-0	1/1	Running	2	1h
zk-1	1/1	Running	0	1h
zk-2	1/1	Running	0	1h

NAME	READY	STATUS	RESTARTS	AGE
zk-0	1/1	Terminating	2	2h
zk-0	0/1	Terminating	2	2h
zk-0	0/1	Terminating	2	2h
zk-0	0/1	Terminating	2	2h
zk-0	0/1	Pending	0	0s
zk-0	0/1	Pending	0	0s
zk-0	0/1	ContainerCreating	0	0s
zk-0	0/1	Running	0	51s
zk-0	1/1	Running	0	1m

Keep watching the StatefulSet 's Pods in the first terminal and drain the node on which zk-1 is scheduled.

```
kubectl drain $(kubectl get pod zk-1 --template {{.spec.nodeName}}) --ignore-daemonsets --for
```

```
WARNING: Deleting pods not managed by ReplicationController, ReplicaSet, Job, or DaemonSet: f
pod "zk-1" deleted
node "kubernetes-node-ixsl" drained
```

The zk-1 Pod cannot be scheduled because the zk StatefulSet contains a PodAntiAffinity rule preventing co-location of the Pods, and as only two nodes are schedulable, the Pod will remain in a Pending state.

```
kubectl get pods -w -l app=zk
```

NAME	READY	STATUS	RESTARTS	AGE
zk-0	1/1	Running	2	1h
zk-1	1/1	Running	0	1h
zk-2	1/1	Running	0	1h

NAME	READY	STATUS	RESTARTS	AGE
zk-0	1/1	Terminating	2	2h
zk-0	0/1	Terminating	2	2h
zk-0	0/1	Terminating	2	2h
zk-0	0/1	Terminating	2	2h
zk-0	0/1	Pending	0	0s
zk-0	0/1	Pending	0	0s
zk-0	0/1	ContainerCreating	0	0s
zk-0	0/1	Running	0	51s
zk-0	1/1	Running	0	1m
zk-1	1/1	Terminating	0	2h
zk-1	0/1	Terminating	0	2h
zk-1	0/1	Terminating	0	2h
zk-1	0/1	Terminating	0	2h
zk-1	0/1	Pending	0	0s
zk-1	0/1	Pending	0	0s

Continue to watch the Pods of the stateful set, and drain the node on which zk-2 is scheduled.

```
kubectl drain $(kubectl get pod zk-2 --template {{.spec.nodeName}}) --ignore-daemonsets --for
```

```
node "kubernetes-node-i4c4" cordoned

WARNING: Deleting pods not managed by ReplicationController, ReplicaSet, Job, or DaemonSet: f
WARNING: Ignoring DaemonSet-managed pods: node-problem-detector-v0.1-dyrog; Deleting pods not
There are pending pods when an error occurred: Cannot evict pod as it would violate the pod's
pod/zk-2
```

Use CTRL-C to terminate to kubectl.

You cannot drain the third node because evicting zk-2 would violate zk-budget . However, the node will remain cordoned.

Use `zkCli.sh` to retrieve the value you entered during the sanity test from `zk-0`.

```
kubectl exec zk-0 zkCli.sh get /hello
```

The service is still available because its `PodDisruptionBudget` is respected.

```
WatchedEvent state:SyncConnected type:None path:null
world
cZxid = 0x200000002
ctime = Wed Dec 07 00:08:59 UTC 2016
mZxid = 0x200000002
mtime = Wed Dec 07 00:08:59 UTC 2016
pZxid = 0x200000002
cversion = 0
dataVersion = 0
aclVersion = 0
ephemeralOwner = 0x0
dataLength = 5
numChildren = 0
```

Use `kubectl uncordon` (</docs/reference/generated/kubectl/kubectl-commands/#uncordon>) to uncordon the first node.

```
kubectl uncordon kubernetes-node-pb41
```

```
node "kubernetes-node-pb41" uncordoned
```

`zk-1` is rescheduled on this node. Wait until `zk-1` is Running and Ready.

```
kubectl get pods -w -l app=zk
```

NAME	READY	STATUS	RESTARTS	AGE
zk-0	1/1	Running	2	1h
zk-1	1/1	Running	0	1h
zk-2	1/1	Running	0	1h
NAME	READY	STATUS	RESTARTS	AGE
zk-0	1/1	Terminating	2	2h
zk-0	0/1	Terminating	2	2h
zk-0	0/1	Terminating	2	2h
zk-0	0/1	Terminating	2	2h
zk-0	0/1	Pending	0	0s
zk-0	0/1	Pending	0	0s
zk-0	0/1	ContainerCreating	0	0s
zk-0	0/1	Running	0	51s
zk-0	1/1	Running	0	1m
zk-1	1/1	Terminating	0	2h
zk-1	0/1	Terminating	0	2h
zk-1	0/1	Terminating	0	2h
zk-1	0/1	Terminating	0	2h
zk-1	0/1	Pending	0	0s
zk-1	0/1	Pending	0	0s
zk-1	0/1	Pending	0	12m
zk-1	0/1	ContainerCreating	0	12m
zk-1	0/1	Running	0	13m
zk-1	1/1	Running	0	13m

Attempt to drain the node on which `zk-2` is scheduled.

```
kubectl drain $(kubectl get pod zk-2 --template {{.spec.nodeName}}) --ignore-daemonsets --for
```

The output:

```
node "kubernetes-node-i4c4" already cordoned
WARNING: Deleting pods not managed by ReplicationController, ReplicaSet, Job, or DaemonSet: f
pod "heapster-v1.2.0-2604621511-wh1r" deleted
pod "zk-2" deleted
node "kubernetes-node-i4c4" drained
```

This time `kubectl drain` succeeds.

Uncordon the second node to allow `zk-2` to be rescheduled.

```
kubectl uncordon kubernetes-node-ixsl
```

```
node "kubernetes-node-ixsl" uncordoned
```

You can use `kubectl drain` in conjunction with `PodDisruptionBudgets` to ensure that your services remain available during maintenance. If drain is used to cordon nodes and evict pods prior to taking the node offline for maintenance, services that express a disruption budget will have that budget respected. You should always allocate additional capacity for critical services so that their Pods can be immediately rescheduled.

Cleaning up

- Use `kubectl uncordon` to uncordon all the nodes in your cluster.
- You must delete the persistent storage media for the `PersistentVolumes` used in this tutorial. Follow the necessary steps, based on your environment, storage configuration, and provisioning method, to ensure that all storage is reclaimed.

6 - Clusters

6.1 - Restrict a Container's Access to Resources with AppArmor

FEATURE STATE: Kubernetes v1.4 [beta]

AppArmor is a Linux kernel security module that supplements the standard Linux user and group based permissions to confine programs to a limited set of resources. AppArmor can be configured for any application to reduce its potential attack surface and provide greater in-depth defense. It is configured through profiles tuned to allow the access needed by a specific program or container, such as Linux capabilities, network access, file permissions, etc. Each profile can be run in either *enforcing* mode, which blocks access to disallowed resources, or *complain* mode, which only reports violations.

AppArmor can help you to run a more secure deployment by restricting what containers are allowed to do, and/or provide better auditing through system logs. However, it is important to keep in mind that AppArmor is not a silver bullet and can only do so much to protect against exploits in your application code. It is important to provide good, restrictive profiles, and harden your applications and cluster from other angles as well.

Objectives

- See an example of how to load a profile on a node
- Learn how to enforce the profile on a Pod
- Learn how to check that the profile is loaded
- See what happens when a profile is violated
- See what happens when a profile cannot be loaded

Before you begin

Make sure:

1. Kubernetes version is at least v1.4 -- Kubernetes support for AppArmor was added in v1.4. Kubernetes components older than v1.4 are not aware of the new AppArmor annotations, and will **silently ignore** any AppArmor settings that are provided. To ensure that your Pods are receiving the expected protections, it is important to verify the Kubelet version of your nodes:

```
kubectl get nodes -o=jsonpath='${range .items[*]}{@.metadata.name}: {@.status.nodeInfo.k
```

```
gke-test-default-pool-239f5d02-gyn2: v1.4.0
gke-test-default-pool-239f5d02-x1kf: v1.4.0
gke-test-default-pool-239f5d02-xwux: v1.4.0
```

2. AppArmor kernel module is enabled -- For the Linux kernel to enforce an AppArmor profile, the AppArmor kernel module must be installed and enabled. Several distributions enable the module by default, such as Ubuntu and SUSE, and many others provide optional support. To check whether the module is enabled, check the `/sys/module/apparmor/parameters/enabled` file:

```
cat /sys/module/apparmor/parameters/enabled
Y
```

If the Kubelet contains AppArmor support (\geq v1.4), it will refuse to run a Pod with AppArmor options if the kernel module is not enabled.

Note: Ubuntu carries many AppArmor patches that have not been merged into the upstream Linux kernel, including patches that add additional hooks and features. Kubernetes has only been tested with the upstream version, and does not promise support for other features.

3. Container runtime supports AppArmor -- Currently all common Kubernetes-supported container runtimes should support AppArmor, like Docker (<https://docs.docker.com/engine/>), CRI-O (<https://cri-o.io/#what-is-cri-o>) or containerd (<https://containerd.io/docs/>). Please refer to the corresponding runtime documentation and verify that the cluster fulfills the requirements to use AppArmor.
4. Profile is loaded -- AppArmor is applied to a Pod by specifying an AppArmor profile that each container should be run with. If any of the specified profiles is not already loaded in the kernel, the Kubelet ($\geq v1.4$) will reject the Pod. You can view which profiles are loaded on a node by checking the `/sys/kernel/security/apparmor/profiles` file. For example:

```
ssh gke-test-default-pool-239f5d02-gyn2 "sudo cat /sys/kernel/security/apparmor/profiles"
```

```
apparmor-test-deny-write (enforce)
apparmor-test-audit-write (enforce)
docker-default (enforce)
k8s-nginx (enforce)
```

For more details on loading profiles on nodes, see [Setting up nodes with profiles](#).

As long as the Kubelet version includes AppArmor support ($\geq v1.4$), the Kubelet will reject a Pod with AppArmor options if any of the prerequisites are not met. You can also verify AppArmor support on nodes by checking the node ready condition message (though this is likely to be removed in a later release):

```
kubectl get nodes -o=jsonpath=${range .items[*]}{@.metadata.name}: {.status.conditions[?(@.r
```

```
gke-test-default-pool-239f5d02-gyn2: kubelet is posting ready status. AppArmor enabled
gke-test-default-pool-239f5d02-x1kf: kubelet is posting ready status. AppArmor enabled
gke-test-default-pool-239f5d02-xwux: kubelet is posting ready status. AppArmor enabled
```

Securing a Pod

Note: AppArmor is currently in beta, so options are specified as annotations. Once support graduates to general availability, the annotations will be replaced with first-class fields (more details in [Upgrade path to GA](#)).

AppArmor profiles are specified *per-container*. To specify the AppArmor profile to run a Pod container with, add an annotation to the Pod's metadata:

```
container.apparmor.security.beta.kubernetes.io/<container_name>: <profile_ref>
```

Where `<container_name>` is the name of the container to apply the profile to, and `<profile_ref>` specifies the profile to apply. The `profile_ref` can be one of:

- `runtime/default` to apply the runtime's default profile
- `localhost/<profile_name>` to apply the profile loaded on the host with the name `<profile_name>`
- `unconfined` to indicate that no profiles will be loaded

See the [API Reference](#) for the full details on the annotation and profile name formats.

Kubernetes AppArmor enforcement works by first checking that all the prerequisites have been met, and then forwarding the profile selection to the container runtime for enforcement. If the prerequisites have not been met, the Pod will be rejected, and will not run.

To verify that the profile was applied, you can look for the AppArmor security option listed in the container created event:

```
kubectl get events | grep Created
```

22s	22s	1	hello-apparmor	Pod	spec.containers{hello}	Normal
-----	-----	---	----------------	-----	------------------------	--------

You can also verify directly that the container's root process is running with the correct profile by checking its proc attr:

```
kubectl exec <pod_name> cat /proc/1/attr/current
```

```
k8s-apparmor-example-deny-write (enforce)
```

Example

This example assumes you have already set up a cluster with AppArmor support.

First, we need to load the profile we want to use onto our nodes. This profile denies all file writes:

```
#include <tunables/global>

profile k8s-apparmor-example-deny-write flags=(attach_disconnected) {
    #include <abstractions/base>

    file,

    # Deny all file writes.
    deny /** w,
}
```

Since we don't know where the Pod will be scheduled, we'll need to load the profile on all our nodes. For this example we'll use SSH to install the profiles, but other approaches are discussed in [Setting up nodes with profiles](#).

```
NODES=(
    # The SSH-accessible domain names of your nodes
    gke-test-default-pool-239f5d02-gyn2.us-central1-a.my-k8s
    gke-test-default-pool-239f5d02-x1kf.us-central1-a.my-k8s
    gke-test-default-pool-239f5d02-xwux.us-central1-a.my-k8s)
for NODE in ${NODES[*]}; do ssh $NODE 'sudo apparmor_parser -q <<EOF
#include <tunables/global>

profile k8s-apparmor-example-deny-write flags=(attach_disconnected) {
    #include <abstractions/base>

    file,

    # Deny all file writes.
    deny /** w,
}
EOF'
done
```

Next, we'll run a simple "Hello AppArmor" pod with the deny-write profile:

```
podsecurity/hello-apparmor.yaml
```

(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/security/hello-apparmor.yaml>)

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-apparmor
  annotations:
    # Tell Kubernetes to apply the AppArmor profile "k8s-apparmor-example-deny-write".
    # Note that this is ignored if the Kubernetes node is not running version 1.4 or greater.
    container.apparmor.security.beta.kubernetes.io/hello: localhost/k8s-apparmor-example-deny-write
spec:
  containers:
  - name: hello
    image: busybox
    command: [ "sh", "-c", "echo 'Hello AppArmor!' && sleep 1h" ]
```

```
kubectl create -f ./hello-apparmor.yaml
```

If we look at the pod events, we can see that the Pod container was created with the AppArmor profile "k8s-apparmor-example-deny-write":

```
kubectl get events | grep hello-apparmor
```

14s	14s	1	hello-apparmor	Pod		Normal
14s	14s	1	hello-apparmor	Pod	spec.containers{hello}	Normal
13s	13s	1	hello-apparmor	Pod	spec.containers{hello}	Normal
13s	13s	1	hello-apparmor	Pod	spec.containers{hello}	Normal
13s	13s	1	hello-apparmor	Pod	spec.containers{hello}	Normal

We can verify that the container is actually running with that profile by checking its proc attr:

```
kubectl exec hello-apparmor cat /proc/1/attr/current
```

```
k8s-apparmor-example-deny-write (enforce)
```

Finally, we can see what happens if we try to violate the profile by writing to a file:

```
kubectl exec hello-apparmor touch /tmp/test
```

```
touch: /tmp/test: Permission denied
error: error executing remote command: command terminated with non-zero exit code: Error executing command
```

To wrap up, let's look at what happens if we try to specify a profile that hasn't been loaded:

```
kubectl create -f /dev/stdin <<EOF
```

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-apparmor-2
  annotations:
    container.apparmor.security.beta.kubernetes.io/hello: localhost/k8s-apparmor-example-allow-write
spec:
  containers:
  - name: hello
```

```
image: busybox
command: [ "sh", "-c", "echo 'Hello AppArmor!' && sleep 1h" ]
EOF
pod/hello-apparmor-2 created
```

```
kubectl describe pod hello-apparmor-2
```

```
Name:          hello-apparmor-2
Namespace:     default
Node:          gke-test-default-pool-239f5d02-x1kf/
Start Time:    Tue, 30 Aug 2016 17:58:56 -0700
Labels:        <none>
Annotations:   container.apparmor.security.beta.kubernetes.io/hello=localhost/k8s-apparmor-example-allow-write
Status:        Pending
Reason:        AppArmor
Message:       Pod Cannot enforce AppArmor: profile "k8s-apparmor-example-allow-write" is not loaded
IP:            <none>
Controllers:   <none>
Containers:
  hello:
    Container ID:
    Image:        busybox
    Image ID:
    Port:
    Command:
      sh
      -c
      echo 'Hello AppArmor!' && sleep 1h
    State:        Waiting
      Reason:      Blocked
    Ready:        False
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-dnz7v (ro)
Conditions:
  Type          Status
  Initialized    True
  Ready         False
  PodScheduled  True
Volumes:
  default-token-dnz7v:
    Type:      Secret (a volume populated by a Secret)
    SecretName: default-token-dnz7v
    Optional:   false
QoS Class:     BestEffort
Node-Selectors: <none>
Tolerations:   <none>
Events:
  FirstSeen    LastSeen    Count   From              SubobjectPath    Type    Reason
  -----
  23s          23s         1       {default-scheduler }              Normal    Scheduled
  23s          23s         1       {kubelet e2e-test-stclair-node-pool-t1f5}              Warning   Failed
```

Note the pod status is Pending, with a helpful error message: Pod Cannot enforce AppArmor: profile "k8s-apparmor-example-allow-write" is not loaded . An event was also recorded with the same message.

Administration

Setting up nodes with profiles

Kubernetes does not currently provide any native mechanisms for loading AppArmor profiles onto nodes. There are lots of ways to setup the profiles though, such as:

- Through a [DaemonSet](/docs/concepts/workloads/controllers/daemonset/) that runs a Pod on each node to ensure the correct profiles are loaded. An example implementation can be found [here](https://git.k8s.io/kubernetes/test/images/apparmor-loader) (https://git.k8s.io/kubernetes/test/images/apparmor-loader).
- At node initialization time, using your node initialization scripts (e.g. Salt, Ansible, etc.) or image.
- By copying the profiles to each node and loading them through SSH, as demonstrated in the [Example](#).

The scheduler is not aware of which profiles are loaded onto which node, so the full set of profiles must be loaded onto every node. An alternative approach is to add a node label for each profile (or class of profiles) on the node, and use a [node selector \(/docs/concepts/scheduling-eviction/assign-pod-node/\)](/docs/concepts/scheduling-eviction/assign-pod-node/) to ensure the Pod is run on a node with the required profile.

Restricting profiles with the PodSecurityPolicy

Note: PodSecurityPolicy is deprecated in Kubernetes v1.21, and will be removed in v1.25. See [PodSecurityPolicy documentation \(/docs/concepts/policy/pod-security-policy/\)](/docs/concepts/policy/pod-security-policy/) for more information.

If the PodSecurityPolicy extension is enabled, cluster-wide AppArmor restrictions can be applied. To enable the PodSecurityPolicy, the following flag must be set on the `apiserver` :

```
--enable-admission-plugins=PodSecurityPolicy[,others...]
```

The AppArmor options can be specified as annotations on the PodSecurityPolicy:

```
apparmor.security.beta.kubernetes.io/defaultProfileName: <profile_ref>
apparmor.security.beta.kubernetes.io/allowedProfileNames: <profile_ref>[,others...]
```

The default profile name option specifies the profile to apply to containers by default when none is specified. The allowed profile names option specifies a list of profiles that Pod containers are allowed to be run with. If both options are provided, the default must be allowed. The profiles are specified in the same format as on containers. See the [API Reference](#) for the full specification.

Disabling AppArmor

If you do not want AppArmor to be available on your cluster, it can be disabled by a command-line flag:

```
--feature-gates=AppArmor=false
```

When disabled, any Pod that includes an AppArmor profile will fail validation with a "Forbidden" error. Note that by default docker always enables the "docker-default" profile on non-privileged pods (if the AppArmor kernel module is enabled), and will continue to do so even if the feature-gate is disabled. The option to disable AppArmor will be removed when AppArmor graduates to general availability (GA).

Upgrading to Kubernetes v1.4 with AppArmor

No action is required with respect to AppArmor to upgrade your cluster to v1.4. However, if any existing pods had an AppArmor annotation, they will not go through validation (or PodSecurityPolicy admission). If permissive profiles are loaded on the nodes, a malicious user could pre-apply a permissive profile to escalate the pod privileges above the docker-default. If this is a concern, it is recommended to scrub the cluster of any pods containing an annotation with `apparmor.security.beta.kubernetes.io` .

Upgrade path to General Availability

When AppArmor is ready to be graduated to general availability (GA), the options currently specified through annotations will be converted to fields. Supporting all the upgrade and downgrade paths through the transition is very nuanced, and will be explained in detail when the transition occurs. We will commit to supporting both fields and annotations for at least 2 releases, and will explicitly reject the annotations for at least 2 releases after that.

Authoring Profiles

Getting AppArmor profiles specified correctly can be a tricky business. Fortunately there are some tools to help with that:

- `aa-genprof` and `aa-logprof` generate profile rules by monitoring an application's activity and logs, and admitting the actions it takes. Further instructions are provided by the [AppArmor documentation \(https://gitlab.com/apparmor/apparmor/wikis/Profiling_with_tools\)](https://gitlab.com/apparmor/apparmor/wikis/Profiling_with_tools).
- [bane \(https://github.com/jfrazelle/bane\)](https://github.com/jfrazelle/bane) is an AppArmor profile generator for Docker that uses a simplified profile language.

It is recommended to run your application through Docker on a development workstation to generate the profiles, but there is nothing preventing running the tools on the Kubernetes node where your Pod is running.

To debug problems with AppArmor, you can check the system logs to see what, specifically, was denied. AppArmor logs verbose messages to `dmesg`, and errors can usually be found in the system logs or through `journalctl`. More information is provided in [AppArmor failures \(https://gitlab.com/apparmor/apparmor/wikis/AppArmor_Failures\)](https://gitlab.com/apparmor/apparmor/wikis/AppArmor_Failures).

API Reference

Pod Annotation

Specifying the profile a container will run with:

- **key:** `container.apparmor.security.beta.kubernetes.io/<container_name>` Where `<container_name>` matches the name of a container in the Pod. A separate profile can be specified for each container in the Pod.
- **value:** a profile reference, described below

Profile Reference

- `runtime/default` : Refers to the default runtime profile.
 - Equivalent to not specifying a profile (without a PodSecurityPolicy default), except it still requires AppArmor to be enabled.
 - For Docker, this resolves to the `docker-default` (<https://docs.docker.com/engine/security/apparmor/>) profile for non-privileged containers, and `unconfined` (no profile) for privileged containers.
- `localhost/<profile_name>` : Refers to a profile loaded on the node (localhost) by name.
 - The possible profile names are detailed in the [core policy reference \(https://gitlab.com/apparmor/apparmor/wikis/AppArmor_Core_Policy_Reference#profile-names-and-attachment-specifications\)](https://gitlab.com/apparmor/apparmor/wikis/AppArmor_Core_Policy_Reference#profile-names-and-attachment-specifications).
- `unconfined` : This effectively disables AppArmor on the container.

Any other profile reference format is invalid.

PodSecurityPolicy Annotations

Specifying the default profile to apply to containers when none is provided:

- **key:** `apparmor.security.beta.kubernetes.io/defaultProfileName`
- **value:** a profile reference, described above

Specifying the list of profiles Pod containers is allowed to specify:

- **key:** `apparmor.security.beta.kubernetes.io/allowedProfileNames`
- **value:** a comma-separated list of profile references (described above)
 - Although an escaped comma is a legal character in a profile name, it cannot be explicitly allowed here.

What's next

Additional resources:

- [Quick guide to the AppArmor profile language \(https://gitlab.com/apparmor/apparmor/wikis/QuickProfileLanguage\)](https://gitlab.com/apparmor/apparmor/wikis/QuickProfileLanguage)
- [AppArmor core policy reference \(https://gitlab.com/apparmor/apparmor/wikis/Policy_Layout\)](https://gitlab.com/apparmor/apparmor/wikis/Policy_Layout)

6.2 - Restrict a Container's Syscalls with seccomp

FEATURE STATE: Kubernetes v1.19 [stable]

Seccomp stands for secure computing mode and has been a feature of the Linux kernel since version 2.6.12. It can be used to sandbox the privileges of a process, restricting the calls it is able to make from userspace into the kernel. Kubernetes lets you automatically apply seccomp profiles loaded onto a Node to your Pods and containers.

Identifying the privileges required for your workloads can be difficult. In this tutorial, you will go through how to load seccomp profiles into a local Kubernetes cluster, how to apply them to a Pod, and how you can begin to craft profiles that give only the necessary privileges to your container processes.

Objectives

- Learn how to load seccomp profiles on a node
- Learn how to apply a seccomp profile to a container
- Observe auditing of syscalls made by a container process
- Observe behavior when a missing profile is specified
- Observe a violation of a seccomp profile
- Learn how to create fine-grained seccomp profiles
- Learn how to apply a container runtime default seccomp profile

Before you begin

Your Kubernetes server must be version v1.22. To check the version, enter `kubectl version`. In order to complete all steps in this tutorial, you must install [kind](https://kind.sigs.k8s.io/docs/user/quick-start/) (<https://kind.sigs.k8s.io/docs/user/quick-start/>) and [kubectl](https://docs/tasks/tools/) ([/docs/tasks/tools/](https://docs/tasks/tools/)). This tutorial will show examples both alpha (new in v1.22) and generally available seccomp functionality. You should make sure that your cluster is [configured correctly](https://kind.sigs.k8s.io/docs/user/quick-start/#setting-kubernetes-version) (<https://kind.sigs.k8s.io/docs/user/quick-start/#setting-kubernetes-version>) for the version you are using.

Note: It is not possible to apply a seccomp profile to a container running with `privileged: true` set in the container's `securityContext`. Privileged containers always run as `Unconfined`.

Enable the use of `RuntimeDefault` as the default seccomp profile for all workloads

FEATURE STATE: Kubernetes v1.22 [alpha]

`SeccompDefault` is an optional kubelet [feature gate](https://docs/reference/command-line-tools-reference/feature-gates) ([/docs/reference/command-line-tools-reference/feature-gates](https://docs/reference/command-line-tools-reference/feature-gates)) as well as corresponding `--seccomp-default` [command line flag](https://docs/reference/command-line-tools-reference/kubelet) ([/docs/reference/command-line-tools-reference/kubelet](https://docs/reference/command-line-tools-reference/kubelet)). Both have to be enabled simultaneously to use the feature.

If enabled, the kubelet will use the `RuntimeDefault` seccomp profile by default, which is defined by the container runtime, instead of using the `Unconfined` (seccomp disabled) mode. The default profiles aim to provide a strong set of security defaults while preserving the functionality of the workload. It is possible that the default profiles differ between container runtimes and their release versions, for example when comparing those from CRI-O and containerd.

Some workloads may require a lower amount of syscall restrictions than others. This means that they can fail during runtime even with the `RuntimeDefault` profile. To mitigate such a failure, you can:

- Run the workload explicitly as `Unconfined`.
- Disable the `SeccompDefault` feature for the nodes. Also making sure that workloads get scheduled on nodes where the feature is disabled.

- Create a custom seccomp profile for the workload.

If you were introducing this feature into production-like cluster, the Kubernetes project recommends that you enable this feature gate on a subset of your nodes and then test workload execution before rolling the change out cluster-wide.

More detailed information about a possible upgrade and downgrade strategy can be found in the [related Kubernetes Enhancement Proposal \(KEP\)](https://github.com/kubernetes/enhancements/tree/a70cc18/keps/sig-node/2413-seccomp-by-default#upgrade--downgrade-strategy) (<https://github.com/kubernetes/enhancements/tree/a70cc18/keps/sig-node/2413-seccomp-by-default#upgrade--downgrade-strategy>).

Since the feature is in alpha state it is disabled per default. To enable it, pass the flags `--feature-gates=SeccompDefault=true --seccomp-default` to the `kubelet` CLI or enable it via the [kubelet configuration file](#) ([/docs/tasks/administer-cluster/kubelet-config-file/](#)). To enable the feature gate in `kind` (<https://kind.sigs.k8s.io>), ensure that `kind` provides the minimum required Kubernetes version and enables the `SeccompDefault` feature [in the kind configuration](#) (<https://kind.sigs.k8s.io/docs/user/quick-start/#enable-feature-gates-in-your-cluster>):

```
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
featureGates:
  SeccompDefault: true
```

Create Seccomp Profiles

The contents of these profiles will be explored later on, but for now go ahead and download them into a directory named `profiles/` so that they can be loaded into the cluster.

[audit.json](#)

[violation.json](#)


[fine-grained.json](#)

[pods/security/seccomp/profiles/audit.json](#)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/security/seccomp/profiles/audit.json>)

```
{
  "defaultAction": "SCMP_ACT_LOG"
}
```

Create a Local Kubernetes Cluster with Kind

For simplicity, `kind` (<https://kind.sigs.k8s.io/>) can be used to create a single node cluster with the seccomp profiles loaded. Kind runs Kubernetes in Docker, so each node of the cluster is a container. This allows for files to be mounted in the filesystem of each container similar to loading files onto a node.

[pods/security/seccomp/kind.yaml](https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/security/seccomp/kind.yaml)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/security/seccomp/kind.yaml>) 

```
apiVersion: kind.x-k8s.io/v1alpha4
kind: Cluster
nodes:
- role: control-plane
  extraMounts:
  - hostPath: "./profiles"
    containerPath: "/var/lib/kubelet/seccomp/profiles"
```

Download the example above, and save it to a file named `kind.yaml` . Then create the cluster with the configuration.

```
kind create cluster --config=kind.yaml
```

Once the cluster is ready, identify the container running as the single node cluster:

```
docker ps
```

You should see output indicating that a container is running with name `kind-control-plane` .

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
6a96207fed4b	kindest/node:v1.18.2	"/usr/local/bin/entr..."	27 seconds ago	Up 24

If observing the filesystem of that container, one should see that the `profiles/` directory has been successfully loaded into the default seccomp path of the kubelet. Use `docker exec` to run a command in the Pod:

```
docker exec -it 6a96207fed4b ls /var/lib/kubelet/seccomp/profiles
```

```
audit.json  fine-grained.json  violation.json
```

Create a Pod with a seccomp profile for syscall auditing

To start off, apply the `audit.json` profile, which will log all syscalls of the process, to a new Pod.

Download the correct manifest for your Kubernetes version:

[v1.19 or Later \(GA\)](#)

[Pre-v1.19 \(alpha\)](#)

[pods/security/seccomp/ga/audit-pod.yaml](https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/security/seccomp/ga/audit-pod.yaml)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/security/seccomp/ga/audit-pod.yaml>)

```
apiVersion: v1
kind: Pod
metadata:
  name: audit-pod
  labels:
    app: audit-pod
spec:
  securityContext:
    seccompProfile:
      type: Localhost
      localhostProfile: profiles/audit.json
  containers:
  - name: test-container
    image: hashicorp/http-echo:0.2.3
    args:
    - "-text=just made some syscalls!"
    securityContext:
      allowPrivilegeEscalation: false
```

Create the Pod in the cluster:

```
kubectl apply -f audit-pod.yaml
```

This profile does not restrict any syscalls, so the Pod should start successfully.

```
kubectl get pod/audit-pod
```

NAME	READY	STATUS	RESTARTS	AGE
audit-pod	1/1	Running	0	30s

In order to be able to interact with this endpoint exposed by this container, create a NodePort Service that allows access to the endpoint from inside the kind control plane container.

```
kubectl expose pod/audit-pod --type NodePort --port 5678
```

Check what port the Service has been assigned on the node.

```
kubectl get svc/audit-pod
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
audit-pod	NodePort	10.111.36.142	<none>	5678:32373/TCP	72s

Now you can `curl` the endpoint from inside the kind control plane container at the port exposed by this Service. Use `docker exec` to run a command in the Pod:

```
docker exec -it 6a96207fed4b curl localhost:32373
```

```
just made some syscalls!
```

You can see that the process is running, but what syscalls did it actually make? Because this Pod is running in a local cluster, you should be able to see those in `/var/log/syslog`. Open up a new terminal window and `tail` the output for calls from `http-echo`:

```
tail -f /var/log/syslog | grep 'http-echo'
```

You should already see some logs of syscalls made by `http-echo`, and if you `curl` the endpoint in the control plane container you will see more written.

```
Jul 6 15:37:40 my-machine kernel: [369128.669452] audit: type=1326 audit(1594067860.484:14530)
Jul 6 15:37:40 my-machine kernel: [369128.669453] audit: type=1326 audit(1594067860.484:14531)
Jul 6 15:37:40 my-machine kernel: [369128.669455] audit: type=1326 audit(1594067860.484:14532)
Jul 6 15:37:40 my-machine kernel: [369128.669456] audit: type=1326 audit(1594067860.484:14533)
Jul 6 15:37:40 my-machine kernel: [369128.669517] audit: type=1326 audit(1594067860.484:14540)
Jul 6 15:37:40 my-machine kernel: [369128.669519] audit: type=1326 audit(1594067860.484:14541)
Jul 6 15:38:40 my-machine kernel: [369188.671648] audit: type=1326 audit(1594067920.488:14559)
Jul 6 15:38:40 my-machine kernel: [369188.671726] audit: type=1326 audit(1594067920.488:14560)
```

You can begin to understand the syscalls required by the `http-echo` process by looking at the `syscall=` entry on each line. While these are unlikely to encompass all syscalls it uses, it can serve as a basis for a seccomp profile for this container.

Clean up that Pod and Service before moving to the next section:

```
kubectl delete pod/audit-pod
kubectl delete svc/audit-pod
```

Create Pod with seccomp Profile that Causes Violation

For demonstration, apply a profile to the Pod that does not allow for any syscalls.

Download the correct manifest for your Kubernetes version:

v1.19 or Later (GA).

Pre-v1.19 (alpha)

[pods/security/seccomp/ga/violation-pod.yaml](https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/security/seccomp/ga/violation-pod.yaml)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/security/seccomp/ga/violation-pod.yaml>)

```
apiVersion: v1
kind: Pod
metadata:
  name: violation-pod
```

```
labels:
  app: violation-pod
spec:
  securityContext:
    seccompProfile:
      type: Localhost
      localhostProfile: profiles/violation.json
  containers:
    - name: test-container
      image: hashicorp/http-echo:0.2.3
      args:
        - "-text=just made some syscalls!"
      securityContext:
        allowPrivilegeEscalation: false
```

```
kubectl apply -f violation-pod.yaml
```

```
kubectl get pod/violation-pod
```

NAME	READY	STATUS	RESTARTS	AGE
violation-pod	0/1	CrashLoopBackOff	1	6s

```
kubectl delete pod/violation-pod
kubectl delete svc/violation-pod
```

If you take a look at the `fine-pod.json`, you will notice some of the syscalls seen in the first example where the profile set `"defaultAction": "SCMP_ACT_LOG"`. Now the profile is setting `"defaultAction": "SCMP_ACT_ERRNO"`, but explicitly allowing a set of syscalls in the `"action": "SCMP_ACT_ALLOW"` block. Ideally, the container will run successfully and you will see no messages sent to `syslog`.

[pods/security/seccomp/ga/fine-pod.yaml](https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/security/seccomp/ga/fine-pod.yaml)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/security/seccomp/ga/fine-pod.yaml>)

```
apiVersion: v1
kind: Pod
metadata:
  name: fine-pod
  labels:
    app: fine-pod
spec:
  securityContext:
    seccompProfile:
      type: Localhost
      localhostProfile: profiles/fine-grained.json
  containers:
  - name: test-container
    image: hashicorp/http-echo:0.2.3
    args:
    - "-text=just made some syscalls!"
    securityContext:
      allowPrivilegeEscalation: false
```

Create the Pod in your cluster:

```
kubectl apply -f fine-pod.yaml
```

The Pod should start successfully.

```
kubectl get pod/fine-pod
```

NAME	READY	STATUS	RESTARTS	AGE
fine-pod	1/1	Running	0	30s

Open up a new terminal window and `tail` the output for calls from `http-echo` :

```
tail -f /var/log/syslog | grep 'http-echo'
```

Expose the Pod with a NodePort Service:

```
kubectl expose pod/fine-pod --type NodePort --port 5678
```

Check what port the Service has been assigned on the node:

```
kubectl get svc/fine-pod
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
fine-pod	NodePort	10.111.36.142	<none>	5678:32373/TCP	72s

`curl` the endpoint from inside the kind control plane container:

```
docker exec -it 6a96207fed4b curl localhost:32373
```

```
just made some syscalls!
```

You should see no output in the `syslog` because the profile allowed all necessary syscalls and specified that an error should occur if one outside of the list is invoked. This is an ideal situation from a security perspective, but required some effort in analyzing the program. It would be nice if there was a simple way to get closer to this security without requiring as much effort.

Clean up that Pod and Service before moving to the next section:

```
kubectl delete pod/fine-pod
kubectl delete svc/fine-pod
```

Create Pod that uses the Container Runtime Default seccomp Profile

Most container runtimes provide a sane set of default syscalls that are allowed or not. The defaults can easily be applied in Kubernetes by using the `runtime/default` annotation or setting the seccomp type in the security context of a pod or container to `RuntimeDefault`.

Download the correct manifest for your Kubernetes version:

[v1.19 or Later \(GA\)](#)

[Pre-v1.19 \(alpha\)](#)

[pods/security/seccomp/ga/default-pod.yaml](#)
(<https://raw.githubusercontent.com/kubernetes/website/main/content/en/examples/pods/security/seccomp/ga/default-pod.yaml>)

```
apiVersion: v1
kind: Pod
metadata:
  name: audit-pod
  labels:
    app: audit-pod
spec:
  securityContext:
    seccompProfile:
      type: RuntimeDefault
  containers:
    - name: test-container
      image: hashicorp/http-echo:0.2.3
      args:
        - "-text=just made some syscalls!"
      securityContext:
        allowPrivilegeEscalation: false
```

The default seccomp profile should provide adequate access for most workloads.

What's next

Additional resources:

- [A seccomp Overview \(https://lwn.net/Articles/656307/\)](https://lwn.net/Articles/656307/)
- [Seccomp Security Profiles for Docker \(https://docs.docker.com/engine/security/seccomp/\)](https://docs.docker.com/engine/security/seccomp/)

7 - Services

7.1 - Using Source IP

Applications running in a Kubernetes cluster find and communicate with each other, and the outside world, through the Service abstraction. This document explains what happens to the source IP of packets sent to different types of Services, and how you can toggle this behavior according to your needs.

Before you begin

Terminology

This document makes use of the following terms:

NAT (https://en.wikipedia.org/wiki/Network_address_translation)

network address translation

Source NAT (https://en.wikipedia.org/wiki/Network_address_translation#SNAT)

replacing the source IP on a packet; in this page, that usually means replacing with the IP address of a node.

Destination NAT (https://en.wikipedia.org/wiki/Network_address_translation#DNAT)

replacing the destination IP on a packet; in this page, that usually means replacing with the IP address of a Pod (</docs/concepts/workloads/pods/>)

VIP (</docs/concepts/services-networking/service/#virtual-ips-and-service-proxies>)

a virtual IP address, such as the one assigned to every Service (</docs/concepts/services-networking/service/>) in Kubernetes

kube-proxy (</docs/concepts/services-networking/service/#virtual-ips-and-service-proxies>)

a network daemon that orchestrates Service VIP management on every node

Prerequisites

You need to have a Kubernetes cluster, and the kubectl command-line tool must be configured to communicate with your cluster. It is recommended to run this tutorial on a cluster with at least two nodes that are not acting as control plane hosts. If you do not already have a cluster, you can create one by using minikube (https://minikube.sigs.k8s.io/docs/tutorials/multi_node/) or you can use one of these Kubernetes playgrounds:

- [Katacoda](https://www.katacoda.com/courses/kubernetes/playground) (<https://www.katacoda.com/courses/kubernetes/playground>)
- [Play with Kubernetes](http://labs.play-with-k8s.com/) (<http://labs.play-with-k8s.com/>)

The examples use a small nginx webserver that echoes back the source IP of requests it receives through an HTTP header. You can create it as follows:

```
kubectl create deployment source-ip-app --image=k8s.gcr.io/echoserver:1.4
```

The output is:

```
deployment.apps/source-ip-app created
```

Objectives

- Expose a simple application through various types of Services
- Understand how each Service type handles source IP NAT
- Understand the tradeoffs involved in preserving source IP

Source IP for Services with Type=ClusterIP

Packets sent to ClusterIP from within the cluster are never source NAT'd if you're running kube-proxy in iptables mode (</docs/concepts/services-networking/service/#proxy-mode-iptables>), (the default). You can query the kube-proxy mode by fetching `http://localhost:10249/proxyMode` on the node where kube-proxy is running.

```
kubectl get nodes
```

The output is similar to this:

NAME		STATUS	ROLES	AGE	VERSION
kubernetes-node-6jst	Ready	<none>	2h	v1.13.0	
kubernetes-node-cx31	Ready	<none>	2h	v1.13.0	
kubernetes-node-jj1t	Ready	<none>	2h	v1.13.0	

Get the proxy mode on one of the nodes (kube-proxy listens on port 10249):

```
# Run this in a shell on the node you want to query.
curl http://localhost:10249/proxyMode
```

The output is:

```
iptables
```

You can test source IP preservation by creating a Service over the source IP app:

```
kubectl expose deployment source-ip-app --name=clusterip --port=80 --target-port=8080
```

The output is:

```
service/clusterip exposed
```

```
kubectl get svc clusterip
```

The output is similar to:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
clusterip	ClusterIP	10.0.170.92	<none>	80/TCP	51s

And hitting the ClusterIP from a pod in the same cluster:

```
kubectl run busybox -it --image=busybox --restart=Never --rm
```

The output is similar to this:

```
Waiting for pod default/busybox to be running, status is Pending, pod ready: false
If you don't see a command prompt, try pressing enter.
```

You can then run a command inside that Pod:

```
# Run this inside the terminal from "kubectl run"
ip addr
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
3: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1460 qdisc noqueue
    link/ether 0a:58:0a:f4:03:08 brd ff:ff:ff:ff:ff:ff
    inet 10.244.3.8/24 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::188a:84ff:feb0:26a5/64 scope link
        valid_lft forever preferred_lft forever
```

...then use `wget` to query the local webserver

```
# Replace "10.0.170.92" with the IPv4 address of the Service named "clusterip"
wget -q0 - 10.0.170.92
```

```
CLIENT VALUES:
client_address=10.244.3.8
command=GET
...
```

The `client_address` is always the client pod's IP address, whether the client pod and server pod are in the same node or in different nodes.

Source IP for Services with Type=NodePort

Packets sent to Services with `Type=NodePort` (</docs/concepts/services-networking/service/#nodeport>) are source NAT'd by default. You can test this by creating a NodePort Service:

```
kubectl expose deployment source-ip-app --name=nodeport --port=80 --target-port=8080 --type=N
```

The output is:

```
service/nodeport exposed
```

```
NODEPORT=$(kubectl get -o jsonpath="{.spec.ports[0].nodePort}" services nodeport)
NODES=$(kubectl get nodes -o jsonpath='{ $$.items[*].status.addresses[?(@.type=="InternalIP")]
```

If you're running on a cloud provider, you may need to open up a firewall-rule for the `nodes:nodeport` reported above. Now you can try reaching the Service from outside the cluster through the node port allocated above.

```
for node in $NODES; do curl -s $node:$NODEPORT | grep -i client_address; done
```

The output is similar to:

```
client_address=10.180.1.1
client_address=10.240.0.5
client_address=10.240.0.3
```

Note that these are not the correct client IPs, they're cluster internal IPs. This is what happens:

- Client sends packet to `node2:nodePort`
- `node2` replaces the source IP address (SNAT) in the packet with its own IP address
- `node2` replaces the destination IP on the packet with the pod IP
- packet is routed to node 1, and then to the endpoint

- the pod's reply is routed back to node2
- the pod's reply is sent back to the client

Visually:

```
graph LR; client(client)-->node2[Node 2]; node2-->client; node2-. SNAT .->node1[Node 1]; node1-. SNAT .->node2; node1-->endpoint(endpoint); classDef plain fill:#ddd,stroke:#fff,stroke-width:4px,color:#000; classDef k8s fill:#326ce5,stroke:#fff,stroke-width:4px,color:#fff; class node1,node2,endpoint k8s; class client plain;
```

To avoid this, Kubernetes has a feature to preserve the client source IP (</docs/tasks/access-application-cluster/create-external-load-balancer/#preserving-the-client-source-ip>). If you set `service.spec.externalTrafficPolicy` to the value `Local`, kube-proxy only proxies proxy requests to local endpoints, and does not forward traffic to other nodes. This approach preserves the original source IP address. If there are no local endpoints, packets sent to the node are dropped, so you can rely on the correct source-ip in any packet processing rules you might apply a packet that make it through to the endpoint.

Set the `service.spec.externalTrafficPolicy` field as follows:

```
kubectl patch svc nodeport -p '{"spec":{"externalTrafficPolicy":"Local"}}'
```

The output is:

```
service/nodeport patched
```

Now, re-run the test:

```
for node in $NODES; do curl --connect-timeout 1 -s $node:$NODEPORT | grep -i client_address; done
```

The output is similar to:

```
client_address=198.51.100.79
```

Note that you only got one reply, with the *right* client IP, from the one node on which the endpoint pod is running.

This is what happens:

- client sends packet to `node2:nodePort`, which doesn't have any endpoints
- packet is dropped
- client sends packet to `node1:nodePort`, which *does* have endpoints
- node1 routes packet to endpoint with the correct source IP

Visually:

```
graph TD; client --> node1[Node 1]; client(client) --x node2[Node 2]; node1 --> endpoint(endpoint); endpoint --> node1; classDef plain fill:#ddd,stroke:#fff,stroke-width:4px,color:#000; classDef k8s fill:#326ce5,stroke:#fff,stroke-width:4px,color:#fff; class node1,node2,endpoint k8s; class client plain;
```

Source IP for Services with Type=LoadBalancer

Packets sent to Services with `Type=LoadBalancer` (</docs/concepts/services-networking/service/#loadbalancer>) are source NAT'd by default, because all schedulable Kubernetes nodes in the `Ready` state are eligible for load-balanced traffic. So if packets arrive at a node without an endpoint, the system proxies it to a node *with* an endpoint, replacing the source IP on the packet with the IP of the node (as described in the previous section).

You can test this by exposing the source-ip-app through a load balancer:

```
kubectl expose deployment source-ip-app --name=loadbalancer --port=80 --target-port=8080 --ty
```

The output is:

```
service/loadbalancer exposed
```

Print out the IP addresses of the Service:

```
kubectl get svc loadbalancer
```

The output is similar to this:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
loadbalancer	LoadBalancer	10.0.65.118	203.0.113.140	80/TCP	5m

Next, send a request to this Service's external-ip:

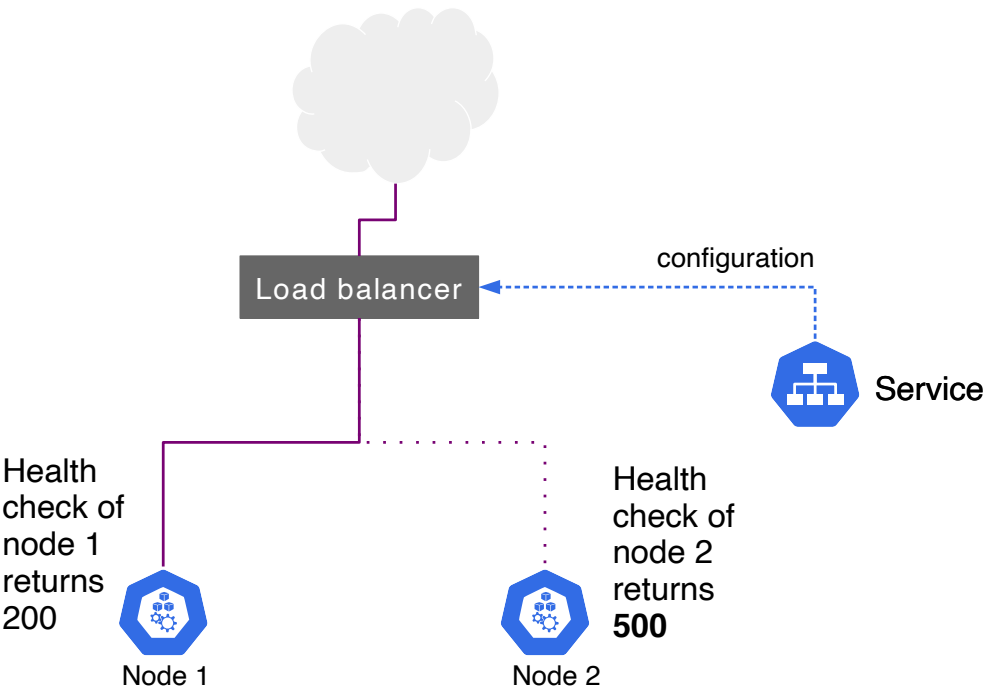
```
curl 203.0.113.140
```

The output is similar to this:

```
CLIENT VALUES:
client_address=10.240.0.5
...
```

However, if you're running on Google Kubernetes Engine/GCE, setting the same `service.spec.externalTrafficPolicy` field to `Local` forces nodes *without* Service endpoints to remove themselves from the list of nodes eligible for loadbalanced traffic by deliberately failing health checks.

Visually:



You can test this by setting the annotation:

```
kubectl patch svc loadbalancer -p '{"spec":{"externalTrafficPolicy":"Local"}}'
```

You should immediately see the `service.spec.healthCheckNodePort` field allocated by Kubernetes:

```
kubectl get svc loadbalancer -o yaml | grep -i healthCheckNodePort
```

The output is similar to this:

```
healthCheckNodePort: 32122
```

The `service.spec.healthCheckNodePort` field points to a port on every node serving the health check at `/healthz`. You can test this:

```
kubectl get pod -o wide -l run=source-ip-app
```

The output is similar to this:

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
source-ip-app-826191075-qehz4	1/1	Running	0	20h	10.180.1.136	kubern...

Use `curl` to fetch the `/healthz` endpoint on various nodes:

```
# Run this locally on a node you choose
curl localhost:32122/healthz
```

```
1 Service Endpoints found
```

On a different node you might get a different result:

```
# Run this locally on a node you choose
curl localhost:32122/healthz
```

```
No Service Endpoints Found
```

A controller running on the [control plane \(/docs/reference/glossary/?all=true#term-control-plane\)](/docs/reference/glossary/?all=true#term-control-plane) is responsible for allocating the cloud load balancer. The same controller also allocates HTTP health checks pointing to this port/path on each node. Wait about 10 seconds for the 2 nodes without endpoints to fail health checks, then use `curl` to query the IPv4 address of the load balancer:

```
curl 203.0.113.140
```

The output is similar to this:

```
CLIENT VALUES:
client_address=198.51.100.79
...
```

Cross-platform support

Only some cloud providers offer support for source IP preservation through Services with `Type=LoadBalancer`. The cloud provider you're running on might fulfill the request for a loadbalancer in a few different ways:

1. With a proxy that terminates the client connection and opens a new connection to your nodes/endpoints. In such cases the source IP will always be that of the cloud LB, not that of the client.
2. With a packet forwarder, such that requests from the client sent to the loadbalancer VIP end up at the node with the source IP of the client, not an intermediate proxy.

Load balancers in the first category must use an agreed upon protocol between the loadbalancer and backend to communicate the true client IP such as the HTTP Forwarded (<https://tools.ietf.org/html/rfc7239#section-5.2>) or X-FORWARDED-FOR (<https://en.wikipedia.org/wiki/X-Forwarded-For>) headers, or the proxy protocol (<https://www.haproxy.org/download/1.8/doc/proxy-protocol.txt>). Load balancers in the second category can leverage the feature described above by creating an HTTP health check pointing at the port stored in the `service.spec.healthCheckNodePort` field on the Service.

Cleaning up

Delete the Services:

```
kubectl delete svc -l app=source-ip-app
```

Delete the Deployment, ReplicaSet and Pod:

```
kubectl delete deployment source-ip-app
```

What's next

- Learn more about connecting applications via services (</docs/concepts/services-networking/connect-applications-service/>).
- Read how to Create an External Load Balancer (</docs/tasks/access-application-cluster/create-external-load-balancer/>).